

Operating Systems Lab:

Context Switch in Pintos



Youjip Won

Overview

- ▣ Process structure
- ▣ Process state
- ▣ Process context
- ▣ `schedule()`
- ▣ switch thread

Process structure: struct thread

pintos/src/threads/thread.h

```
struct thread
{
    /* Owned by thread.c. */
    tid_t tid;
    enum thread_status status;
    char name[16];
    uint8_t *stack;
    int priority;
    struct list_elem allelem;
/

    /* Shared between thread.c and synch.c. */
    struct list_elem elem;

#ifdef USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir;
#endif

    /* Owned by thread.c. */
    unsigned magic;
};
```

enum thread_status

```
{
    THREAD_RUNNING, /* Running thread. */
    THREAD_READY,   /* Not running but ready to run. */
    THREAD_BLOCKED, /* Waiting for an event to trigger. */
    THREAD_DYING    /* About to be destroyed. */
};
```

```
/* Thread identifier. */
/* Thread state. */
/* Name (for debugging purposes). */
/* Saved stack pointer. */
/* Priority. */
/* List element for all threads list. */

/* List element. */

/* Page directory. */

/* Detects stack overflow. */
```

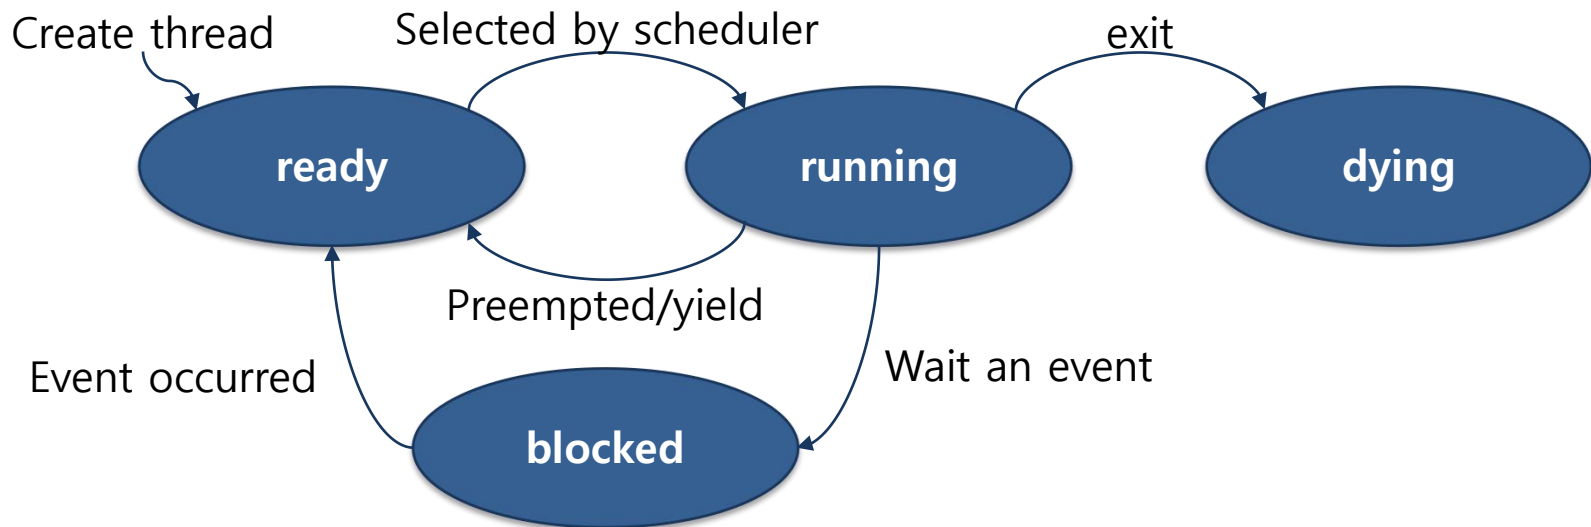
Process state

▣ Process state

`/* States in a thread's life cycle. */`

`enum thread_status`

```
{  
    THREAD_RUNNING,    /* Running thread. */  
    THREAD_READY,      /* Not running but ready to run. */  
    THREAD_BLOCKED,    /* Waiting for an event to trigger. */  
    THREAD_DYING       /* About to be destroyed. */  
};
```



Creating a thread

```
tid_t
```

```
thread_create (const char *name, int priority,  
               thread_func *function, void *aux)
```

- ▣ Creates a new kernel thread named NAME with the given PRIORITY, which executes FUNCTION passing AUX as the argument, and adds it to the ready queue. Returns the thread identifier for the new thread, or TID_ERROR if creation fails.

Creating a thread

Allocating a struct thread object

```
static void
init_thread (struct thread *t, const char *name, int priority)
{
    ASSERT (t != NULL);
    ASSERT (PRI_MIN <= priority && priority <= PRI_MAX);
    ASSERT (name != NULL);

    memset (t, 0, sizeof *t);
    t->status = THREAD_BLOCKED;
    strncpy (t->name, name, sizeof t->name);
    t->stack = (uint8_t *) t + PGSIZE;
    t->priority = priority;
    t->magic = THREAD_MAGIC;
    list_push_back (&all_list, &t->allelem);
}
```

Process list

pintos/src/threads/thread.c


```
/* List of processes in THREAD_READY state, that is, processes that
are ready to run but not actually running. */
static struct list ready_list;

/* List of all processes. Processes are added to this list
when they are first scheduled and removed when they exit. */
static struct list all_list;
```

- ▣ ready_list: a set of threads that are ready for execution
- ▣ all_list: A set of all threads in the system.

Creating a thread

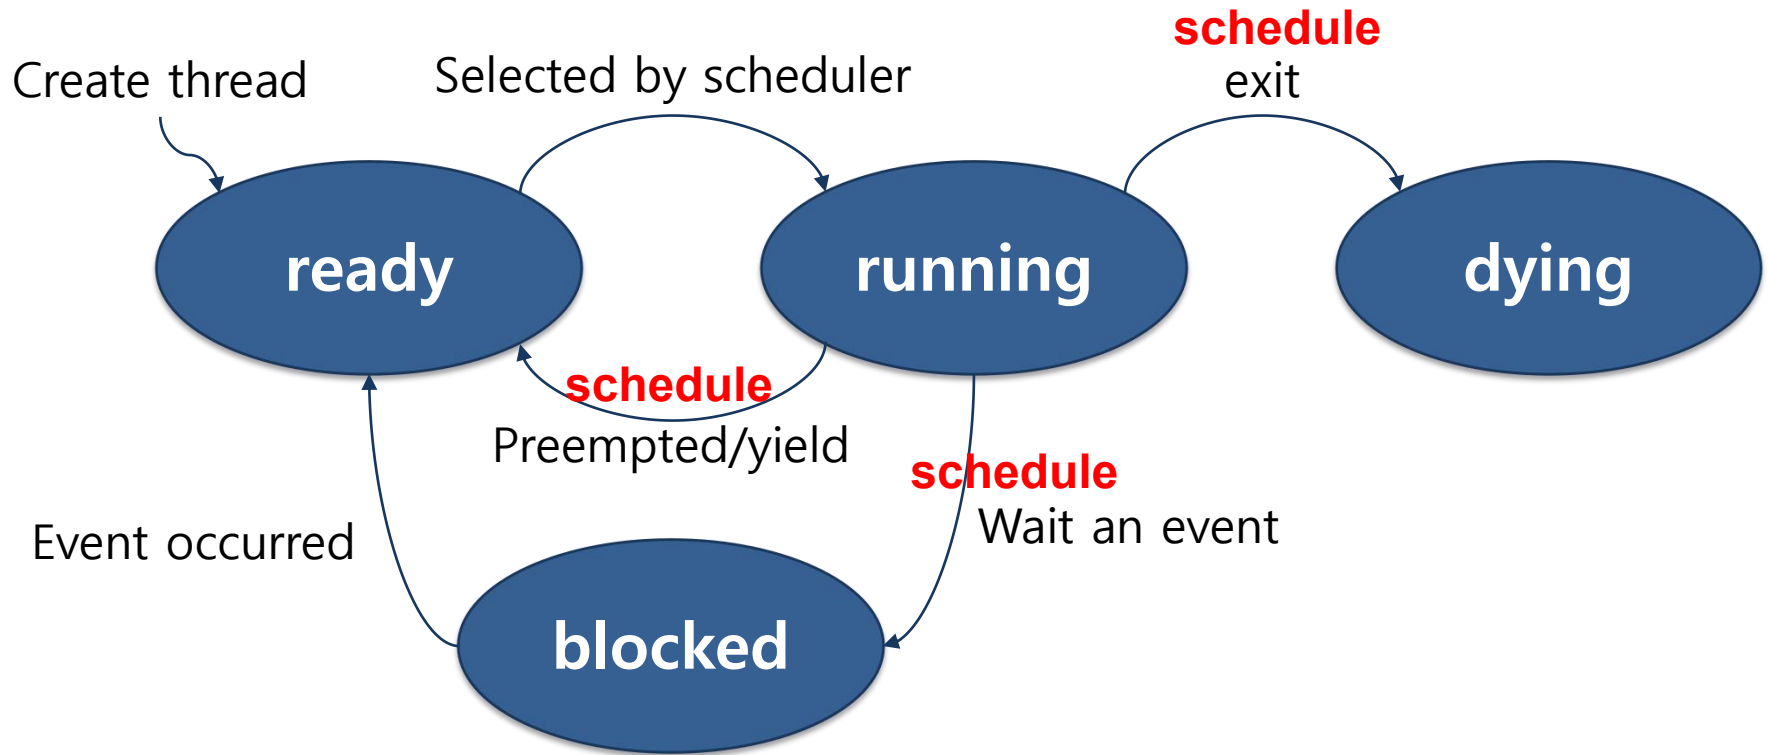
- ▣ Newly created thread is inserted at the end of the list `all_list` (`list_push_back()`)
 - ◆ `thread_create()`
 - ◆ `init_thread`
- ▣ After the thread is ready, it is inserted at the `ready_list`.
 - ◆ `thread_unblock`



`schedule()`

- ❑ Schedule a new process.
 - ◆ Get the currently running process.
 - ◆ Get the next process to run.
 - ◆ Switch context from current to next.
- ❑ Who calls `schedule()`?
 - ◆ `exit`, `block`, `yield`
 - ◆ *Or pre-empted, (when the time quantum expires...)*
- ❑ Before calling `schedule`
 - ◆ Disable interrupt.
 - ◆ Change the state of the running thread from running to something else.

schedule()



```

void
thread_block (void)
{
    ASSERT (!intr_context ());
    ASSERT (intr_get_level () == INTR_OFF);

    thread_current ()->status = THREAD_BLOCKED;
    schedule ();
}

```

/ Yields the CPU. The current thread is not put to sleep and may be scheduled again immediately at the scheduler's whim. */*

```

void
thread_yield (void)
{
    struct thread *cur = thread_current ();
    enum intr_level old_level;

    ASSERT (!intr_context ());

    old_level = intr_disable ();
    if (cur != idle_thread)
        list_push_back (&ready_list, &cur->elem);
    cur->status = THREAD_READY;
    schedule ();
    intr_set_level (old_level);
}

```

```

void
thread_exit (void)
{
    ASSERT (!intr_context ());

#ifdef USERPROG
    process_exit ();
#endif

    /* Remove thread from all threads list, set our status to dying,
       and schedule another process. That process will destroy us
       when it calls thread_schedule_tail(). */
    intr_disable ();
    list_remove (&thread_current()->allelem);
    thread_current ()->status = THREAD_DYING;
    schedule ();
    NOT_REACHED ();
}

```


schedule (void)

```
static void
schedule (void)
{
    struct thread *cur = running_thread ();
    struct thread *next = next_thread_to_run ();
    struct thread *prev = NULL;

    ASSERT (intr_get_level () == INTR_OFF);
    ASSERT (cur->status != THREAD_RUNNING);
    ASSERT (is_thread (next));

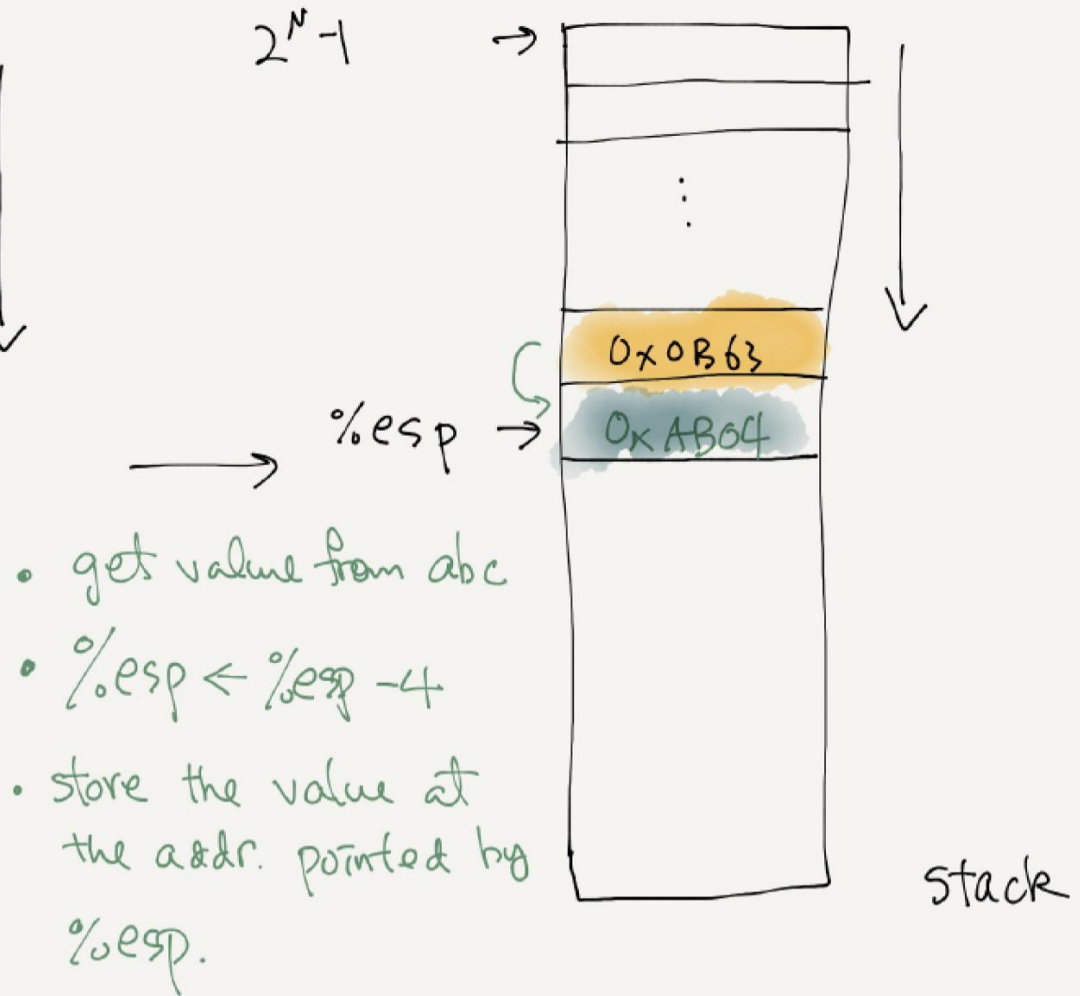

    if (cur != next)
        prev = switch_threads (cur, next);
    thread_schedule_tail (prev);
}
```

Stack




Process context

push abc



Process context

pop abc



User stack vs. kernel stack

Executing in User Mode

- user stack

executing in the kernel


- use kernel stack

- System call

- entering the kernel

① switch from user stack to kernel stack

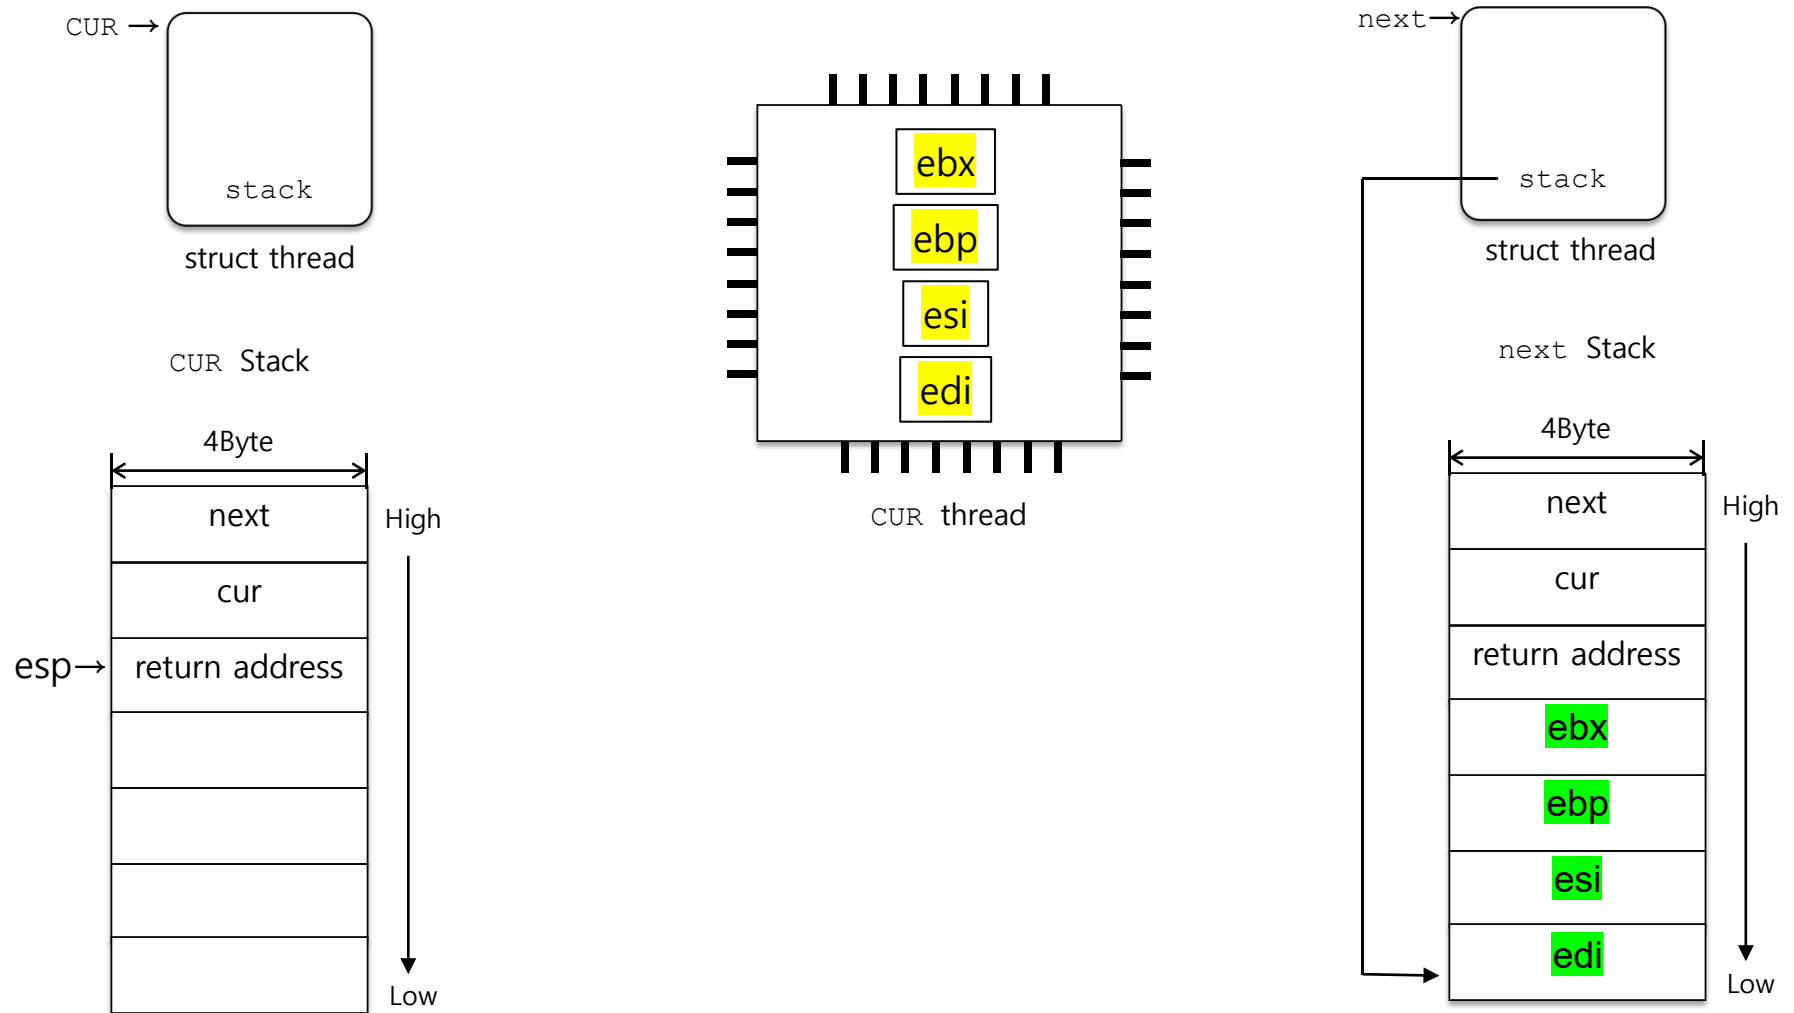
② raise privilege level




```
switch_threads(struct thread *cur,  
               struct thread *next)
```

- ▣ Save the registers on the kernel stack of `cur`.
- ▣ Save the location of the current stack top at the current `struct thread`'s `stack` member.
- ▣ Restore the new thread's stack top (kernel stack) into CPU's stack pointer (`esp`).
- ▣ Restore registers from the stack (kernel stack).

Call switch_threads




```
switch_threads(struct thread *cur,  
               struct thread *next)
```

▣ Two tasks

- ◆ Switch the CPU context from CUR thread to NEXT thread.
- ◆ Return `cur` (as `prev` below).

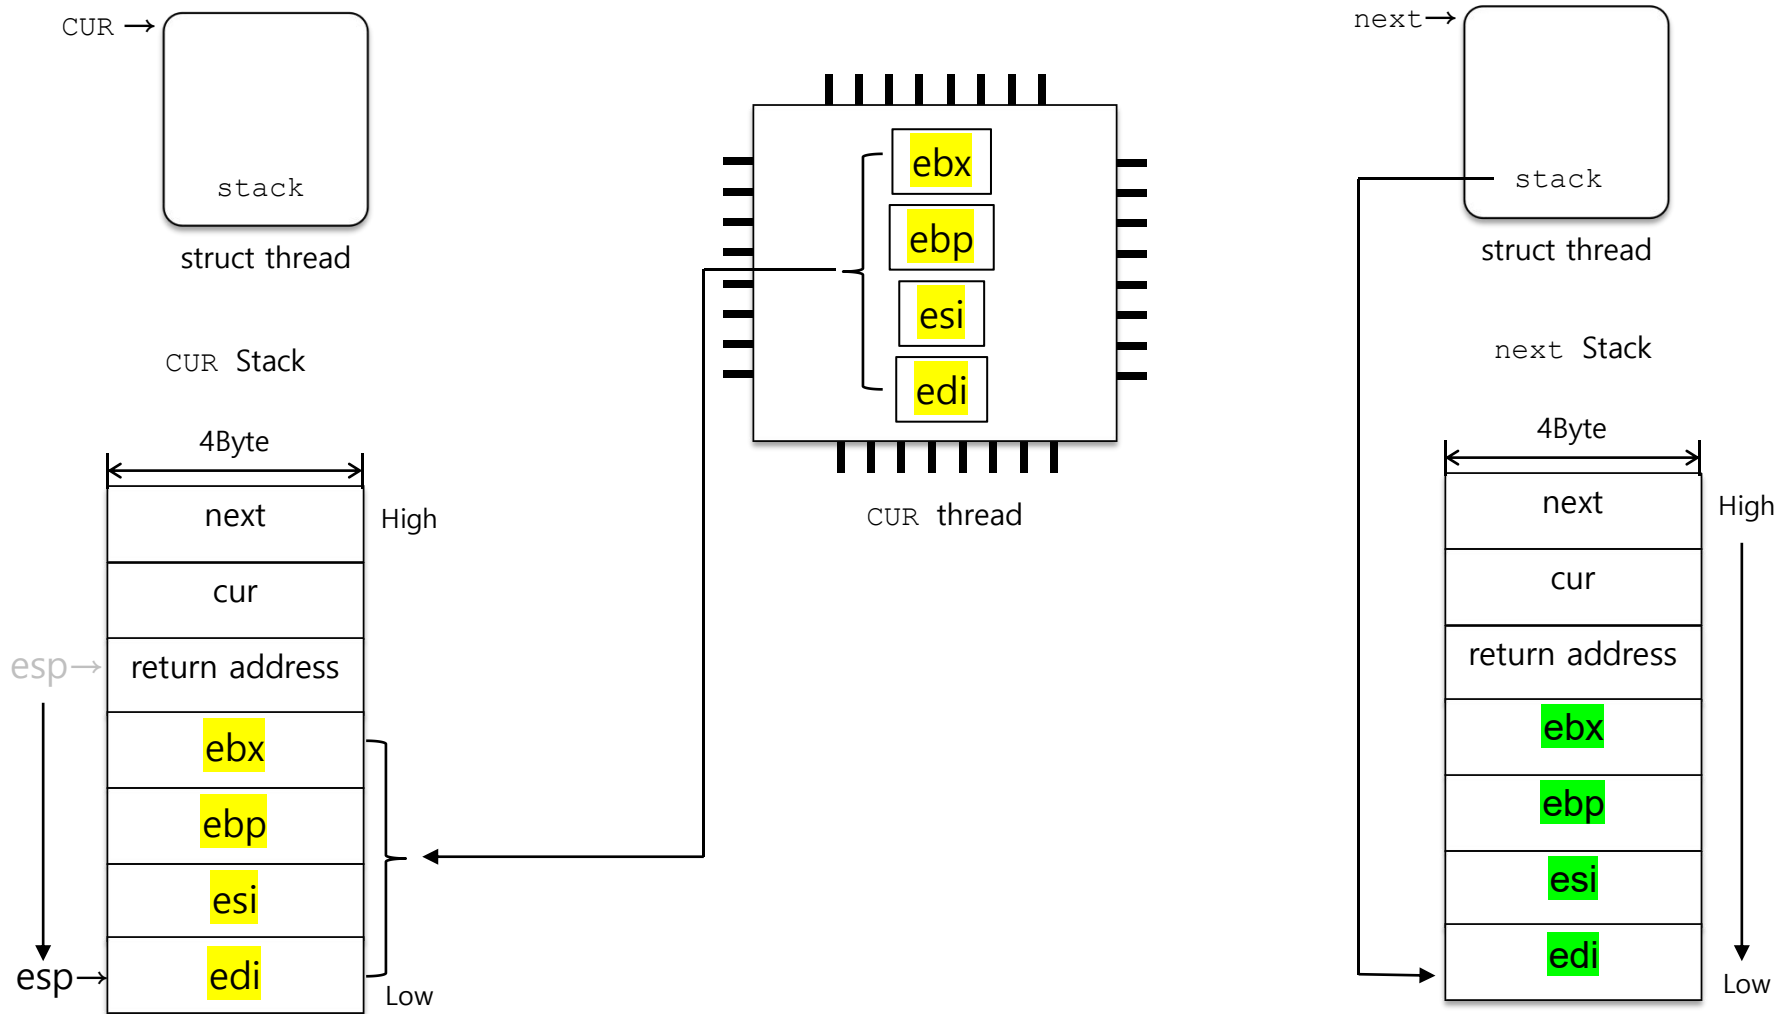
thread/thread.c

```
static void  
schedule (void)  
{  
    ...  
  
    prev = switch_threads(cur, next);  
  
    ...  
}
```



Just before jump into
`switch_thread`

Store current context



switch_threads: I

- Save current thread's registers at its kernel stack.

thread/switch.S

```
switch_threads:
```

```
    pushl %ebx  
    pushl %ebp  
    pushl %esi  
    pushl %edi
```

```
.globl thread_stack_ofs
```

```
    mov thread_stack_ofs, %edx
```

```
    movl SWITCH_CUR(%esp), %eax
```

```
    movl %esp, (%eax,%edx,1)
```

```
    movl SWITCH_NEXT(%esp), %ecx
```

```
    movl (%ecx,%edx,1), %esp
```

```
    popl %edi
```


```
    popl %esi
```

```
    popl %ebp
```

```
    popl %ebx
```

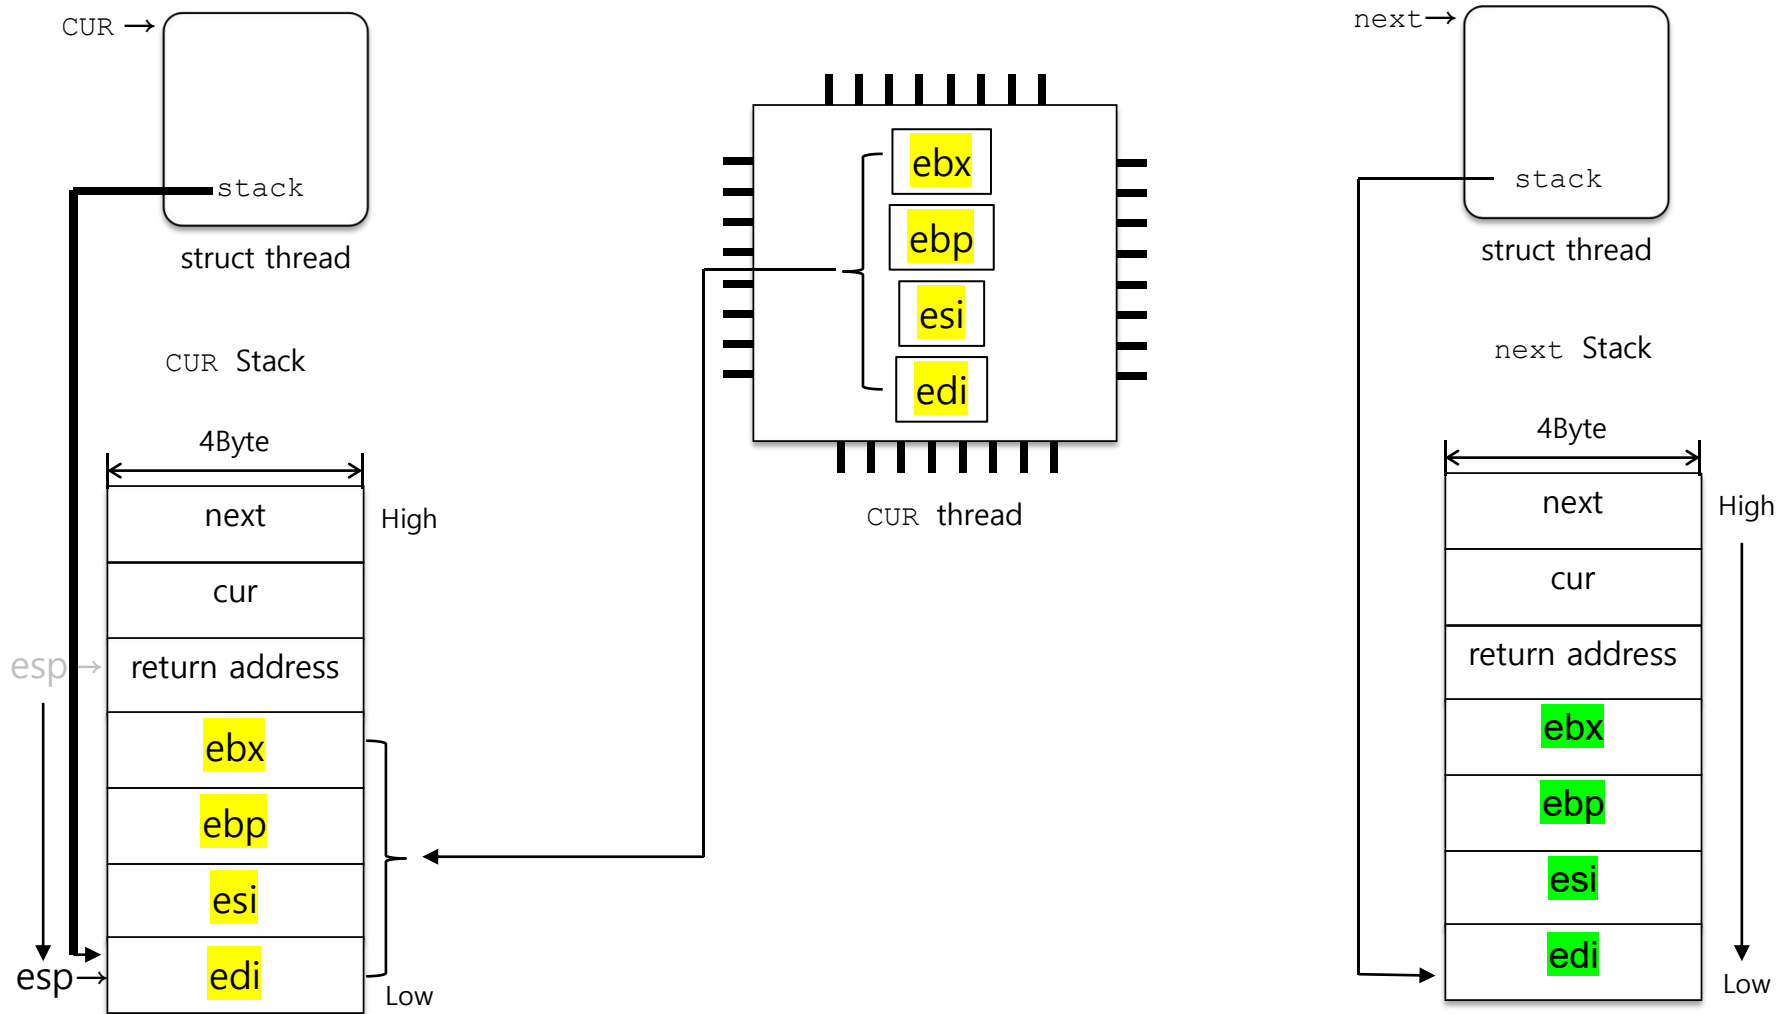
```
    ret
```

CUR Stack



After storing the context of `cur` at its kernel stack

Store current context



switch_threads

1. Get the offset of stack pointer in thread structure.

thread/thread.c

```
uint32_t thread_stack_ofs = offsetof (struct thread, stack)
```

lib/stddef.h

```
#define offsetof(TYPE, MEMBER) ((size_t) &((TYPE*)0) ->MEMBER)
```

2. Load the offset to edx.


```
.globl thread_stack_ofs
mov thread_stack_ofs, %edx

movl SWITCH_CUR(%esp), %eax
movl %esp, (%eax,%edx,1)

movl SWITCH_NEXT(%esp), %ecx
movl (%ecx,%edx,1), %esp

popl %edi
popl %esi
popl %ebp
popl %ebx
ret
```

struct thread



```
switch_threads(struct thread *cur,  
               struct thread *next)
```

3. Save the location of the current struct thread to `eax`. (return value)

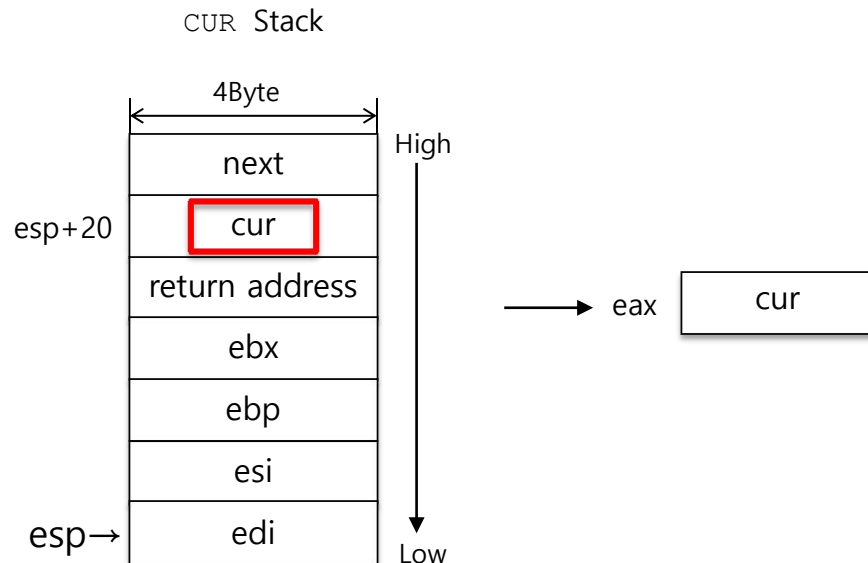
thread/switch.S

```
switch_threads:  
    pushl %ebx  
    pushl %ebp  
    pushl %esi  
    pushl %edi  
  
.globl thread_stack_ofs  
    mov thread_stack_ofs, %edx  
  
    movl SWITCH_CUR(%esp), %eax  
    movl %esp, (%eax,%edx,1)  
  
    movl SWITCH_NEXT(%esp), %ecx  
    movl (%ecx,%edx,1), %esp  
  
    popl %edi  
    popl %esi  
    popl %ebp  
    popl %ebx  
    ret
```

thread/switch.h

```
#define SWITCH_CUR 20
```

- $\text{SWITCH_CUR}(\%esp) = \%esp + 20$



switch_threads

- Save current thread's stack top address at the struct thread (stack field).

$(\%eax, \%edx, 1) = \%eax + \%edx * 1$

$\%eax$: location of the struct thread of CUR

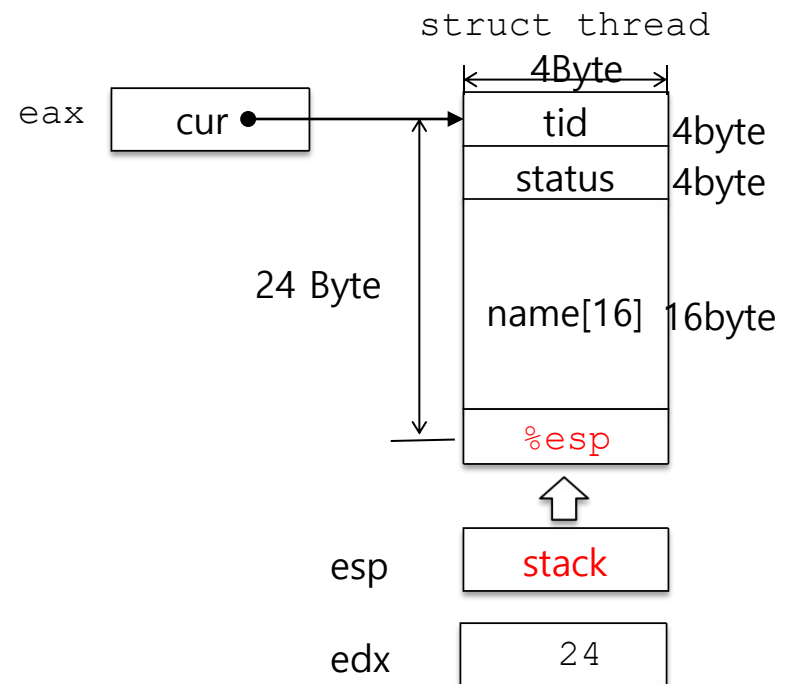
thread/switch.S

```
.globl thread_stack_ofs
mov thread_stack_ofs, %edx

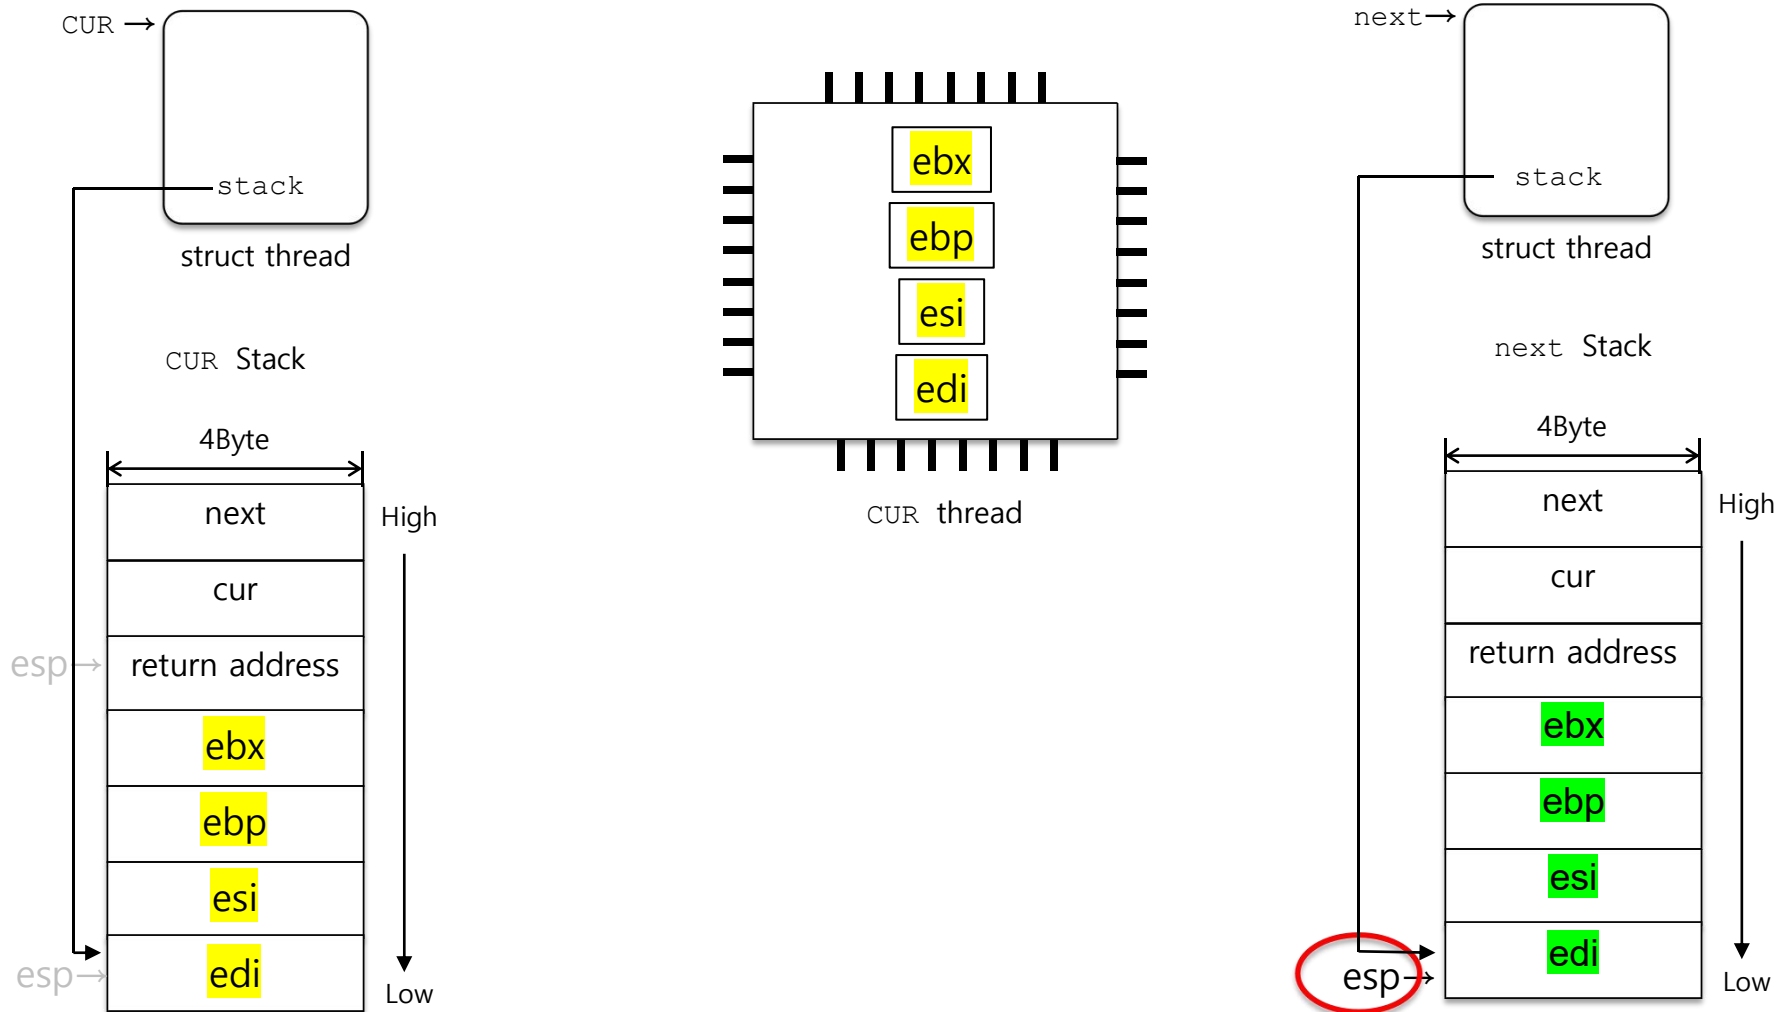
movl SWITCH_CUR(%esp), %eax
movl %esp, (%eax, %edx, 1)

movl SWITCH_NEXT(%esp), %ecx
movl (%ecx, %edx, 1), %esp

popl %edi
popl %esi
popl %ebp
popl %ebx
ret
```



Switch kernel stack



switch_threads

- Load address of the struct thread of NEXT to ecx.

thread/switch.S

```
switch_threads:
    pushl %ebx
    pushl %ebp
    pushl %esi
    pushl %edi

    .globl thread_stack_ofs
    mov thread_stack_ofs, %edx

    movl SWITCH_CUR(%esp), %eax
    movl %esp, (%eax,%edx,1)

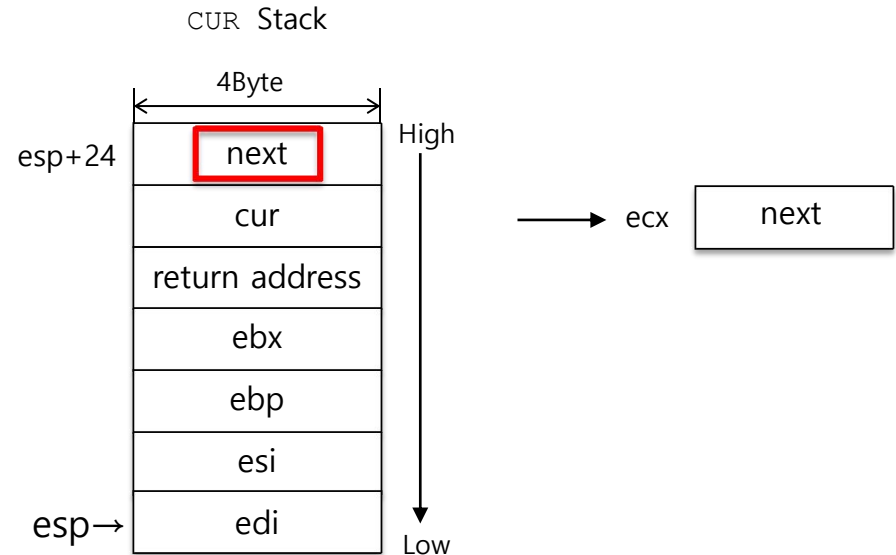
    movl SWITCH_NEXT(%esp), %ecx
    movl (%ecx,%edx,1), %esp

    popl %edi
    popl %esi
    popl %ebp
    popl %ebx
    ret
```

thread/switch.h

```
#define SWITCH_NEXT 24
```

$\text{SWITCH_NEXT}(\%esp) = \%esp + 24$



switch_threads

- Stack switch: Load address of the kernel stack of NEXT to esp.

thread/switch.S

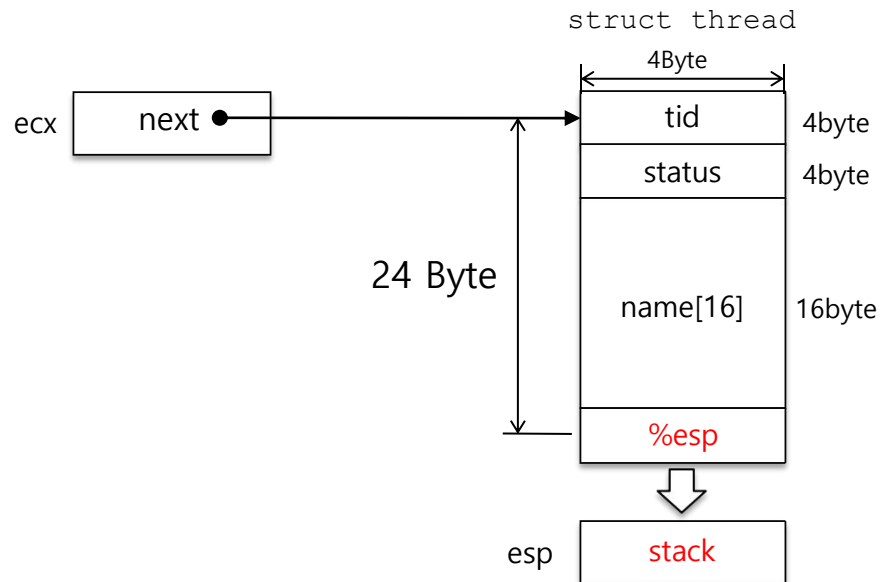
$$(\%ecx, \%edx, 1) = \%ecx + \%edx * 1$$

```
.globl thread_stack_ofs
mov thread_stack_ofs, %edx

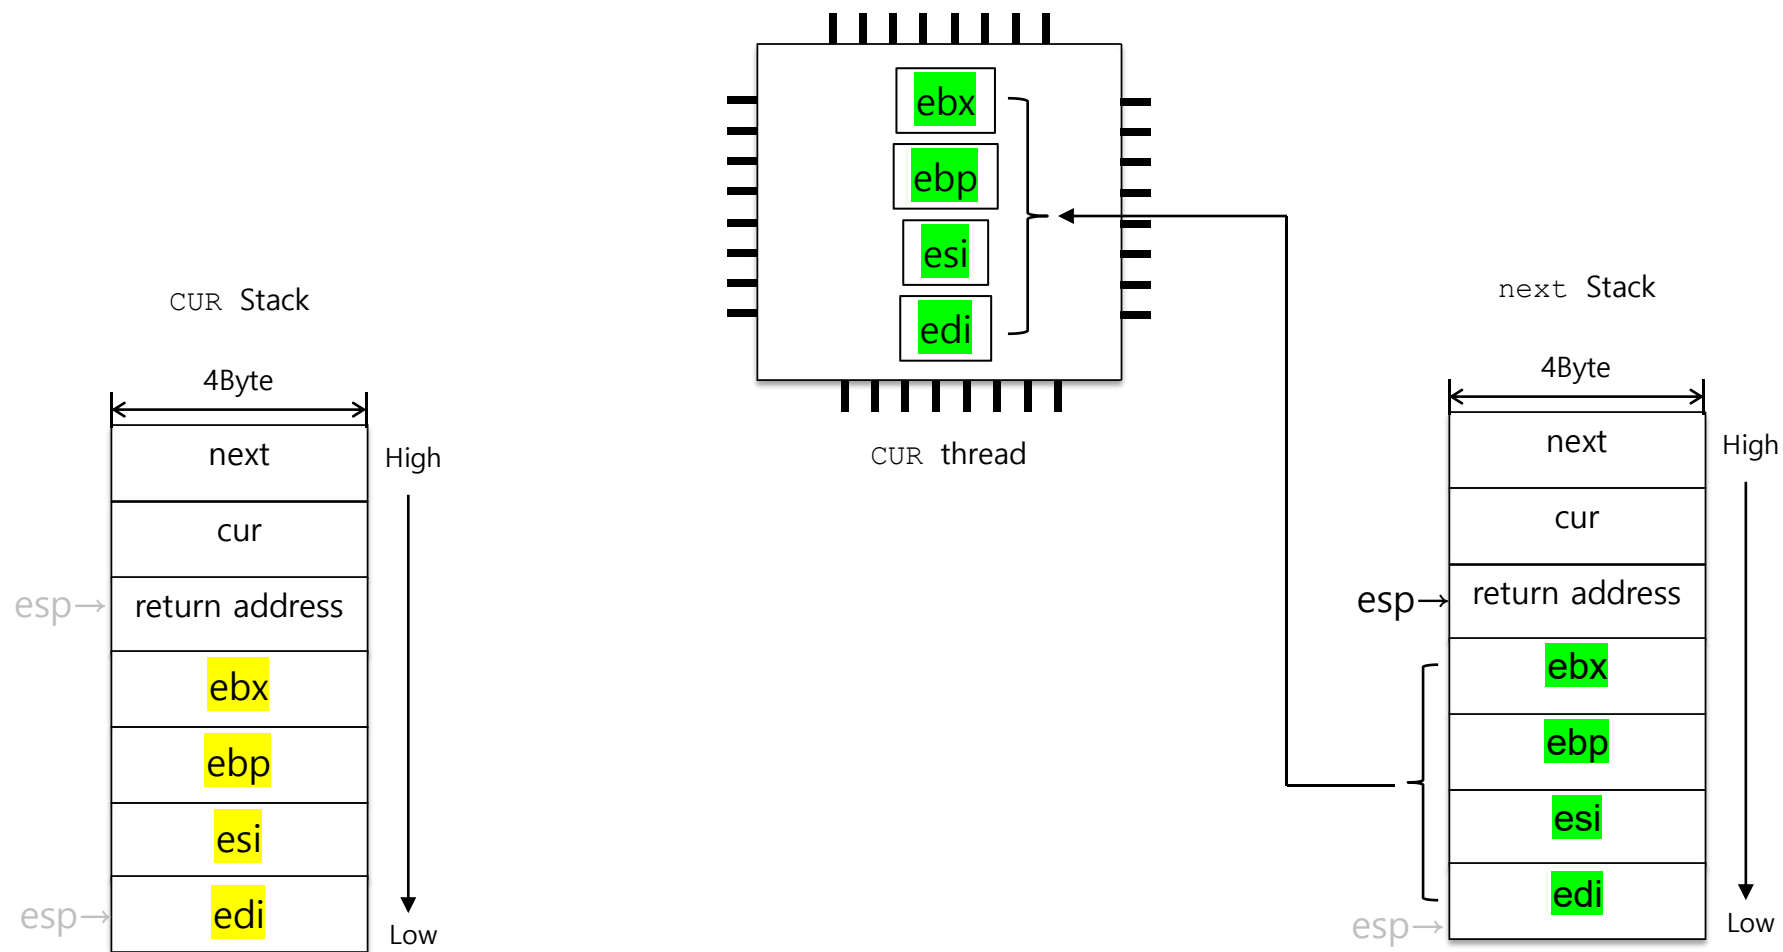
movl SWITCH_CUR(%esp), %eax
movl %esp, (%eax, %edx, 1)

movl SWITCH_NEXT(%esp), %ecx
movl (%ecx, %edx, 1), %esp

popl %edi
popl %esi
popl %ebp
popl %ebx
ret
```



Restore the new context



switch_threads

- ▣ Restore next thread's registers.
- ▣ Return and run instruction of return address.

thread/switch.S


```
switch_threads:
    pushl %ebx
    pushl %ebp
    pushl %esi
    pushl %edi

    .globl thread_stack_ofs
    mov thread_stack_ofs, %edx

    movl SWITCH_CUR(%esp), %eax (return value)
    movl %esp, (%eax,%edx,1)

    movl SWITCH_NEXT(%esp), %ecx
    movl (%ecx,%edx,1), %esp

    popl %edi
    popl %esi
    popl %ebp
    popl %ebx
    ret
```



`thread_schedule_tail()`

- ▣ Mark the new thread as running.
- ▣ If the previous thread was in the dying state, then it also frees the page.

Change the state of new current

```
void thread_schedule_tail (struct thread *prev) {
    struct thread *cur = running_thread ();

    ASSERT (intr_get_level () == INTR_OFF);

    cur->status = THREAD_RUNNING;

    thread_ticks = 0;

#ifdef USERPROG
    process_activate ();
#endif

    if (prev != NULL && prev->status == THREAD_DYING && prev != initial_
        thread)
    {
        ASSERT (prev != cur);
        pallocc_free_page (prev);
    }
}
```


Summary

- ▣ `schedule()`
 - ◆ Called in `exit`, `yield` and `block`.
 - ◆ Get the new process to the CPU.
- ▣ Context switch
 - ◆ Save the context of the currently running thread to the stack.
 - ◆ Save the current stack top at the currently running `struct thread`.
 - ◆ Restore the stack top of the next thread to `esp` register.
 - ◆ Restore the context from the stack of the next thread to run.
- ▣ Change the state of the next process to running and frees the memory from the dying process.