

Operating Systems Lab

Part 1: Threads



Youjip Won

4.4BSD like scheduler

Outline of Advanced Scheduler

- ▣ Main Goal
 - ◆ Implement 4.4 BSD scheduler MLFQ like scheduler without the queues.
 - Give priority to the processes with interactive nature.
 - Priority based scheduler
 - ◆ Use “equation”.
- ▣ Files to modify
 - ◆ threads/thread.*
 - ◆ devices/timer.c

- ▣ Nice value
 - ◆ Represents the 'niceness' of a thread.
 - ◆ If a thread is nicer, it is willing to give up some of its CPU time.
- ▣ Value (from -20 to 20)
 - ◆ Nice (0) : not influence on priority. (initial value)
 - ◆ Nice (positive) : decrease priority.
 - ◆ Nice (negative) : increase priority.
- ▣ Function
 - ◆ `thread_get_nice()`
 - ◆ `thread_set_nice(int new_nice)`

□ Priority

- ◆ From 0 (PRI_MIN) to 63 (PRI_MAX)
- ◆ The larger the number, the higher the priority.
- ◆ It initialized when thread is created (default: 31)

```
priority = PRI_MAX - (recent_cpu / 4) - (nice * 2)
```

▣ Philosophy

- ◆ If the thread is nicer, lower the priority.
- ◆ If the thread have been using lots of CPU recently, lower the priority.
- ◆ For all threads, priority is recalculated once in every fourth clock tick.
- ◆ The result is truncated to its nearest integer.

CPU usage

▣ Update the cpu usage

`recent_cpu`

- ◆ Increase the `recent_cpu` of the currently running process by 1 in every timer interrupt.

- ◆ Decay `recent_cpu` by decay factor in every second.

$$\text{recent_cpu} = \text{decay} * \text{recent_cpu}$$

- ◆ Adjust `recent_cpu` by nice in every second.

$$\text{recent_cpu} = \text{recent_cpu} + \text{nice}$$

- ◆ Putting them together,

$$\text{recent_cpu} = \text{decay} * \text{recent_cpu} + \text{nice}$$

Decay factor

- ▣ In SVR3

$$\text{decay} = \frac{1}{2}$$

- ▣ In BSD4.4

- ◆ In heavy load, decay is nearly 1.
- ◆ In light load, decay is 0.

$$\text{decay} = (2 * \text{load_average}) / (2 * \text{load_average} + 1)$$

load_average

▣ load_average

$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_threads}$$

- ◆ At booting, `load_avg` is initially set to 0.
- ◆ `ready_threads`: the number of threads in `ready_list` and threads in executing at the time of update. (except idle thread)

In summary,

In every fourth tick, recompute the priority of all threads

$$\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$$

In every clock tick, increase the running thread's `recent_cpu` by one.

In every second, update every thread's `recent_cpu`

$$\text{recent_cpu} = \text{decay} * \text{recent_cpu} + \text{nice}, \text{ where}$$

$$\text{decay} = (2 * \text{load_average}) / (2 * \text{load_average} + 1)$$

and

$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_threads}$$

$$\text{Priority} = \text{PRI_MAX} - (\text{recent_cpu}/4) - (\text{nice} * 2)$$

$$\text{recent_cpu} = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1) * \text{recent_cpu} + \text{nice}$$

$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_threads}$$

nice = 0,
load_avg = 0


Sec	Tick	P1		P2		P3	
		Priority	recent_cpu	Priority	recent_cpu	Priority	recent_cpu
0	0	63	0	63	0	63	0
	1	63	1	63	0	63	0
	2	63	2	63	0	63	0
	3	63	3	63	0	63	0
	4	62	4	63	0	63	0
	5	62	4	63	1	63	0
	6	62	4	63	2	63	0
	7	62	4	63	3	63	0
1	8	63	0	62	4	63	0
	9	63	0	62	4	63	1
	10	63	0	62	4	63	2
	11	63	0	62	4	63	3
	12	63	0	63	0	62	4
	13	63	1	63	0	62	4
	14	63	2	63	0	62	4
	15	63	3	63	0	62	4

Priority
=PRI_MAX(63)-(4/4)-(0*2)
=62

 Running Priority recalculate Recent CPU recalculate


Implement fixed point arithmetic

- ▣ Inside kernel, you can do only integer arithmetic.
 - ◆ Kernel does not save floating point register when switching the context.
- ▣ We need to implement fixed point arithmetic using integer arithmetic.
 - ◆ `priority`, `nice`, `ready_threads` value is integer, and `recent_cpu`, `load_avg` value is real number.
 - ◆ Implement the fixed-point arithmetic using 17.14 fixed-point number representation.
 - Decimal point is 14 right-most bits.
 - Integer is 17 next bits to the left.
 - Last of left 1bit is sign bit.



Operations to implement

- n : integer x, y : fixed-point numbers f : 1 in 17.14 format



Convert n to fixed point:	$n * f$
Convert x to integer (rounding toward zero):	x / f
Convert x to integer (rounding to nearest):	$(x + f / 2) / f$ if $x \geq 0$, $(x - f / 2) / f$ if $x < 0$.
Add x and y:	$x + y$
Subtract y from x:	$x - y$
Add x and n:	$x + n * f$
Subtract n from x:	$x - n * f$
Multiply x by y:	$((\text{int64_t}) x) * y / f$
Multiply x by n:	$x * n$
Divide x by y:	$((\text{int64_t}) x) * f / y$
Divide x by n:	x / n

Examples

- Convert n to fixed point: $n * f$: shift n by 14 bits to the left.
- Convert x to integer: shift x by 14 to the right.

Implementation

- ▣ Each thread maintains
 - ◆ `nice` and `recent_cpu`
 - ◆ Add `nice`, `recent_cpu` to thread structure.
- ▣ Functions to be added
 - ◆ The function that calculate priority using `recent_cpu` and `nice`.
 - ◆ The function that calculate `recent_cpu`.
 - ◆ The function that calculate `load_avg`.
 - ◆ The function that increase `recent_cpu` by 1.
 - ◆ The function that recalculate priority and `recent_cpu` of all threads.
- ▣ Multiple ready queues vs. single ready queue
 - ◆ You can use single in implementing BSD scheduler.
 - ◆ **If you use multiple ready queues, you will get 10 extra mark for 100 mark.**

Functions to modify

- ◆ `static void init_thread(struct thread *t, const char *name, int priority)`
 - Initialize `nice`, `recent_cpu`
- ◆ `void thread_set_priority(int new_priority)`
 - Disable the priority setting when using the advanced scheduler.
- ◆ `static void timer_interrupt(struct intr frame *args UNUSED)`
 - Recalculate `load_avg`, `recent_cpu` of all threads, priority every 1 sec.
 - Recalculate priority of all threads every 4th tick.

Functions to modify

- ◆ `void lock_acquire(struct lock *lock)`
 - Forbid the priority donation when using the advanced scheduler.
- ◆ `void lock_release(struct lock *lock)`
 - Forbid the priority donation when using the advanced scheduler.

Functions to modify

- ◆ `void thread_set_nice(int nice UNUSED)`
 - Set `nice` value of the current thread.
- ◆ `int thread_get_nice(void)`
 - Return `nice` value of the current thread.
- ◆ `int thread_get_load_avg(void)`
 - Return `load_avg` multiplied by 100
 - `timer_ticks() % TIMER_FREQ == 0`
- ◆ `int thread_get_recent_cpu(void)`
 - Return `recent_cpu` multiplied by 100

Result

```
$ make check
```

```
pass tests/threads/mlfqs-block
pass tests/threads/alarm-single
pass tests/threads/alarm-multiple
pass tests/threads/alarm-simultaneous
pass tests/threads/alarm-priority
pass tests/threads/alarm-zero
pass tests/threads/alarm-negative
pass tests/threads/priority-change
pass tests/threads/priority-donate-one
pass tests/threads/priority-donate-multiple
pass tests/threads/priority-donate-multiple2
pass tests/threads/priority-donate-nest
pass tests/threads/priority-donate-sema
pass tests/threads/priority-donate-lower
pass tests/threads/priority-fifo
pass tests/threads/priority-preempt
pass tests/threads/priority-sema
pass tests/threads/priority-condvar
pass tests/threads/priority-donate-chain
pass tests/threads/mlfqs-load-1
pass tests/threads/mlfqs-load-60
pass tests/threads/mlfqs-load-avg
pass tests/threads/mlfqs-recent-1
pass tests/threads/mlfqs-fair-2
pass tests/threads/mlfqs-fair-20
pass tests/threads/mlfqs-nice-2
pass tests/threads/mlfqs-nice-10
pass tests/threads/mlfqs-block
All 27 tests passed.
```