

# Operating Systems Lab

## Part 2: User Programs

---



**Youjip Won**

# Overview

- ▣ Objective

Execute a user program in Pintos.

- ▣ Background

- ▣ Topics

- ◆ Parameter Passing
- ◆ System call infrastructure
- ◆ File manipulation

# Background

---

# To run a program

- ▣ Read the executable file from the disk.
  - ◆ Filesystem issue
- ▣ Allocate memory for the program to run.
  - ◆ Virtual memory allocation
- ▣ Pass the parameters to the program.
  - ◆ Set up user stack.
- ▣ Context switch to the user program
  - ◆ OS should wait for the program to exit.

# Pintos filesystem

- ❑ Create virtual disk: in `userprog/build`

```
pintos-mkdisk filesys.dsk --filesys-size=2
```

- ◆ `filesys.dsk`: partition name
- ◆ Filesystem size: 2MByte

- ❑ Format the disk

```
pintos -f -q
```

- ❑ Copy the file to the pintos filesystem

- ◆ `-p`: put, `-g`: get, `-a`: target filename

```
pintos -p ../../examples/echo -a echo -- -q
```

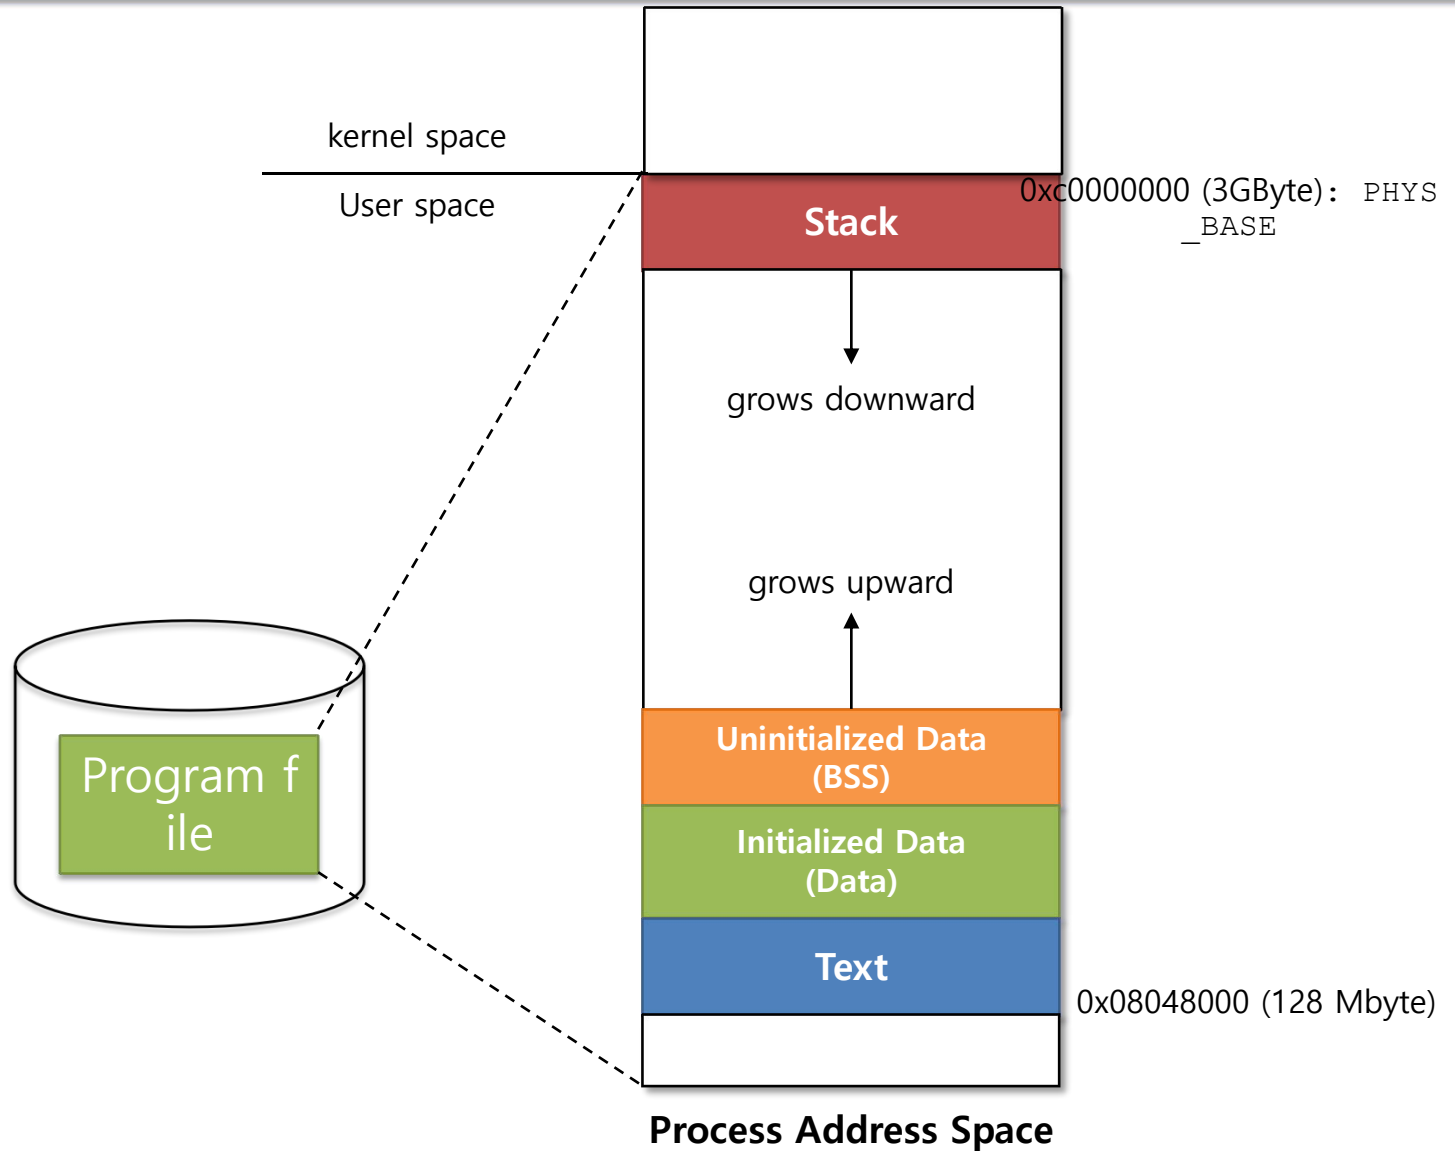
- ❑ Run the program

```
pintos -q run 'echo x'
```

- ❑ Merge the last three lines into one

```
pintos -p ../../examples/echo -a echo -- -f -q run 'echo x'
```

# Pintos VM layout



# Running a program in pintos

## Calling "process\_execute"

```
static void run_task(char ** argv)
{
    ...
    process_wait(process_execute(argv));
    ...
}
```

## Create thread and start running a program

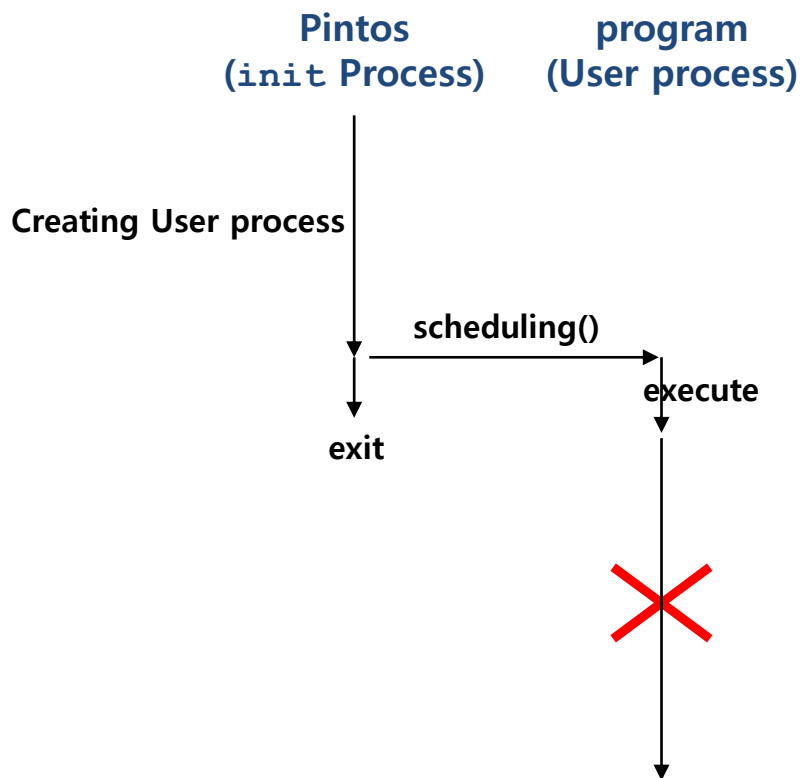
```
tid_t process_execute (const char
                        *file_name)
{
    ...
    tid = thread_create (... , start_process, .);
    ...
    return tid;
}
```

```
int process_wait (tid_t child_tid UNUSED)
{
    return -1;
}
```

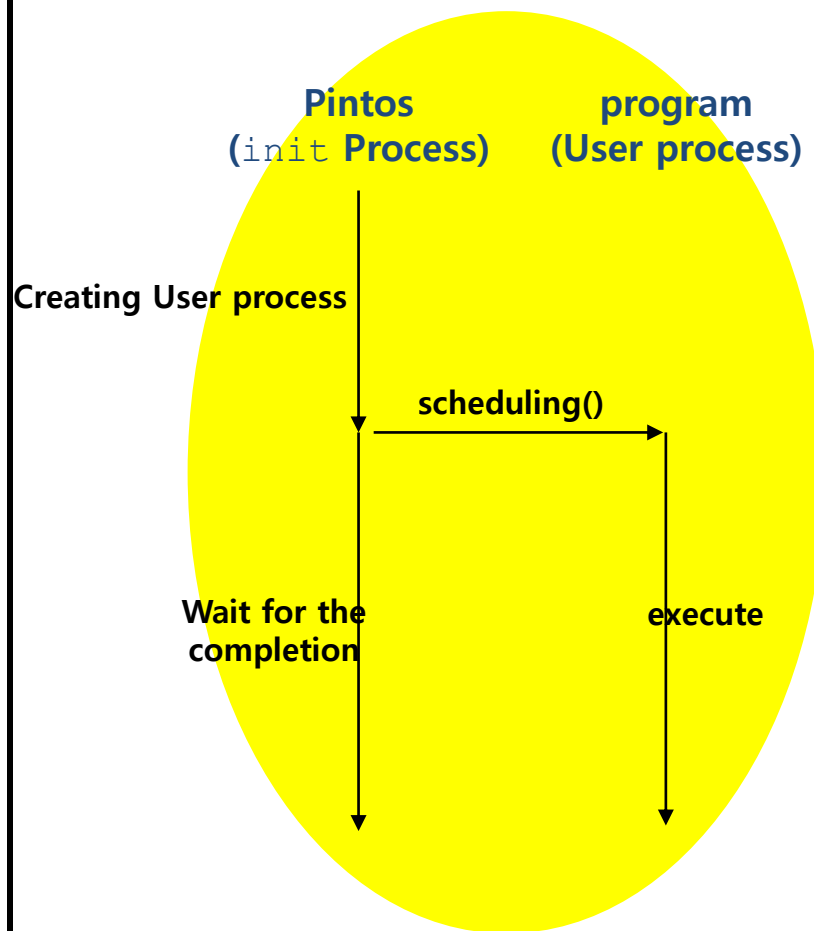
**The OS quits without waiting for the process to finish!!!**

# Executing a program

## Current Pintos



## Final Goal



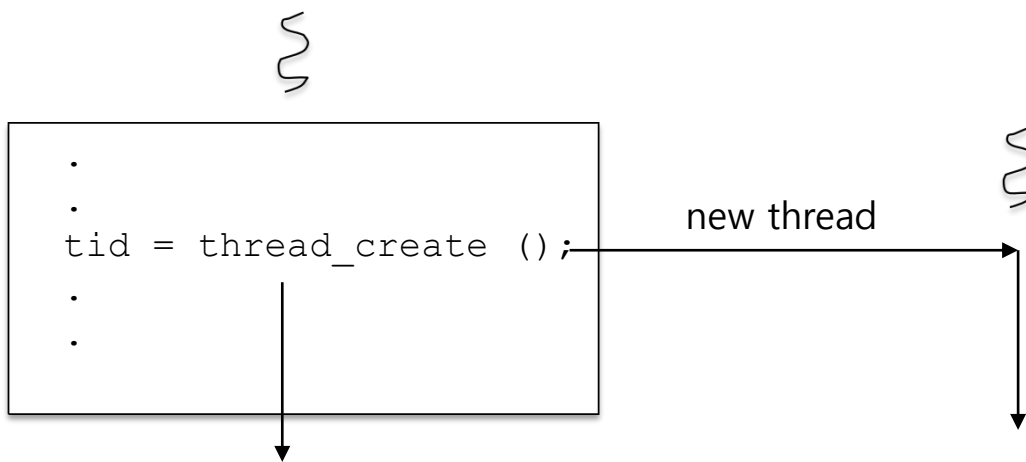


# Executing a program

- Execute “file\_name”.

pintos/src/userprog/process.c

```
tid_t process_execute (const char *file_name)
{
    char *fn_copy;
    tid_t tid;
    ...
    tid = thread_create (file_name, PRI_DEFAULT, start_process,
                        fn_copy);
    ...
    return tid;
}
```



# Creating a thread

- ▣ `thread_create()`
  - ◆ Create "struct thread" and initialize it.
  - ◆ Allocate the kernel stack.
  - ◆ Register the function to run: `start_process`.
  - ◆ Add it to ready list.

# Creating a thread

pintos/src/threads/thread.c – thread\_create()

```
tid_t thread_create (const char *name, int priority,
                    thread_func *function, void *aux)
{
    struct thread *t;
    struct kernel_thread_frame *kf;
    ...
    t = palloc_get_page (PAL_ZERO); /* allocating one page*/
    init_thread (t, name, priority); /* initialize thread structure*/
    tid = t->tid = allocate_tid (); /* allocate tid */
    /* Stack frame for kernel_thread(). */
    kf = alloc_frame (t, sizeof *kf); /* allocate stack */
    kf->eip = NULL;
    kf->function = function; /* function to run*/
    kf->aux = aux;           /* parameters for the function to run */
    ...
    /* Add to run queue. */
    thread_unblock (t);
    return tid;
}
```

# Starting a process

- ◆ `load()`: load the program of name `'file_name'`
- ◆ If it successfully loads the program, run it. Otherwise, `exit()`.
- ◆ `thread_exit()`: quit the thread.

pintos/src/userprog/process.c

```
static void start_process (void *file_name_)
{
    char *file_name = file_name_;
    ...
    /* if_.esp: address of the top of the user stack */
    success = load (file_name, &if_.eip, &if_.esp);
    if (!success)
        thread_exit ();
    /* start user program */
    asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g"
                  (&if_) : "memory");
}
```

# start\_process

```
static void start_process (void *file_name_)
{
    char *file_name = file_name_;
    struct intr_frame if_;
    bool success;
    ...
    success = load (file_name, &if_.eip, &if_.esp);
    if (!success)
        thread_exit ();
    /* Start the user process */
    asm volatile ("movl %0, %%esp; jmp    intr_exit" : : "g" (&if_) : "memory");
}
```

```
bool load (const char *file_name, void (*eip) (void), void **esp)
{
    ...
    struct file *file = NULL;
    ...
    file = filesys_open (file_name);
    ...
    /* Set up stack. */
    if (!setup_stack (esp))
        ...
    success = true;
    return success;
}
```

```
void thread_exit (void)
{
    ...
    process_exit ();
    intr_disable ();
    list_remove (&thread_current()->allelem);
    thread_current ()->status = THREAD_DYING;
    schedule ();
}
```

# Loading a program.

- ▣ Load a ELF file.
  - ◆ Create page table (2 level paging).
  - ◆ Open the file, read the ELF header.
  - ◆ Parse the file, load the 'data' to the data segment.
  - ◆ Create user stack and initialize it.

```

bool load (const char *file_name, void (**eip) (void), void **esp) {
    struct thread *t = thread_current ();
    struct Elf32_Ehdr ehdr;
    struct file *file = NULL;
    ...
    t->pagedir = pagedir_create (); /* create page directory */
    process_activate (); /* set cr3 register*/
    file = filesys_open (file_name); /* Open the file*/
    /* parse the ELF file and get the ELF header*/
    if (file_read (file, &ehdr, sizeof ehdr) != sizeof ehdr
        || memcmp (ehdr.e_ident, "\177ELF\1\1\1", 7)
        || ehdr.e_type != 2
        || ehdr.e_machine != 3
        || ehdr.e_version != 1
        || ehdr.e_phentsize != sizeof (struct Elf32_Phdr)
        || ehdr.e_phnum > 1024)
    /* load segment information */
    struct Elf32_Phdr phdr;
    if (file_ofs < 0 || file_ofs > file_length (file))
        file_seek (file, file_ofs);
    if (file_read (file, &phdr, sizeof phdr) != sizeof phdr)
    ...
    /* load the executable file */
    if (!load_segment (file, file_page, (void *) mem_page,
                      read_bytes, zero_bytes, writable))
    ...
    if (!setup_stack (esp)) /* initializing user stack*/
    *eip = (void (*) (void)) ehdr.e_entry; /*initialize entry point*/
}

```

# Passing the arguments and creating a thread

---



# Overview

- ▣ For `"echo x y z"`
  - ◆ Original:
    - Thread name: `"echo x y z"`
    - Find program with file name `"echo x y z"`
    - Arguments `"echo"`, `"x"`, `"y"`, and `"z"` are not passed
  - ◆ After modification
    - Thread name: `"echo"`
    - Find program with file name `"echo"`
    - Push the arguments to user stack.
- ▣ Files to modify
  - ◆ `pintos/src/userprog/process.*`

# Parse the arguments and push them to the user stack

## ▣ pintos/src/userprog/process.c

```
tid_t process_execute() (const char *file_name)
```

- ◆ Parse the string of `file_name`
- ◆ Forward first token as name of new process to `thread_create()` function

```
static void start_process() (void *file_name_)
```

- ◆ Parse `file_name`
- ◆ Save tokens on user stack of new process.

# Tokenizing

```
char *strtok_r (char *s, const char *delimiters,  
               char **save_ptr) /* string.h */
```

- ◆ Receive a string (s) and delimiters and parse them by delimiters

ex) Parsing a string by the first space

```
char s[] = "String to tokenize.";
char *token, *save_ptr;
for (token = strtok_r (s, " ", &save_ptr); token != NULL;
     token = strtok_r (NULL, " ", &save_ptr))
    printf ("%s\n", token);
```

Result

```
`String'  
'to'  
'tokenize.'
```

# Program Name

## ▣ Thread name

- ◆ Before: Entire command line is passed to `thread_create()`
- ◆ After modification: Forward only first token of command line to first argument of `thread_create()`
  - "echo x y z" → only use "echo" for name of process

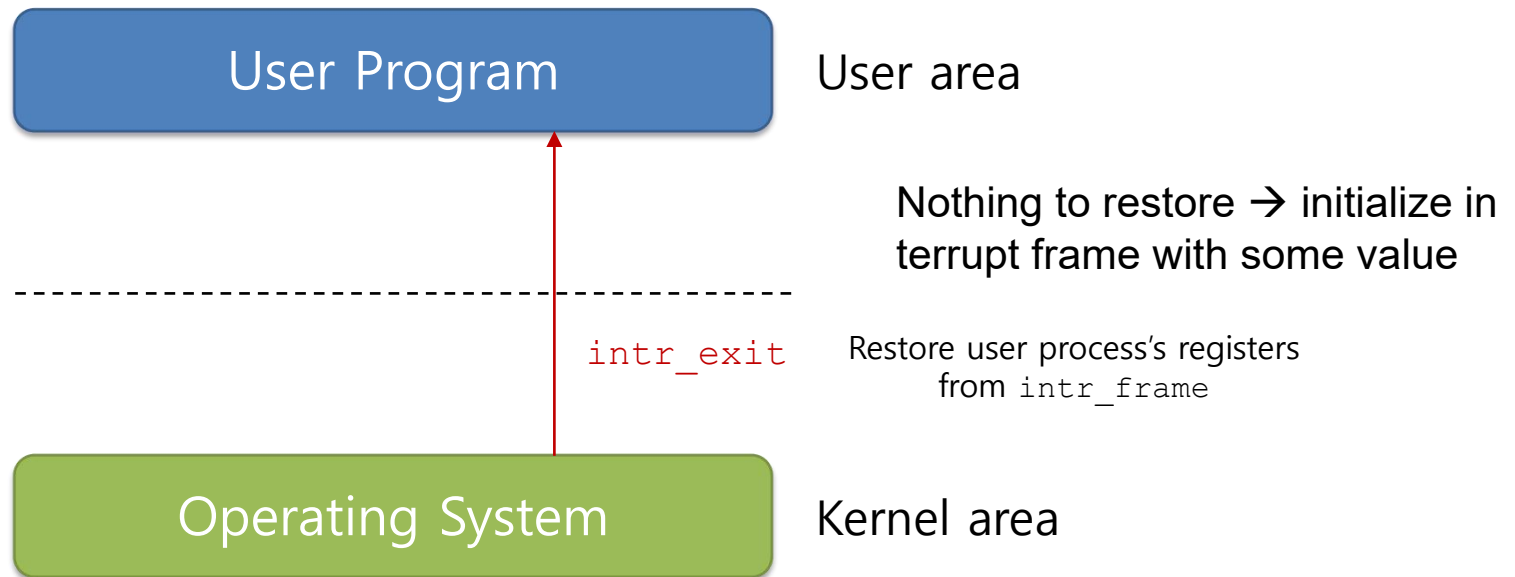
pintos/src/userprog/process.c

```
tid_t process_execute (const char *file_name)
{
    ...
    /* Parse command line and get program name */
    ...
    /* Create a new thread to execute FILE_NAME. */
    tid = thread_create (file_name, PRI_DEFAULT, start_process,
fn_copy);
    ...
}
```

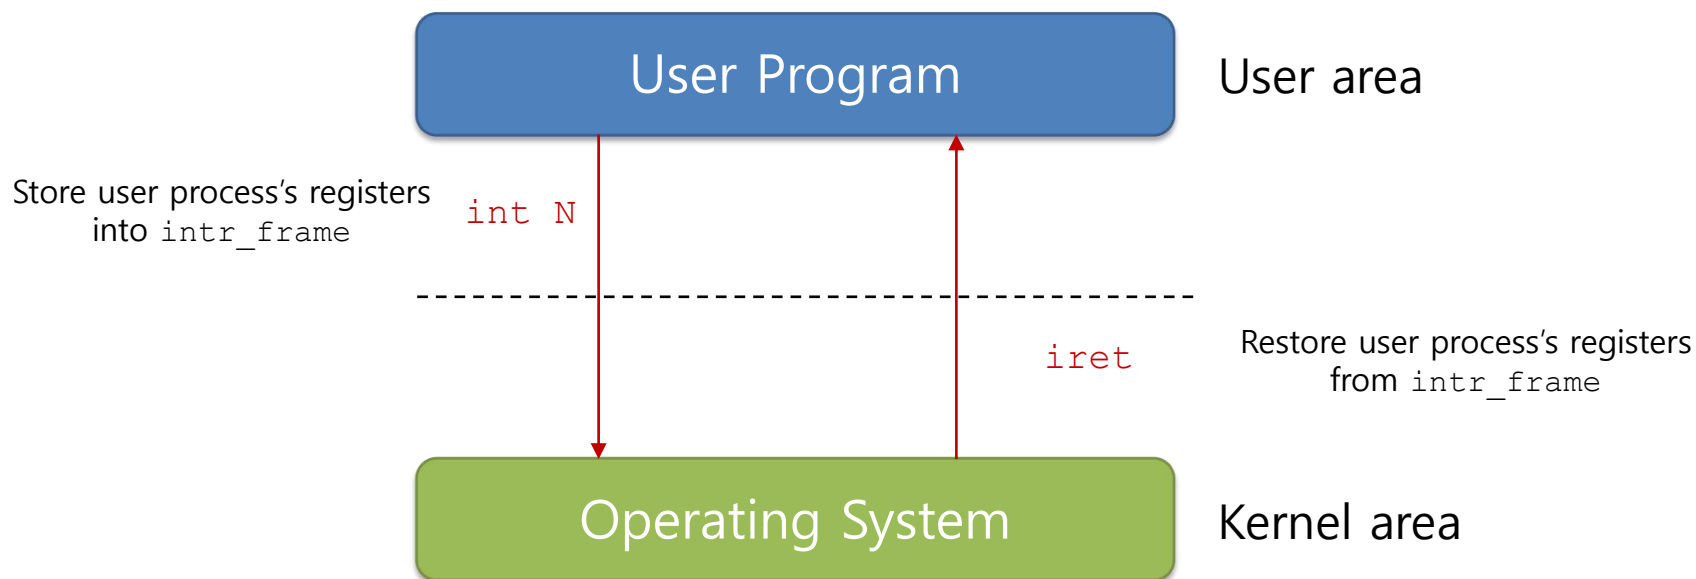
Name of thread

# start\_process

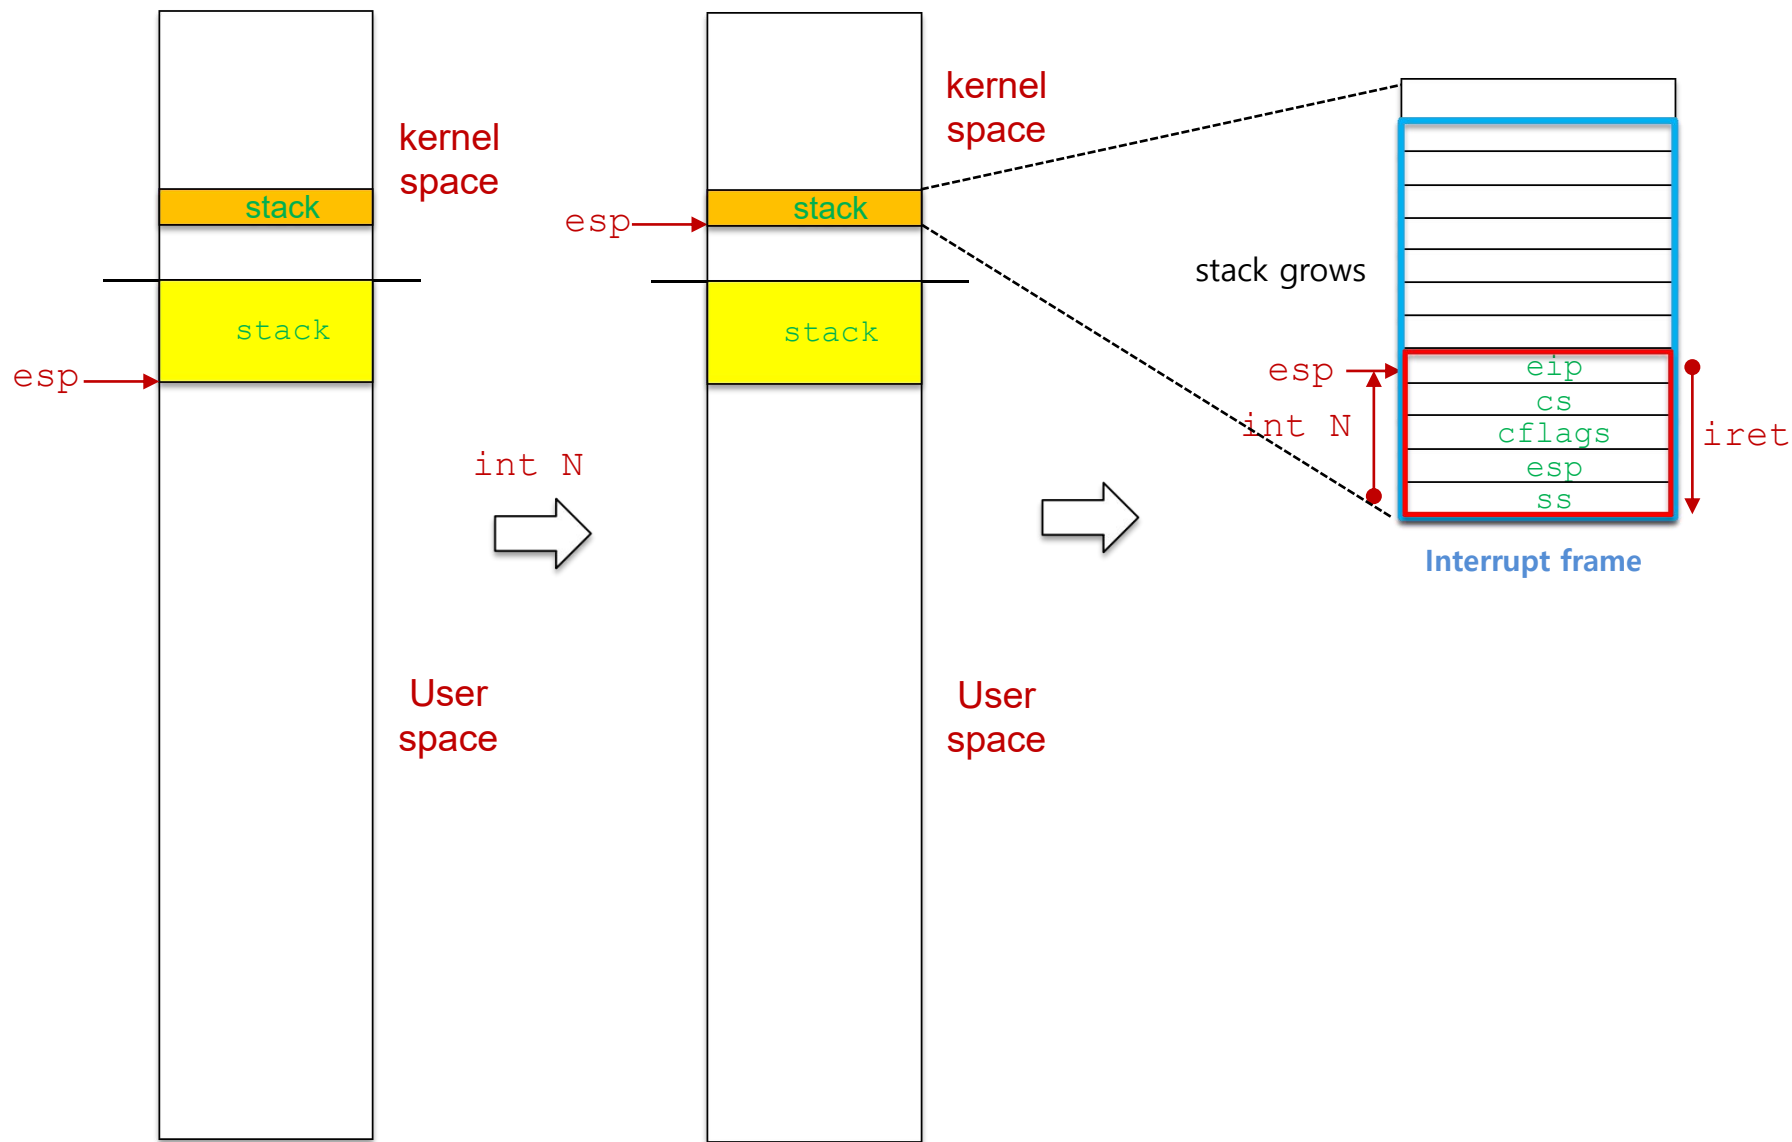
- Allocate interrupt frame.
- Load program and initialize interrupt frame and user stack.
- Setup arguments at the user stack.
- Jump to the user program through `interrupt_exit`.



# Getting into and out of kernel



# Getting into and out of kernel



# struct intr\_frame

```
struct intr_frame {
    /* Pushed by intr_entry in intr-stubs.S.
       These are the interrupted task's saved registers. */
    uint32_t edi;           /* Saved EDI. */
    uint32_t esi;           /* Saved ESI. */
    uint32_t ebp;           /* Saved EBP. */
    uint32_t esp_dummy;     /* Not used. */
    uint32_t ebx;           /* Saved EBX. */
    uint32_t edx;           /* Saved EDX. */
    uint32_t ecx;           /* Saved ECX. */
    uint32_t eax;           /* Saved EAX. */
    uint16_t gs, :16;       /* Saved GS segment register. */
    uint16_t fs, :16;       /* Saved FS segment register. */
    uint16_t es, :16;       /* Saved ES segment register. */
    uint16_t ds, :16;       /* Saved DS segment register. */
```

```
    /* Pushed by intrNN_stub in intr-stubs.S. */
    uint32_t vec_no;        /* Interrupt vector number. */

    /* Sometimes pushed by the CPU,
       otherwise for consistency pushed as 0 by intrNN_stub.
       The CPU puts it just under `eip', but we move it here. */
    uint32_t error_code;    /* Error code. */

    /* Pushed by intrNN_stub in intr-stubs.S.
       This frame pointer eases interpretation of backtraces. */
    void *frame_pointer;    /* Saved EBP (frame pointer). */
```

```
    /* Pushed by the CPU.
       These are the interrupted task's saved registers. */
    void (*eip) (void);     /* Next instruction to execute. */
    uint16_t cs, :16;        /* Code segment for eip. */
    uint32_t eflags;        /* Saved CPU flags. */
    void *esp;              /* Saved stack pointer. */
    uint16_t ss, :16;       /* Data segment for esp. */
```

```
};
```

- It is in the kernel stack.
- It stores user process' registers.

Stack grows.



# Getting into kernel.

```
int n
```

- ▣ when execute the kernel function, e.g. interrupt handler, system call, the OS save s the registers of currently executing process.
- ▣ Where: at the kernel stack of the executing process.
- ▣ execution
  1. Set the `esp` to point to kernel stack
  2. Pushes registers.

# Entering the kernel

```
struct intr_frame {
```

```
    /* Pushed by intr_entry in intr-stubs.S.
```

```
    These are the interrupted task's saved registers. */
```

```
    uint32_t edi;           /* Saved EDI. */
    uint32_t esi;           /* Saved ESI. */
    uint32_t ebp;           /* Saved EBP. */
    uint32_t esp_dummy;     /* Not used. */
    uint32_t ebx;           /* Saved EBX. */
    uint32_t edx;           /* Saved EDX. */
    uint32_t ecx;           /* Saved ECX. */
    uint32_t eax;           /* Saved EAX. */
    uint16_t gs, :16;       /* Saved GS segment register. */
    uint16_t fs, :16;       /* Saved FS segment register. */
    uint16_t es, :16;       /* Saved ES segment register. */
    uint16_t ds, :16;       /* Saved DS segment register. */
```

← esp

After interrupt handler, `intr_entry`

```
    /* Pushed by intrNN_stub in intr-stubs.S. */
```

```
    uint32_t vec_no;        /* Interrupt vector number. */
```

```
    /* Sometimes pushed by the CPU,
```

```
    otherwise for consistency pushed as 0 by intrNN_stub.
```

```
    The CPU puts it just under `eip', but we move it here. */
```

```
    uint32_t error_code;    /* Error code. */
```

```
    /* Pushed by intrNN_stub in intr-stubs.S.
```

```
    This frame pointer eases interpretation of backtraces. */
```

```
    void *frame_pointer;    /* Saved EBP (frame pointer). */
```

← esp

After interrupt handler of `intr N`

```
    /* Pushed by the CPU.
```

```
    These are the interrupted task's saved registers. */
```

```
    void (*eip) (void);     /* Next instruction to execute. */
    uint16_t cs, :16;        /* Code segment for eip. */
    uint32_t eflags;        /* Saved CPU flags. */
    void *esp;              /* Saved stack pointer. */
    uint16_t ss, :16;        /* Data segment for esp. */
```

← esp

After `int` instruction.

time ↑

# Loading

## ▣ Load the program

- ◆ Pass the program name to 'load()'.
- ◆ "Load()" find executable file, using name of file and load it onto memory.

pintos/src/userprog/process.c

```
static void start_process (void *file_name_)
{
    char *file_name = file_name_;
    struct intr_frame if_;
    bool success;
    ...
    /* Parse the command line (Use strtok_r()) */

    /* Initialize interrupt frame and load executable. */
    memset (&if_, 0, sizeof if_);
    ...
    success = load (file_name, &if_.eip, &if_.esp);
    ...
}
```

program name    Function entry point    Stack top (user stack)

```

static void
start_process (void *file_name_)
{
    char *file_name = file_name_;
    struct intr_frame if_;
    bool success;

    /* Initialize interrupt frame and load executable. */
    memset (&if_, 0, sizeof if_);
    if_.gs = if_.fs = if_.es = if_.ds = if_.ss = SEL_UDSEG;
    if_.cs = SEL_UCSEG;
    if_.eflags = FLAG_IF | FLAG_MBS;
    success = load (file_name, &if_.eip, &if_.esp);

    /* If load failed, quit. */
    palloc_free_page (file_name);
    if (!success)
        thread_exit ();
    /*missing parts!!! set up stack */
    /* Start the user process by simulating a return from an
       interrupt, implemented by intr_exit (in
       threads/intr-stubs.S). Because intr_exit takes all of its
       arguments on the stack in the form of a `struct intr_frame',
       we just point the stack pointer (%esp) to our stack frame
       and jump to it. */
    asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g" (&if_) : "memory");
    NOT_REACHED ();
}

```

# Getting out of the kernel

```
asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g" (&if_) : "memory");
```

```
movl %0, %%esp
```

Set the esp to the top of the interrupt frame.

```
jmp intr_exit
```

executes intr\_exit

```
.globl intr_exit
.func intr_exit
intr_exit:
    /* Restore caller's registers. */
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds

    /* Discard `struct intr_frame' vec_no, error_code,
       frame_pointer members. */
    addl $12, %esp

    /* Return to caller. */
    iret
.endfunc
```

# Getting out of the kernel

```
struct intr_frame {
    /* Pushed by intr_entry in intr-stubs.S.
       These are the interrupted task's saved registers. */
    uint32_t edi;           /* Saved EDI. */
    uint32_t esi;           /* Saved ESI. */
    uint32_t ebp;           /* Saved EBP. */
    uint32_t esp_dummy;     /* Not used. */
    uint32_t ebx;           /* Saved EBX. */
    uint32_t edx;           /* Saved EDX. */
    uint32_t ecx;           /* Saved ECX. */
    uint32_t eax;           /* Saved EAX. */
    uint16_t gs, :16;       /* Saved GS segment register. */
    uint16_t fs, :16;       /* Saved FS segment register. */
    uint16_t es, :16;       /* Saved ES segment register. */
    uint16_t ds, :16;       /* Saved DS segment register. */

```

← esp

```
/* Pushed by intrNN_stub in intr-stubs.S. */
uint32_t vec_no;           /* Interrupt vector number. */

/* Sometimes pushed by the CPU,
   otherwise for consistency pushed as 0 by intrNN_stub.
   The CPU puts it just under `eip', but we move it here. */
uint32_t error_code;       /* Error code. */

/* Pushed by intrNN_stub in intr-stubs.S.
   This frame pointer eases interpretation of backtraces. */
void *frame_pointer;       /* Saved EBP (frame pointer). */

```

← esp After intr\_exit

```
/* Pushed by the CPU.
   These are the interrupted task's saved registers. */
void (*eip) (void);        /* Next instruction to execute. */
uint16_t cs, :16;           /* Code segment for eip. */
uint32_t eflags;           /* Saved CPU flags. */
void *esp;                 /* Saved stack pointer. */
uint16_t ss, :16;          /* Data segment for esp. */

```

← esp After iret instruction

time



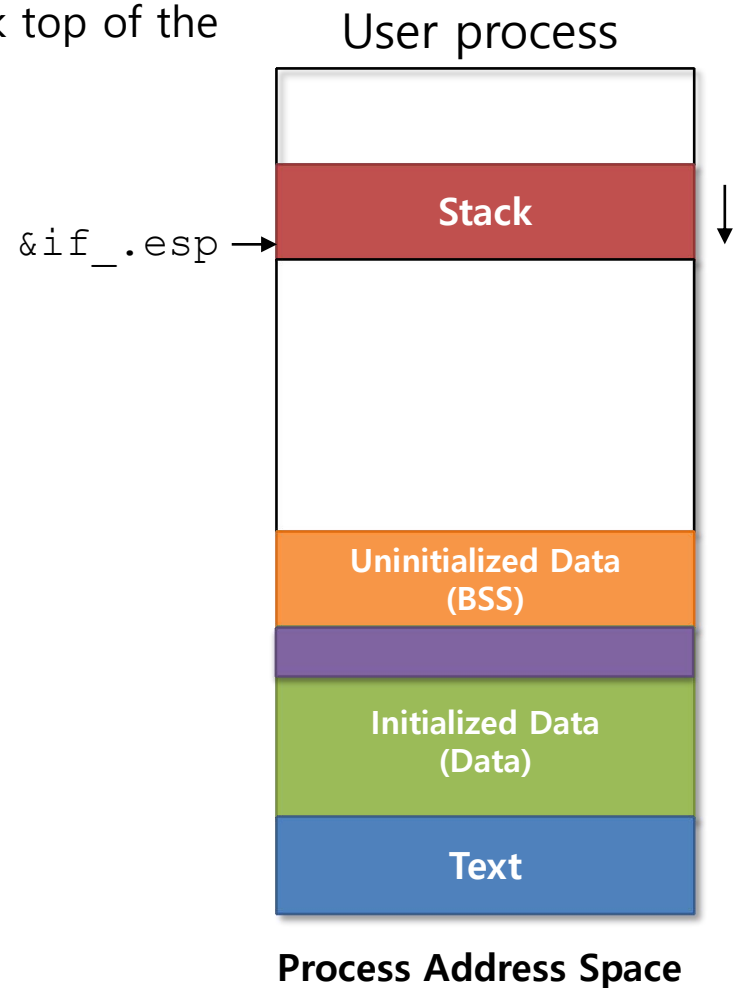
# Write a function that sets up a stack.

"esp" field of the interrupt frame contains the stack top of the user stack.

```
sample_function(int argc,  
                char* argv[],  
                void **stackpointer)
```

Current stack top: `&if_.esp`

Start from `&if_.esp - 4`



# 80x86 Calling Convention

```
%bin/ls -l foo bar
```

```
argc=4
```

```
argv[0] = "bin/ls", argv[1] = "-l", argv[2] = "foo", argv[3] = "bar"
```

1. Push arguments
  1. Push character strings from left to right.
  2. Place padding if necessary to align it by 4 Byte.
  3. Push start address of the character strings.
2. Push argc and argv
  1. Push argv
  2. Push argc
3. Push the address of the next instruction (return address).



# User stack layout in function call

```
%bin/ls -l foo bar
```

	Address	Name	Data	Type	
	0xbfffffffcc	argv[3][...]	'bar\0'	char[4]	Argument(string): 19 B
	0xbfffffff8	argv[2][...]	'foo\0'	char[4]	
	0xbffffff5	argv[1][...]	'-l\0'	char[3]	
	0xbffffffed	argv[0][...]	'/bin/ls\0'	char[8]	
	0xbffffffec	word-align	0	uint8_t	1Byte padding
grows ↓	0xbffffffe8	argv[4]	0	char *	Argument's address
	0xbffffffe4	argv[3]	0xbffffffc	char *	
	0xbffffffe0	argv[2]	0xbffffff8	char *	
	0xbffffffdc	argv[1]	0xbffffff5	char *	
	0xbffffffd8	argv[0]	0xbffffffed	char *	main(int argc , char **argv)
	0xbffffffd4	argv	0xbffffffd8	char **	
	0xbffffffd0	argc	4	int	fake address(0)
stack top →	0xbffffffcc	return address	0 (fake address)	void (*) ()	

Why is “return address” here is 0?

# Interim Check

- ▣ Print the program's stack by using `hex_dump()` (`stdio.h`)
  - ◆ Print memory dump in hexadecimal form
  - ◆ Check if arguments are correctly pushed on user stack.

pintos/src/userprog/process.c

```
static void start_process (void *file_name_)
{
    ...
    success = load (file_name, &if_.eip, &if_.esp);
    ...
    argument_stack(parse , count , &if_.esp);
    hex_dump(if_.esp , if_.esp , PHYS_BASE - if_.esp , true);
    e);

    asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g"
(&if_) : "memory");
    NOT_REACHED ();
}
```

추가

# Intermediate Check (Cont.)

## ▣ Result

```
$pintos -v -- run 'echo x'
```

```
Execution of 'echo x' complete.  
'echo'  
'x'  
Success : 1  
esp : bffffffe0  
bffffffe0 00 00 00 00 02 00 00 00-ec ff ff bf f9 ff ff bf |.....|  
bffffff0 fe ff ff bf 00 00 00 00-00 65 63 68 6f 00 78 00 |.....echo.x.|  
system call!
```

return address  
(fake)

argc

argv

echo

x

# System Calls and Handlers

---

# Overview

## ▣ Main goal

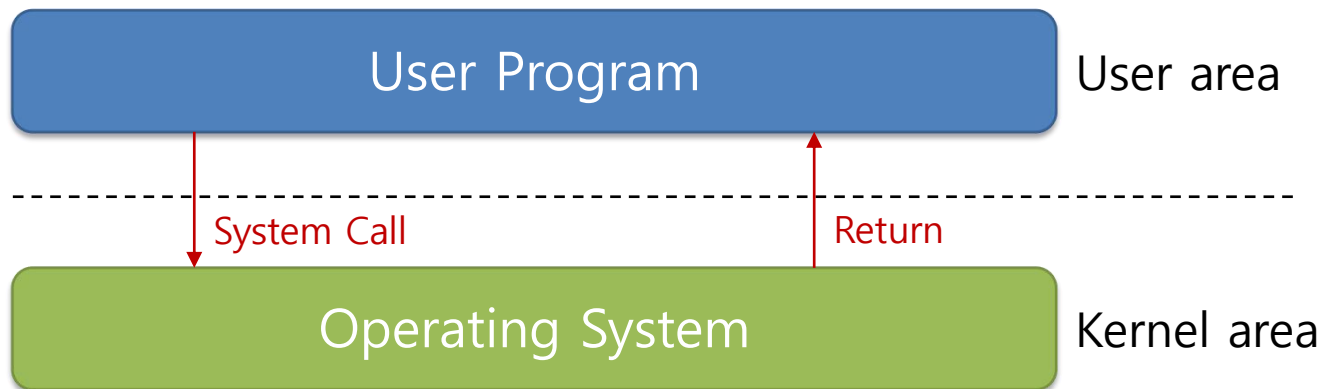
- ◆ Original: system call handler table is empty.
- ◆ After modification:
  - Fill system call handler of pintos out.
  - Add system calls to provide services to users
    - Process related: `halt`, `exit`, `exec`, `wait`
    - File related: `create`, `remove`, `open`, `filesize`, `read`, `write`, `seek`, `tell`, `close`

## ▣ Files to modify

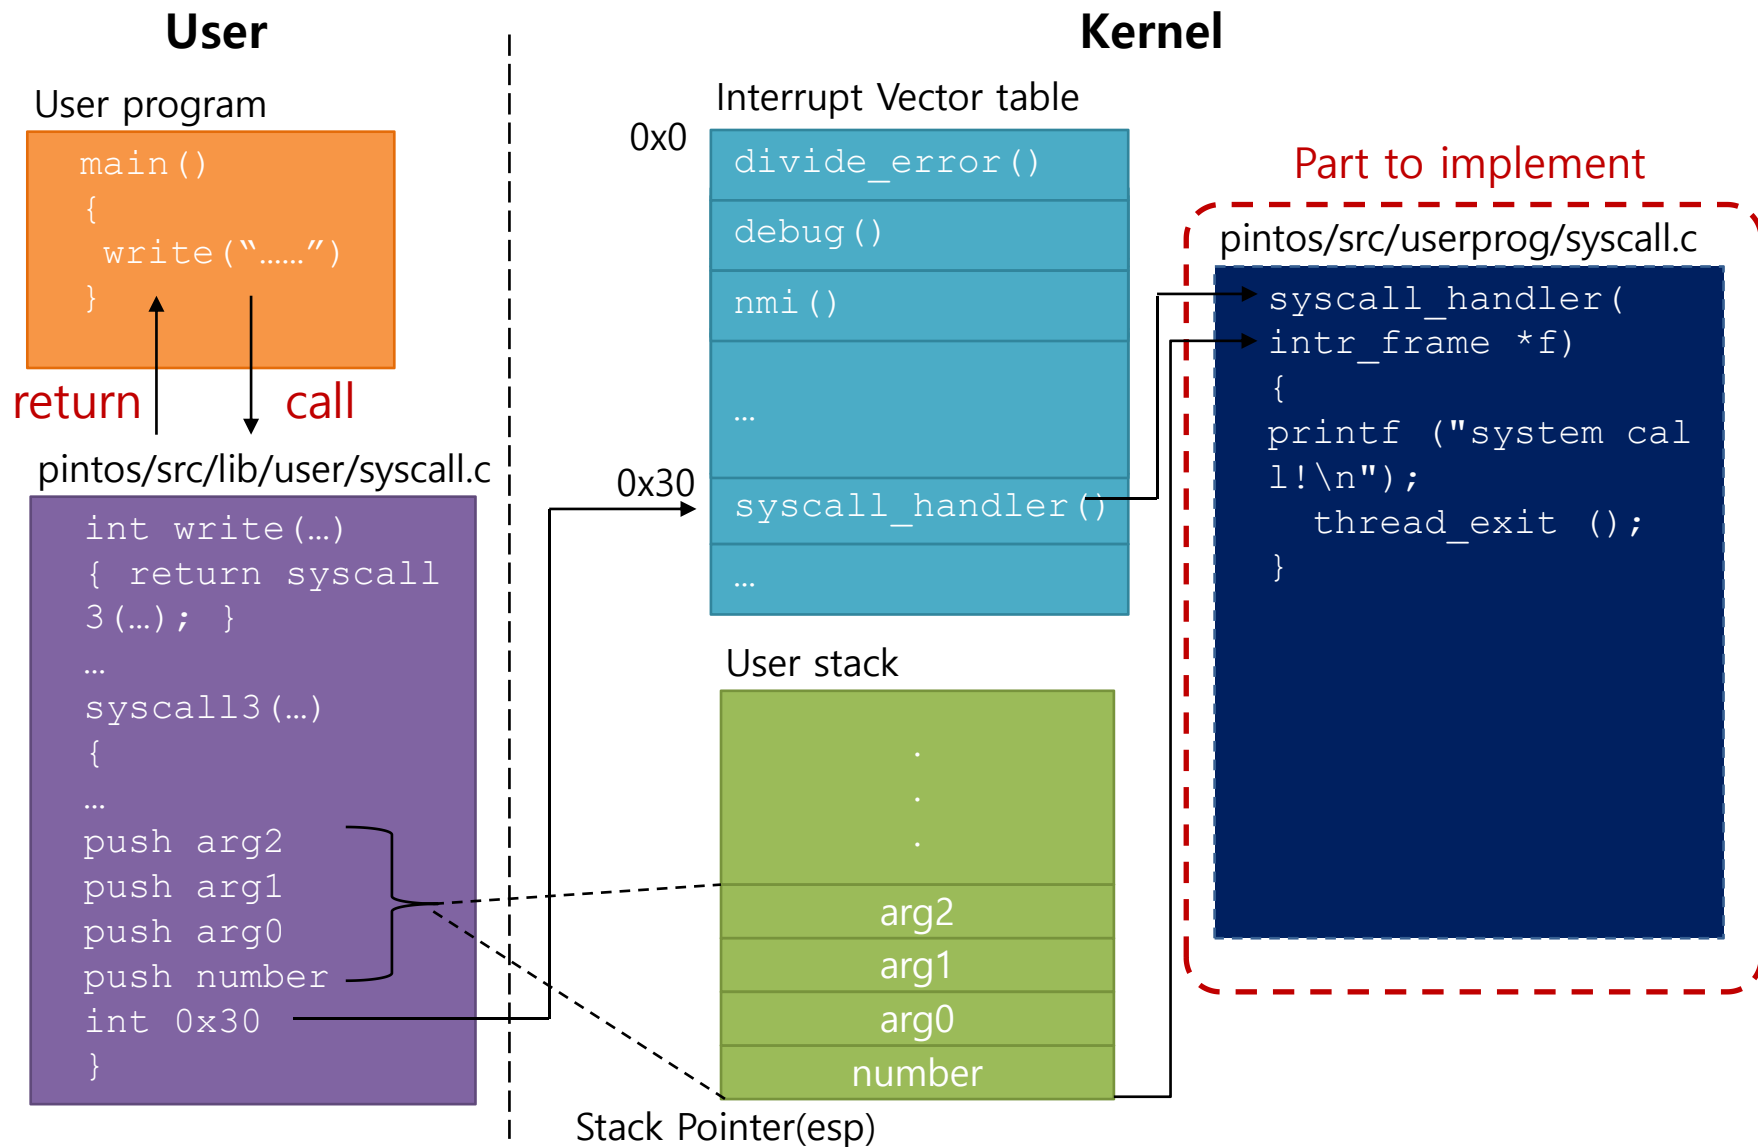
- ◆ `pintos/src/threads/thread.*`
- ◆ `pintos/src/userprog/syscall.*`
- ◆ `pintos/src/userprog/process.*`

# System call

- ▣ Programming interface for services provided by the operating system
- ▣ Allow user mode programs to use kernel features
- ▣ System calls run on kernel mode and return to user mode
- ▣ Key point of system call is that priority of execution mode is raised to the special mode as hardware interrupts are generated to call system call

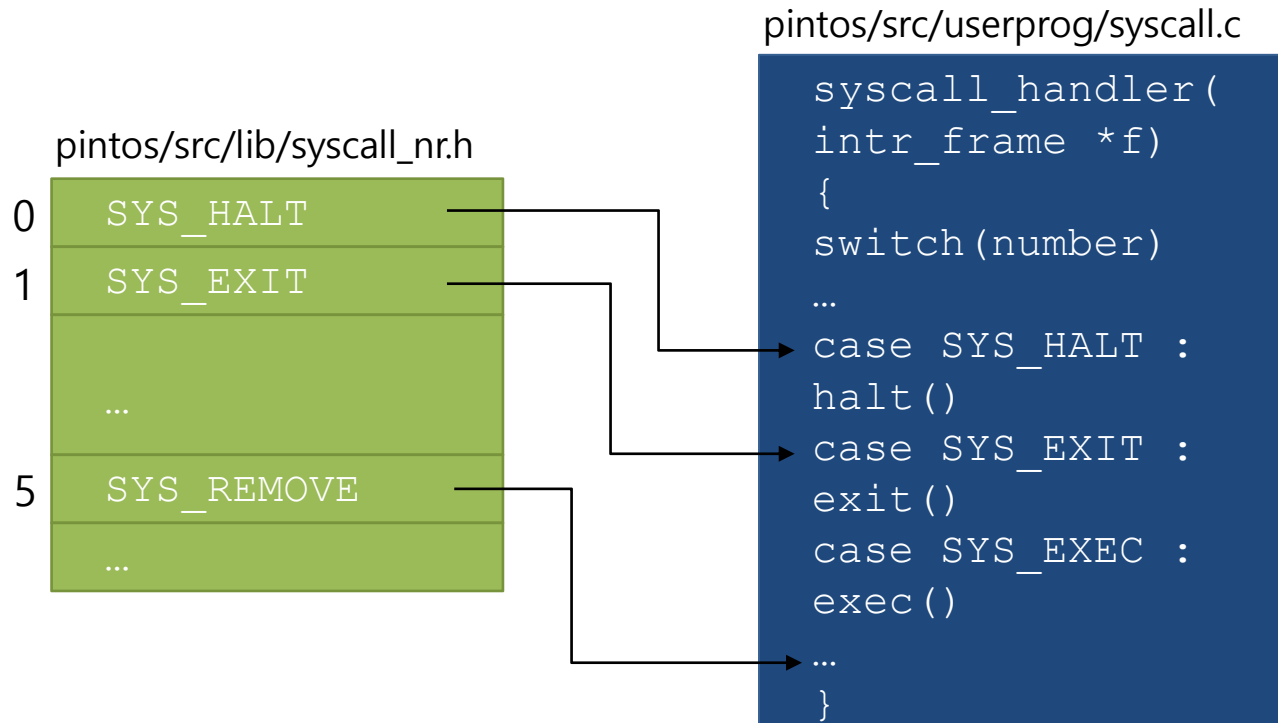


# Call process of System call (Pintos)



# System call handler

- Call the system call from the system call handler using the system call number.
  - The system call number is defined in `pintos/src/lib/syscall_nr.h`





# Requirement for System Call handler

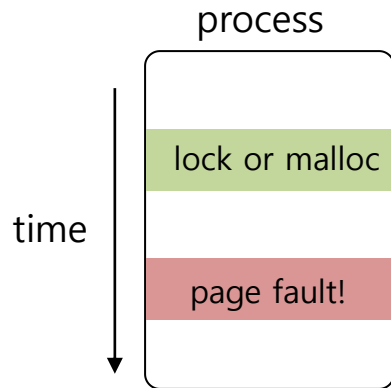
- ▣ Implement system call handler
  - ◆ Make system call handler call system call using system call number
  - ◆ Check validation of the pointers in the parameter list.
    - These pointers must point to user area, not kernel area.
    - If these pointers don't point the valid address, it is page fault
  - ◆ Copy arguments on the user stack to the kernel.
  - ◆ Save return value of system call at `eax` register.

# Address Validation

- ▣ User can pass invalid pointers through the syscall.
  - ◆ A null pointer / A pointer to unmapped virtual memory
  - ◆ A pointer to kernel virtual memory address space (above **PHYS\_BASE**)
- ▣ Kernel need to detect invalidity of pointers and terminating process without harm to the kernel or other running processes.
- ▣ How to detect?
  - ◆ Method 1: Verify the validity of a user-provided pointer.
    - The simplest way to handle user memory access.
    - Use the functions in 'userprog/pagedir.c' and in 'threads/vaddr.h'
  - ◆ Method 2: Check only that a user points below PHYS\_BASE.
    - An invalid pointer will cause 'page\_fault'. You can handle by modifying the code for `page_fault()`.
    - Normally faster than first one, Because it takes advantage of the MMU.
    - It tends to be used in real kernel.

# Accessing User Memory (cont.)

- ▣ In either case, make sure not to “leak” resource.



In the case, before terminating, we need to be sure release the lock or free the page.

- ▣ The first technique is straightforward.
  - ◆ Lock or allocate the page only after verifying the validity of pointers.
- ▣ The second one is more difficult.
  - ◆ Because there's no way to return an error code from a memory access.
  - ◆ You can use provided functions to handle these cases. (functions are in next slide.)

# Accessing User Memory (cont.)

```
/* Reads a byte at user virtual address UADDR.
   UADDR must be below PHYS_BASE.
   Returns the byte value if successful, -1 if a segfault
   occurred. */
static int
get_user (const uint8_t *uaddr)
{
    int result;
    asm ("movl $1f, %0; movzbl %1, %0; 1:"
        : "=a" (result) : "m" (*uaddr));
    return result;
}
```

```
/* Writes BYTE to user address UDEST.
   UDEST must be below PHYS_BASE.
   Returns true if successful, false if a segfault occurred.*/
static bool
put_user (uint8_t *udst, uint8_t byte)
{
    int error_code;
    asm ("movl $1f, %0; movb %b2, %1; 1:"
        : "=a" (error_code), "=m" (*udst) : "q" (byte));
    return error_code != -1;
}
```

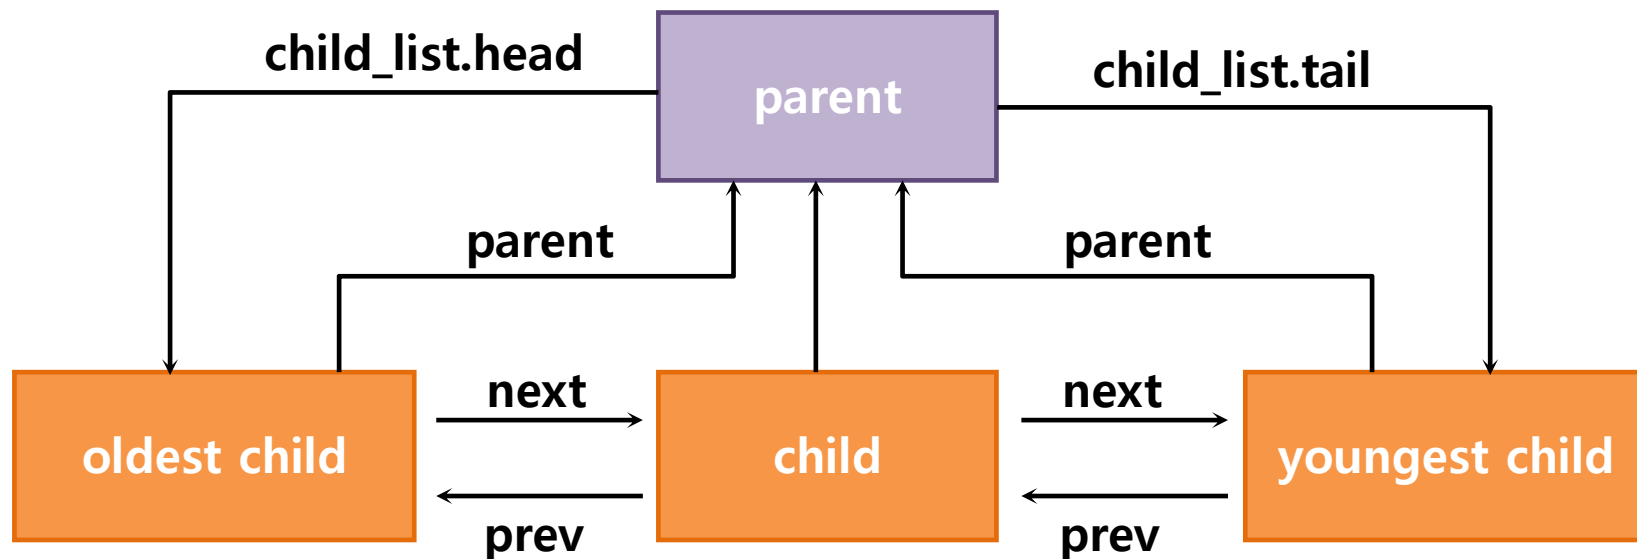
- ◆ You also modify the `page_fault ()`: set `eax` to `0xffffffff` and copies its former value into `eax`.

# Add system calls: Process related system calls

- ▣ `void halt(void)`
  - Shutdown pintos
  - Use `void shutdown_power_off(void)`
- ▣ `void exit(int status)`
  - Exit process
  - Use `void thread_exit(void)`
  - It should print message "Name of process: exit(status)".
- ▣ `pid_t exec (const char *cmd_line)`
  - Create child process and execute program corresponds to `cmd_line` on it
- ▣ `int wait (pid_t pid)`
  - Wait for termination of child process whose process id is `pid`

# Process Hierarchy

- ▣ Augment the existing process with the process hierarchy.
- ▣ To represent the relationship between parent & child,
  - ◆ Pointer to parent process: `struct thread*`
  - ◆ Pointers to the sibling. `struct list`
  - ◆ Pointers to the children: `struct list_elem`



# wait() system call

▣ `int wait(pid_t pid)`

- ◆ Wait for a child process `pid` to exit and retrieve the child's exit status.
- ◆ If `pid` is alive, wait till it terminates. Returns the status that `pid` passed to `exit`.
- ◆ If `pid` did not call `exit`, but was terminated by the kernel, return -1.
- ◆ A parent process can call `wait` for the child process that has terminated.
  - return exit status of the terminated child process.
- ◆ After the child terminates, the parent should deallocate its process descriptor
- ◆ `wait` fails and return -1 if
  - `pid` does not refer to a direct child of the calling process.
  - The process that calls `wait` has already called `wait` on `pid`.

# Kernel function for wait – process\_wait

```
int process_wait (tid_t child_tid UNUSED)
```

- ◆ It is currently empty.

```
int  
process_wait (tid_t child_tid UNUSED)  
{  
    return -1;  
}
```

- ◆ Insert the infinite loop so that the kernel does not finish. For now...



# Correct implementation: `process_wait()`

## ▣ `process_wait()`

- ◆ Search the descriptor of the child process by using `child_tid`.
- ◆ The caller blocks until the child process exits.
- ◆ Once child process exits, deallocate the descriptor of child process and returns exit status of the child process.

## ▣ Semaphore

- ◆ Add a semaphore for “wait” to thread structure.
- ◆ Semaphore is initialized to 0 when the thread is first created.
- ◆ In `wait(tid)`, call `sema_down` for the semaphore of `tid`.
- ◆ In `exit()` of process `tid`, call `sema_up`.
- ◆ Where do we need to place `sema_down` and `sema_up`?

## ▣ Exit status

- ◆ Add a field to denote the exit status to the thread structure.

# Flow of parent calling wait and child

## ▣ Flow of user program execution

—→ Flow  
- - -→ Scheduling

### Init Process

### User Process

Kernel Space

```
run_action()
  run_task()
    process_wait(process_excute())
      sema_down()
      waiting..
      waiting..
      waiting..
      waiting..
      waiting..
      Return exit status
  ...
shutdown_power_off()
Shutdown Pintos
```

Kernel Space

—  
—  
—

```
start_process()
  load()
  Run user program
...
exit()
  thread_exit()
  sema_up()
  exit process
```

Kernel Space

User Space

Kernel Space

# exec() system call

```
pid_t exec(const *cmd_line)
```

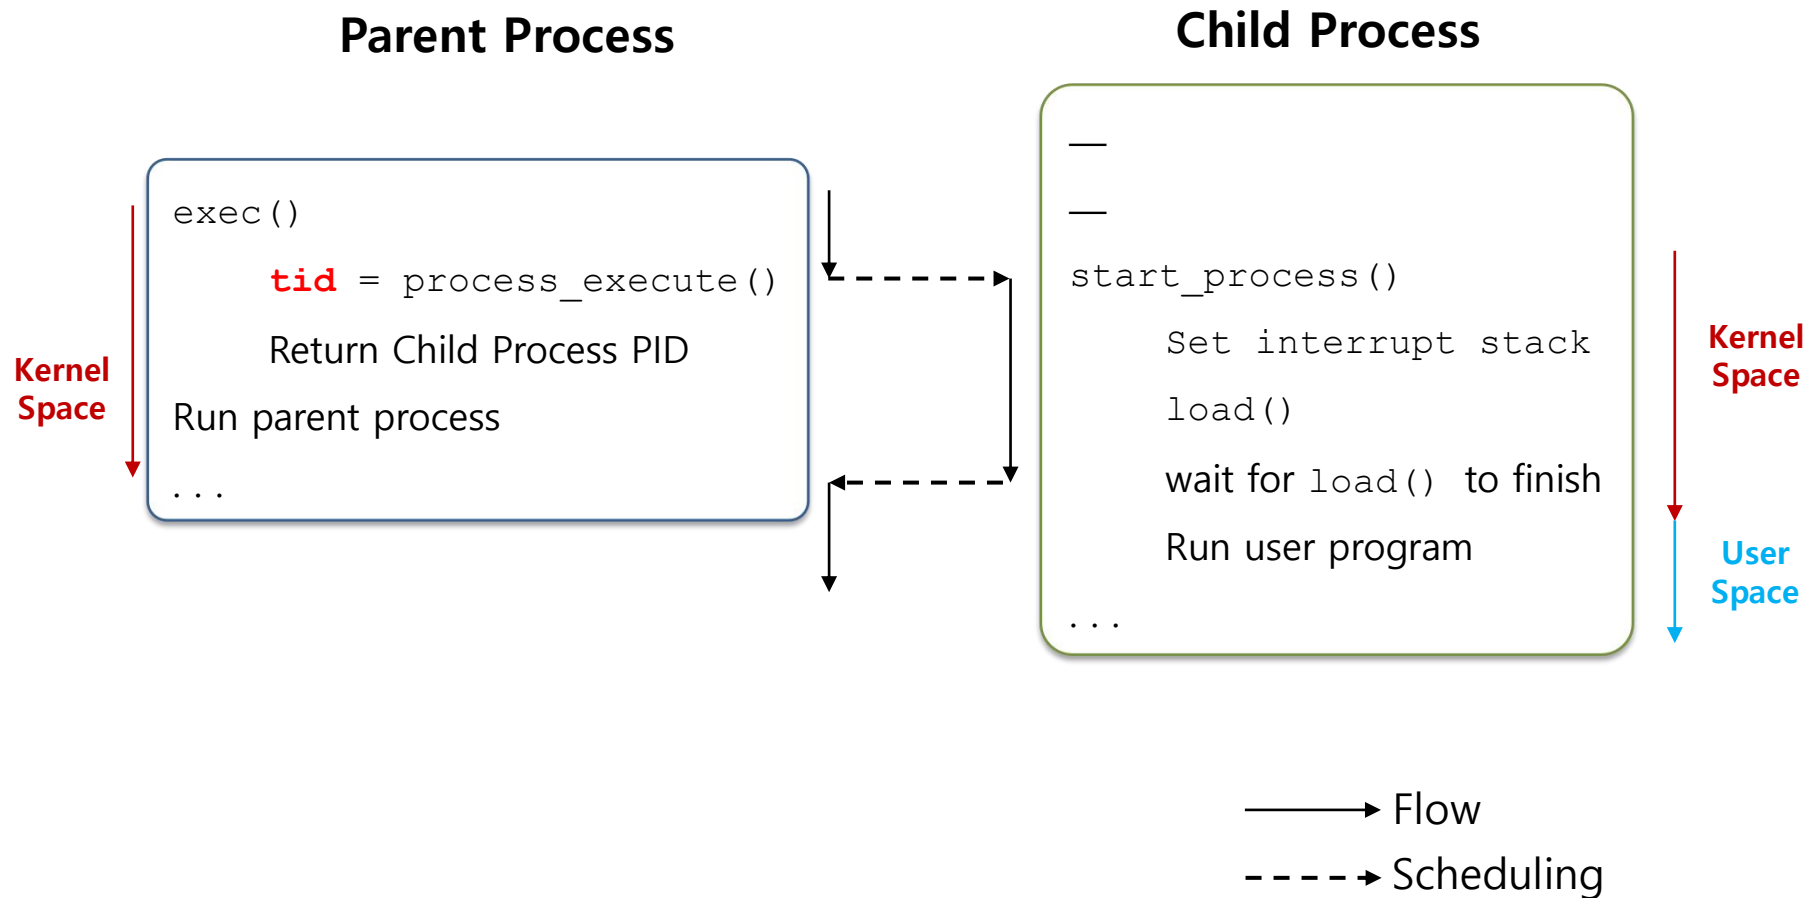
- ♦ Run program which execute `cmd_line`.
- ♦ Create thread and run. `exec()` in pintos is equivalent to `fork()+exec()` in Unix.
- ♦ Pass the arguments to program to be executed.
- ♦ Return `pid` of the new child process.
- ♦ If it fails to load the program or to create a process, return -1.
- ♦ Parent process calling `exec` should wait until child process is created and loads the executable completely.

# Kernel function for `exec()` : `process_execute()`

- ▣ Parent should wait until it knows the child process has successfully created and the binary file is successfully loaded.
- ▣ Semaphore
  - ◆ Add a semaphore for “`exec()`” to thread structure.
  - ◆ Semaphore is initialized by 0 when the thread is first created.
  - ◆ Call `sema_down` to wait for the successful load of the executable file of the child process.
  - ◆ Call `sema_up` when the executable file is successfully loaded.
  - ◆ Where do we need to place `sema_down` and `sema_up`?
- ▣ load status
  - ◆ In the thread structure, we need a field to represent whether the file is successfully loaded or not.

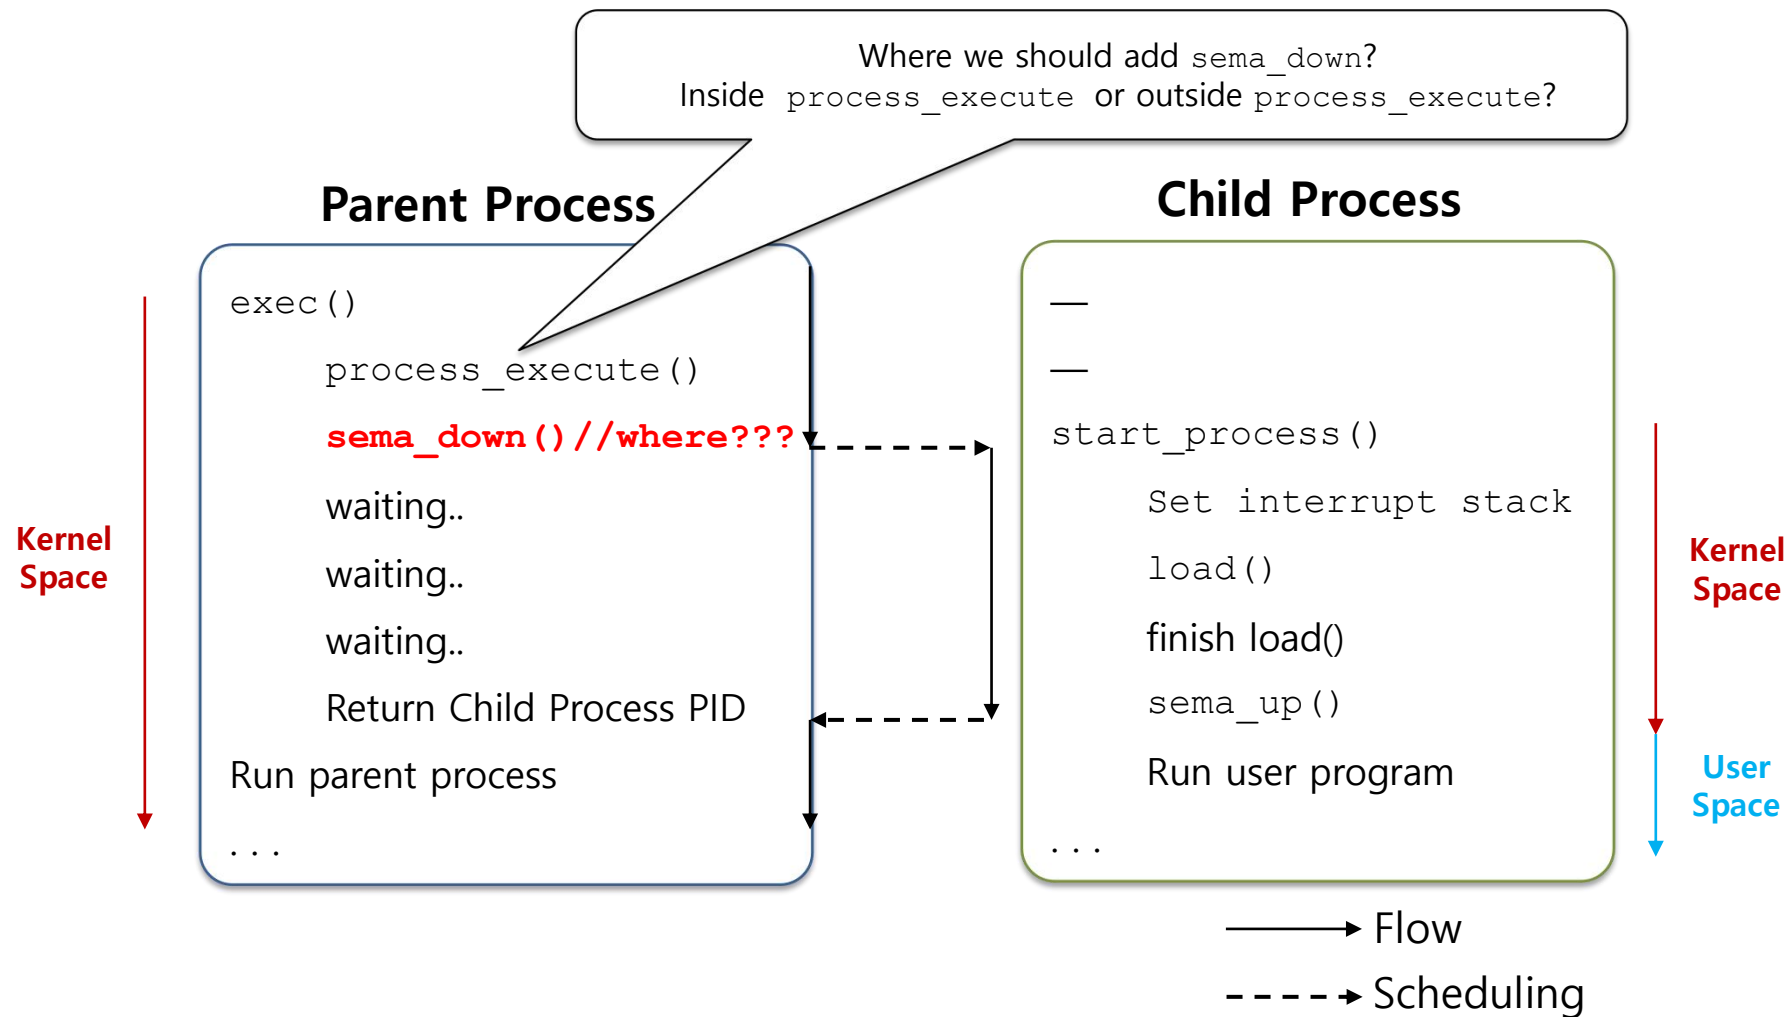
# Current flow of the parent calling exec and the child

- ▣ `exec()` return itself only after child is completely loaded.
- ▣ `tid` can have valid value even the load has failed.



# Correct Flow of the parent calling exec and the child

- ▣ `exec()` return itself only after child is completely loaded.



# exit()

- ▣ Terminate the current user program, returning status to the kernel.
- ▣ If the process' parent waits for it, this is the status that will be returned.

```
void exit (int status)
{
    struct thread *cur = thread_current ();
    /* Save exit status at process descriptor */
    printf("%s: exit(%d)\n" , cur -> name , status);
    thread_exit();
}
```

# Kernel function for `exit()`: `thread_exit`

- ▣ Exit status
  - ◆ Store the status to the status of process.
- ▣ Semaphore
  - ◆ Call `sema_up` for the current process.

```
/* Deschedules the current thread and destroys it.  Never
   returns to the caller. */
void
thread_exit (void)
{
    ASSERT (!intr_context ());

#ifdef USERPROG
    process_exit ();
#endif

    /* Remove thread from all threads list, set our status to dying,
       and schedule another process.  That process will destroy us
       when it calls thread_schedule_tail(). */
    intr_disable ();
    list_remove (&thread_current()->allelem);
    thread_current ()->status = THREAD_DYING;
    schedule ();
    NOT_REACHED ();
}
```



# File Manipulation

---

# File Descriptor in Unix

- Access to File by using File Descriptor

Process Descriptor Table




Part to implement

File Descriptor Table


0	: STDIN
1	: STDOUT
2	: STDERR
3	: File
4	: File
	.
	.
	.

Keyboard File Object



keyboard


Monitor File Object



Monitor

Monitor File Object

I/O File Object



Disk

# File Descriptor Table

## □ Implement File Descriptor Table.

- ◆ Each process has its own file descriptor table (Maximum size: 64 entry).
- ◆ File descriptor table is an array of pointer to struct file.
- ◆ FD is index of the file descriptor table, and it is allocated sequentially.
  - FD 0 and 1 are allocated for `stdin` and `stdout`, respectively.
- ◆ `open()` returns `fd`.
- ◆ `close()` set 0 at file descriptor entry at index `fd`.

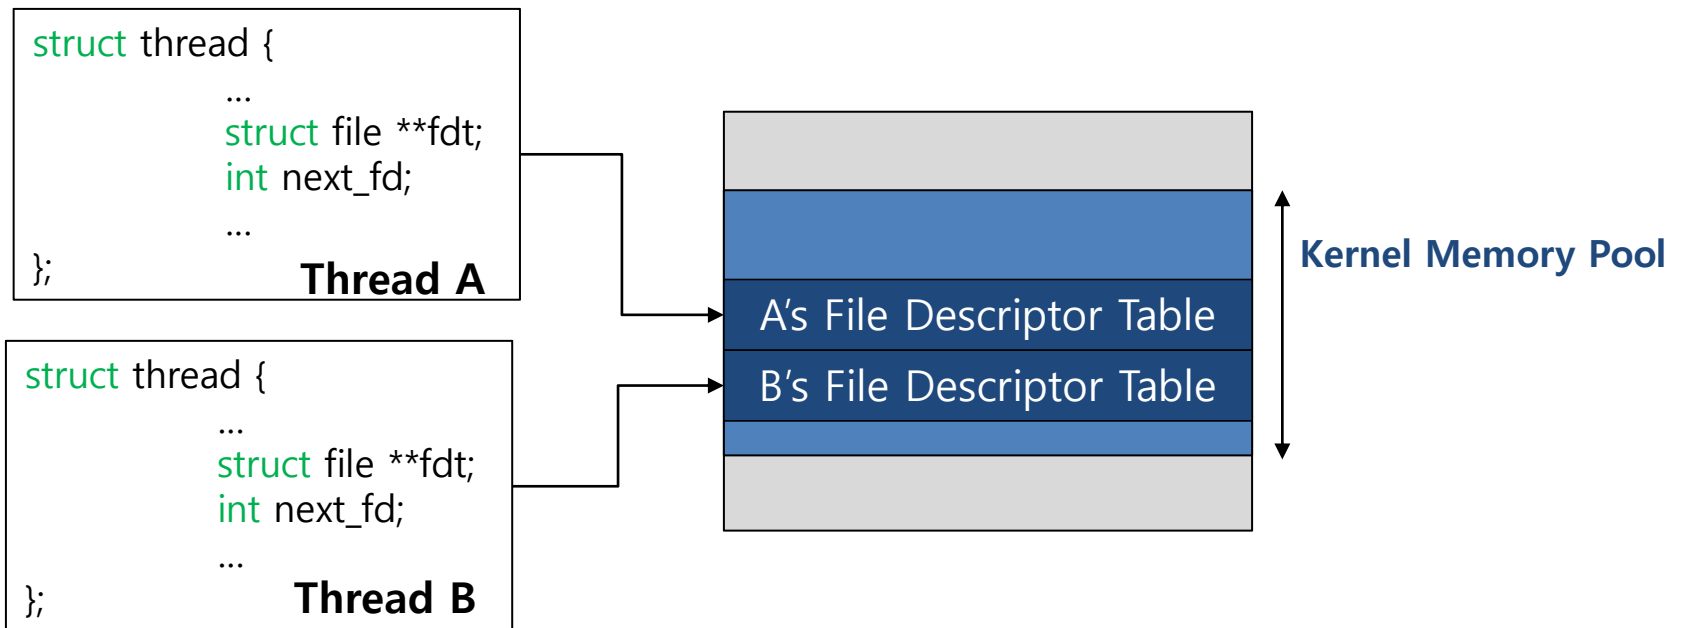
# Allocate File Descriptor Table

- Define FDT as a part of thread structure.

```
struct thread {  
    ...  
    struct file fdt[64];  
    int next_fd;  
    ...  
};
```

**Thread A**

- Allocate FDT at kernel memory area, and add the associated pointer to at the thread structure.



# File Descriptor Table

- ▣ When the thread is created,
  - ◆ Allocate File Descriptor table.
  - ◆ Initialize pointer to file descriptor table.
  - ◆ Reserve fd0, fd1 for stdin and stdout.
- ▣ When thread is terminated,
  - ◆ Close all files.
  - ◆ Deallocate the file descriptor table.
- ▣ Use global lock to avoid race condition on file,
  - ◆ Define a global lock on syscall.h (`struct lock filesys_lock`).
  - ◆ Initialize the lock on `syscall_init()` (Use `lock_init()`).
  - ◆ Protect filesystem related code by global lock.

# Modify page\_fault() for test

- Some tests check whether your kernel handles the bad process properly.
- Pintos needs to kill the process and print the thread name and the exit status -1 when page fault occurs.
- We have to modify `page_fault()` to satisfy test's requirements.

pintos/src/userprog/exception.c

```
static void page_fault (struct intr_frame *f)
{
    ...
    not_present = (f->error_code & PF_P) == 0;
    write = (f->error_code & PF_W) != 0;
    user = (f->error_code & PF_U) != 0;

    /* Call exit(-1) */
    ...
}
```

# Add system calls: File related system calls

- ▣ `bool create(const char *file, unsigned initial_size)`
  - ◆ Create file which have size of `initial_size`.
  - ◆ Use `bool filesystem_create(const char *name, off_t initial_size)`.
  - ◆ Return true if it is succeeded or false if it is not.
- ▣ `bool remove(const char *file)`
  - ◆ Remove file whose name is `file`.
  - ◆ Use `bool filesystem_remove(const char *name)`.
  - ◆ Return true if it is succeeded or false if it is not.
  - ◆ File is removed regardless of whether it is open or closed.
- ▣ `int open(const char *file)`
  - ◆ Open the file corresponds to path in “`file`”.
  - ◆ Return its fd.
  - ◆ Use `struct file *filesystem_open(const char *name)`.

# Add system calls: File related system calls (Cont.)

▣ `int filesize(int fd)`

- ◆ Return the size, in bytes, of the file open as `fd`.
- ◆ Use `off_t file_length(struct file *file)`.

▣ `int read(int fd, void *buffer, unsigned size)`

- ◆ Read `size` bytes from the file open as `fd` into `buffer`.
- ◆ Return the number of bytes actually read (0 at end of file), or -1 if fails.
- ◆ If `fd` is 0, it reads from keyboard using `input_getc()`, otherwise reads from file using `file_read()` function.
  - `uint8_t input_getc(void)`
  - `off_t file_read(struct file *file, void *buffer, off_t size)`



# Add system calls: File related system calls (Cont.)

▣ `int write(int fd, const void *buffer, unsigned size)`

- ◆ Writes `size` bytes from `buffer` to the open file `fd`.
- ◆ Returns the number of bytes actually written.
- ◆ If `fd` is 1, it writes to the console using `putbuf()`, otherwise write to the file using `file_write()` function.
  - `void putbuf(const char *buffer, size_t n)`
  - `off_t file_write(struct file *file, const void *buffer, off_t size)`

▣ `void seek(int fd, unsigned position)`

- ◆ Changes the next byte to be read or written in open file `fd` to `position`.
- ◆ Use `void file_seek(struct file *file, off_t new_pos)`.

# Add system calls: File related system calls (Cont.)

▣ `unsigned tell(int fd)`

- ◆ Return the position of the next byte to be read or written in open file `fd`.
- ◆ Use `off_t file_tell(struct file *file)`.

▣ `void close(int fd)`

- ◆ Close file descriptor `fd`.
- ◆ Use `void file_close(struct file *file)`.

# Denying writes to executable

- ▣ What if the OS tries to execute the file that is being modified?
- ▣ Do not allow the file to be modified when it is opened for execution.
- ▣ Approach
  - ◆ When the file is loaded for execution, call `file_deny_write()`.
  - ◆ When the file finishes execution, call `file_allow_write()`.

```
static bool load (const char *cmdline, void (**eip) (void), void **esp)
```

- ◆ Call `file_deny_write()` when program file is opened.
- ◆ Add a running file structure to thread structure.

```
void process_exit (void)
```

- ◆ Modify current process to close the running file.

# Result

- ▣ Check the result of all tests.
  - ◆ Path: pintos/src/userprog

\$make grade

## SUMMARY BY TEST SET

Test Set	Pts	Max	% Ttl	% Max
tests/userprog/Rubric.functionality	108/108		35.0%/	35.0%
tests/userprog/Rubric.robustness	88/	88	25.0%/	25.0%
tests/userprog/no-vm/Rubric	1/	1	10.0%/	10.0%
tests/filesys/base/Rubric	30/	30	30.0%/	30.0%
Total			100.0%/	100.0%

# Summary

