

# Operating Systems Lab

## Part 3: Virtual Memory

---

**KAIST EE**

**Youjip Won**

# Overview of Virtual Memory

- ▣ Background of Virtual Memory in Pintos
- ▣ Requirements
  - ◆ Paging(swapping)
  - ◆ Growing stack
  - ◆ Memory mapped file
  - ◆ Accessing user memory

# Swapping


---

# To Do's

- ❑ Implement data structure to represent physical page frame.
- ❑ Implement page replacement policy such as LRU, clock, second-chance
- ❑ swapping
  - ◆ Store victim pages in swap space when they belong to data segment or stack segment.
  - ◆ swap-out pages are reloaded into memory by demand paging.

# Hardware Support

- ❑ The dirty bit of page table is set to “1” by hardware when writing to the memory space
- ❑ The accessed bit in page table is set to ‘1’ by hardware each time the page is referenced



- ❑ When page with dirty bit “1” is selected as victim, the changes must always be stored on disk
- ❑ Hardware does not re-zero the accessed bit.

# Page Table Manipulation in Pintos (userprog/pagedir.c)

- ◆ `bool pagedir_is_dirty (uint32_t *pd, const void *vpage)`
  - Return dirty bit of pte for vpage in pd
- ◆ `void pagedir_set_dirty (uint32_t *pd, const void *vpage, bool dirty)`
  - Set the dirty bit to `dirty` in the pte for vpage in pd
- ◆ `bool pagedir_is_accessed (uint32_t *pd, const void *vpage)`
  - Return access bit of pte for vpage in pd
- ◆ `void pagedir_set_accessed (uint32_t *pd, const void *vpage, bool accessed)`
  - Set the access bit to `accessed` in the pte for vpage in pd

# struct page: New data structure required

Select the physical page frame for replacement.

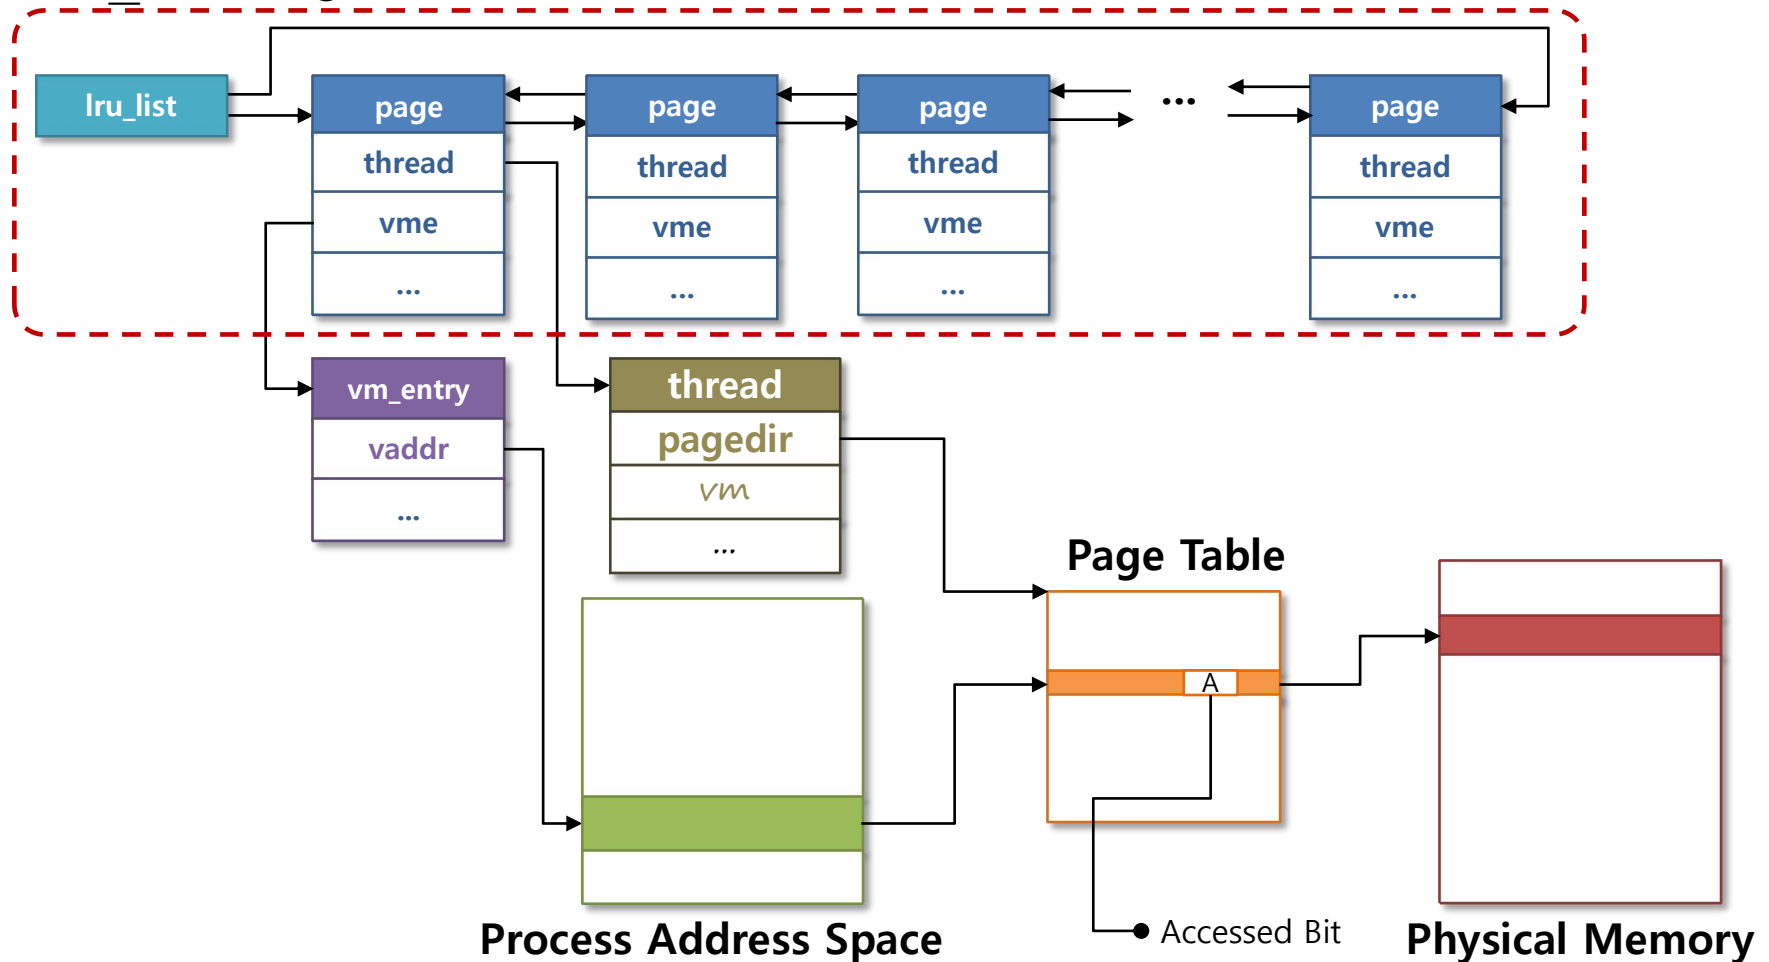
- ▣ Data structure representing each physical page that contains a user page
  - ◆ physical address of page
  - ◆ reference to the virtual page object to which physical page is mapped
  - ◆ Reference to the thread structure to which it belongs
  - ◆ lru: field for list

pintos/src/vm/page.h

```
struct page {  
  
    // fill this out  
  
};
```

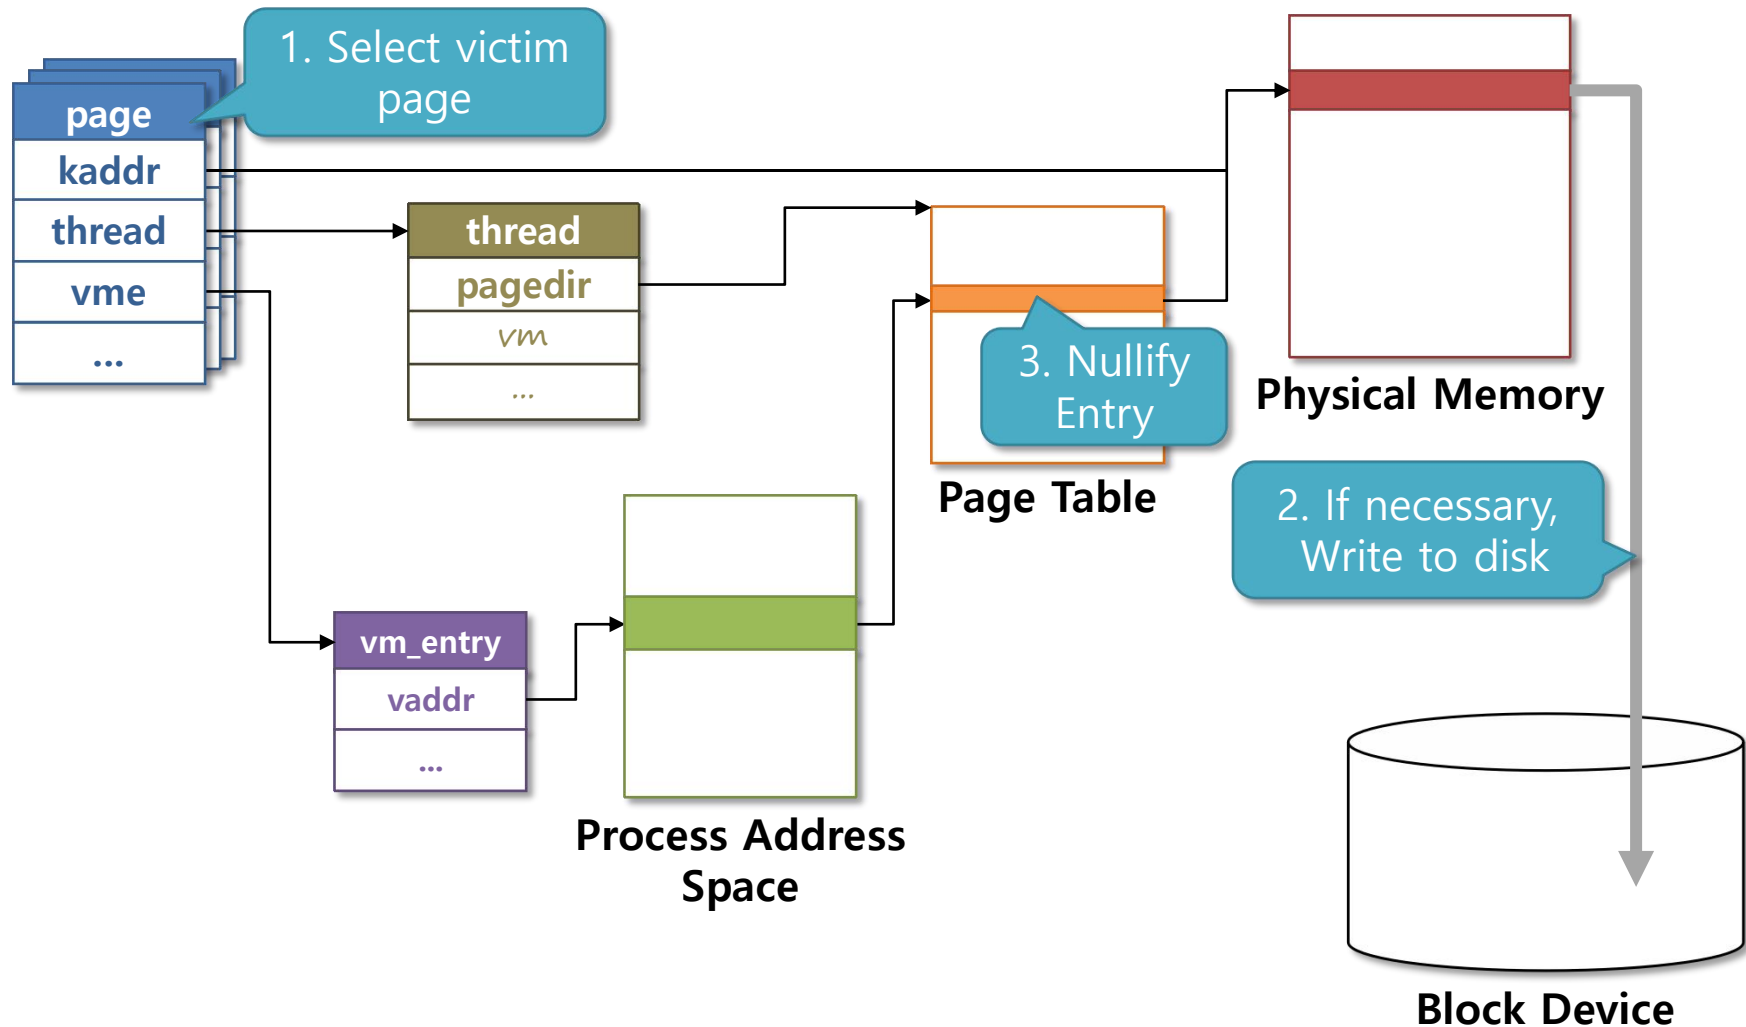
# A page pool for swapping

- Manage physical pages in use as a list of pages.
- `lru_list` : global variable



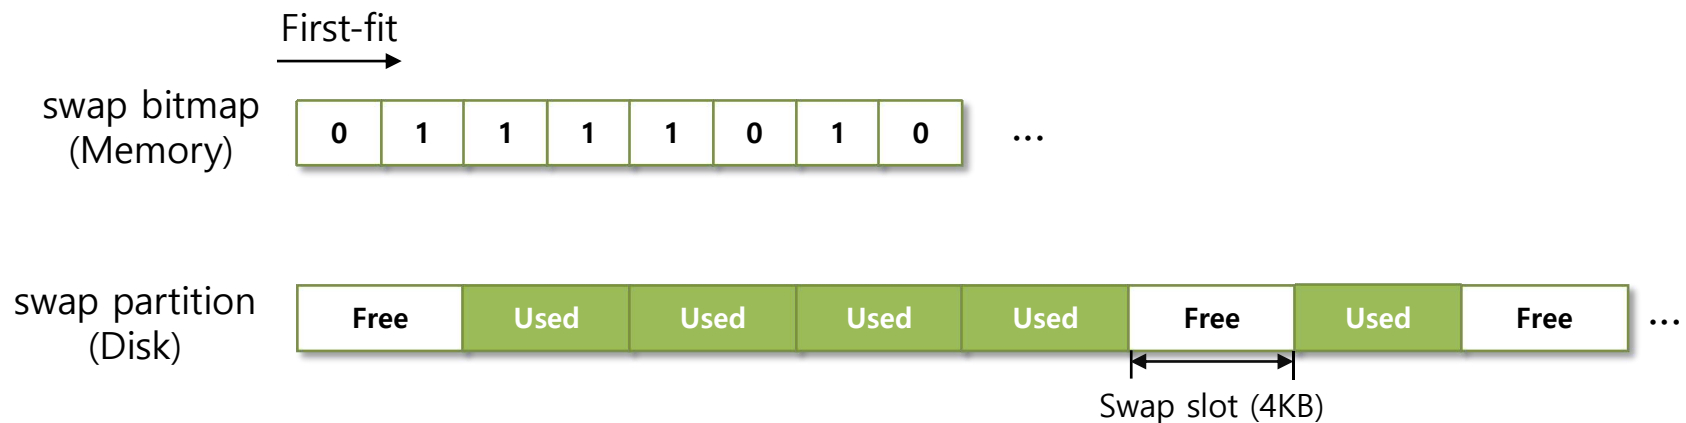


# Swap-out



# Managing swap partition

- ❑ Swap partition is managed per swap slot unit(4 Kbyte).
- ❑ Maintaining a swap partition: swap bitmap (global variable in memory)
- ❑ Search bitmap for free slot.
- ❑ What happens to swap bitmap if the system crashes?



# Functions offered by pintos for swap space manipulation

- ▣ Swap partition is provided as block device in pintos.
- ▣ Functions for block device (`src/block/block.c`)
  - ◆ `struct block *block_get_role (enum block_type role)`
    - Return the block device (`struct block *`) fulfilling the given ROLE.
    - ROLES defined in pintos now (`devices/block.h`)
      - `BLOCK_KERNEL`: OS Partition
      - `BLOCK_FILESYS`: File system
      - `BLOCK_SCRATCH`: Scratch partition
      - `BLOCK_SWAP`: Swap partition
  - ◆ `void block_read (struct block *block, block_sector_t sector, void *buffer)`
    - Read contents at sector on block and save them at buffer
  - ◆ `void block_write (struct block *block, block_sector_t sector, const void *buffer)`
    - Write contents at buffer at sector on block

# Implementation

- ❑ LRU list for physical page frame
  - ◆ List of `struct page`
  - ◆ List of physical pages allocated to user process
- ❑ functions for allocate/release physical page frame from the list
  - ◆ When there runs out of physical page frame, select a victim and swap it out.
- ❑ Modify page fault handler for swapping.
  - ◆ Before: Physical page is allocated directly when page fault occurs.  
When there is no page to allocate, pintos is finished.
  - ◆ After: Physical page is allocated from LRU list when page fault occurs.  
When there is no page to allocate, pintos swap in the page.

# Functions to write

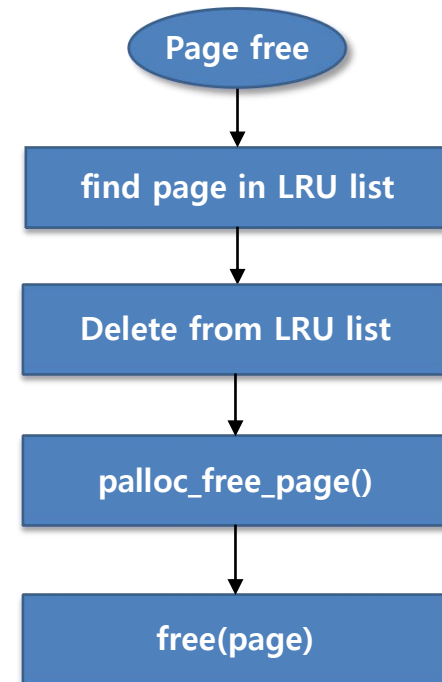
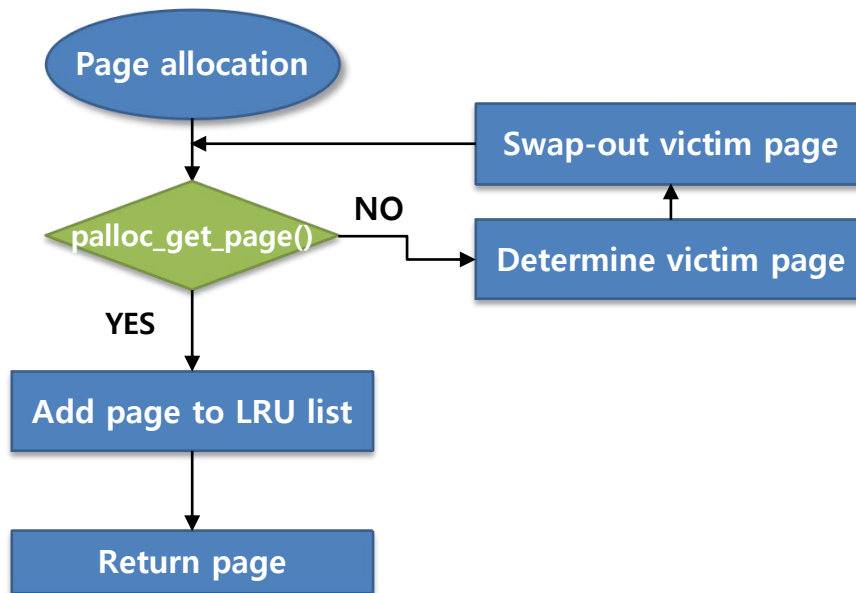
- ❑ Function about LRU list (initializing, insert, remove).
- ❑ Function to allocate a page from LRU list.
- ❑ Function to free page from LRU list.
- ❑ Function to select victim page and swap-out the page.
  - ◆ e.g.: Clock algorithm, Second chance algorithm
- ❑ Function about swapping (initializing, swap in, swap out).

# Functions to modify

- ▣ `bool handle_mm_fault(struct vm_entry *vme)`
  - ◆ Modify to allocate physical pages from LRU list when page fault occurs
  - ◆ Modify to swap-in if `vm_entry` type is `VM_ANON`
  
- ▣ `static bool setup_stack(void **esp)`
  - ◆ Modify to allocate pages from LRU list when page fault occurs
  
- ▣ `int main(void)`
  - ◆ Initialize LRU list.

# Functions for allocation/free page

- Try to obtain free space when memory cannot be allocated through `palloc_get_page()` within the page allocation function.



# Swap-out

## ▣ Type of a page in the physical page frame

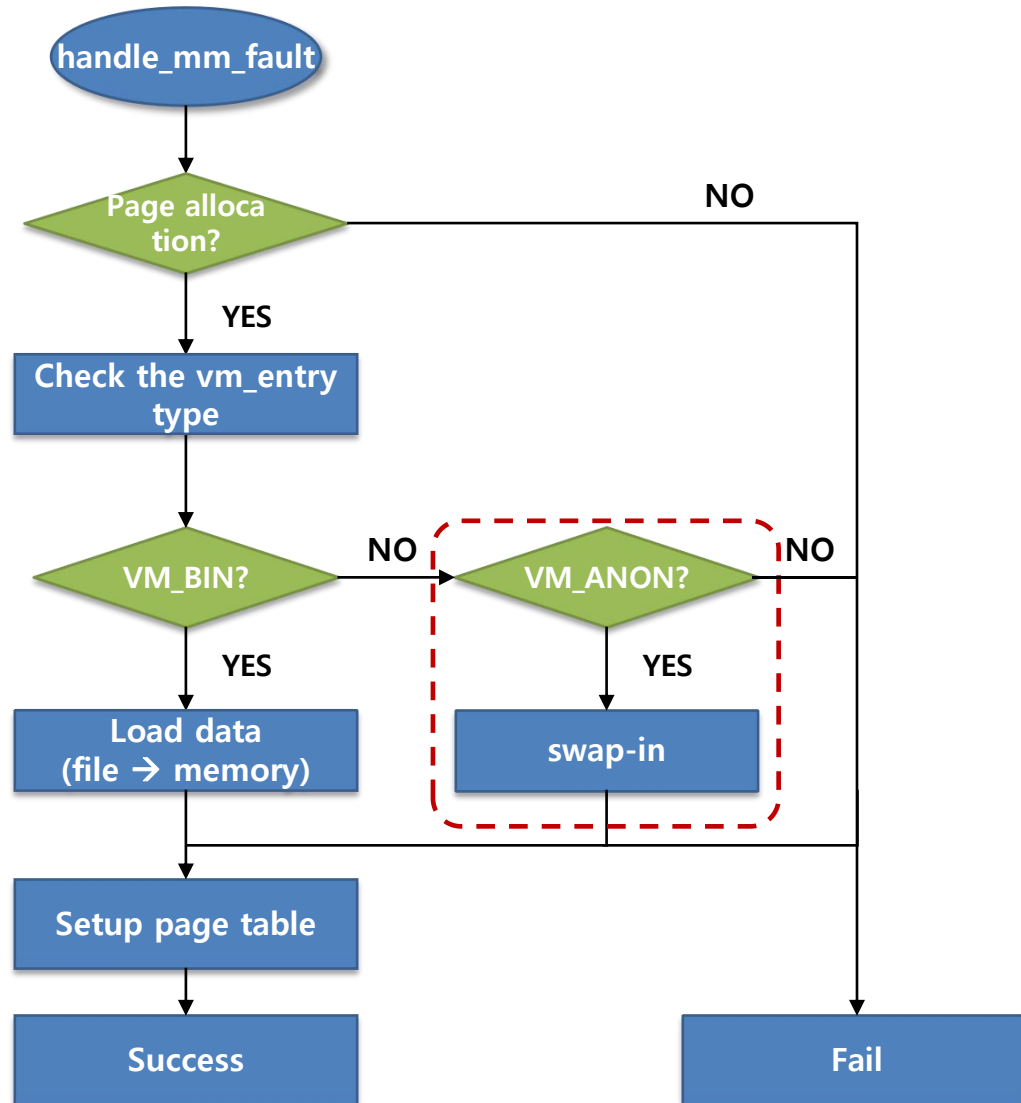
- ◆ VM\_BIN
  - If dirty bit is “1”, write to the swap partition and free the page frame.
  - Change type to VM\_ANON for demand paging
- ◆ VM\_FILE
  - If dirty bit is “1”, write the page to the file and free the page frame.
  - If dirty bit is “0”, free the page frame.
- ◆ VM\_ANON
  - Write to the swap partition.

## ▣ Mark the page “not present” in pd (page directory).

```
void pagedir_clear_page (uint32_t *pd, void *upage)
```



# Demand paging for anonymous page (stack or heap)



# Modify `handle_mm_fault()`

- ▣ If `vm_entry` type is `VM_ANON`, modify code to swap in

`pintos/src/userprog/process.c`

```
bool handle_mm_fault(struct vm_entry *vme) {
    bool success = false;
    void *kaddr;
    ...
    switch(vme->type) {
        case VM_BIN:
            success = load_file(kaddr, vme);
            break;

        case VM_ANON:
            /* insert swap in code */
            break;
    }
    ...
}
```

# Growing Stack

---


# Expandable Stack

- Implement expandable stack

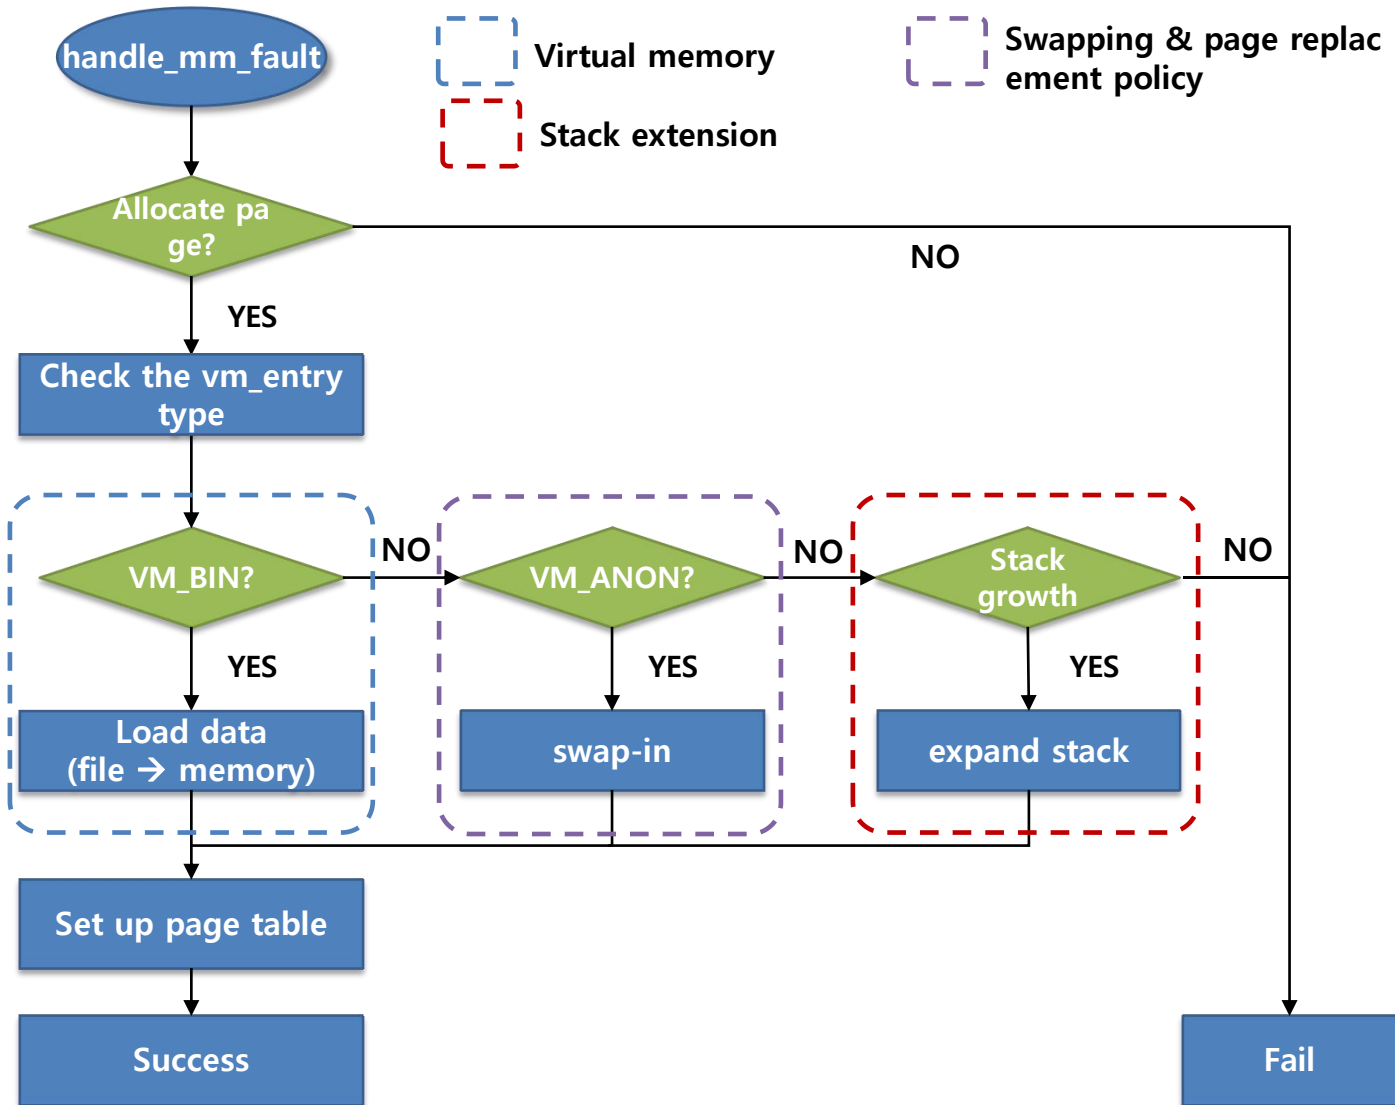
- ◆ In current pintos, stack size is fixed to 4KB.
- ◆ Make the stack expandable.
  - If a process accesses the address that lies outside the stack and that can be handled by expanding the stack, expand the stack.
    - e.g. (access address < stack pointer – 32) Expand stack
- ◆ maximum size of stack is 8MB.

# When to expand stack

- Expand the stack when the memory access is within 32 Byte of stack top.
  - “PUSHA” instruction in 80x86 pushes 32 bytes at once.



# Stack extension mechanism




# Memory Mapped File

---

# mmap vs. munmap

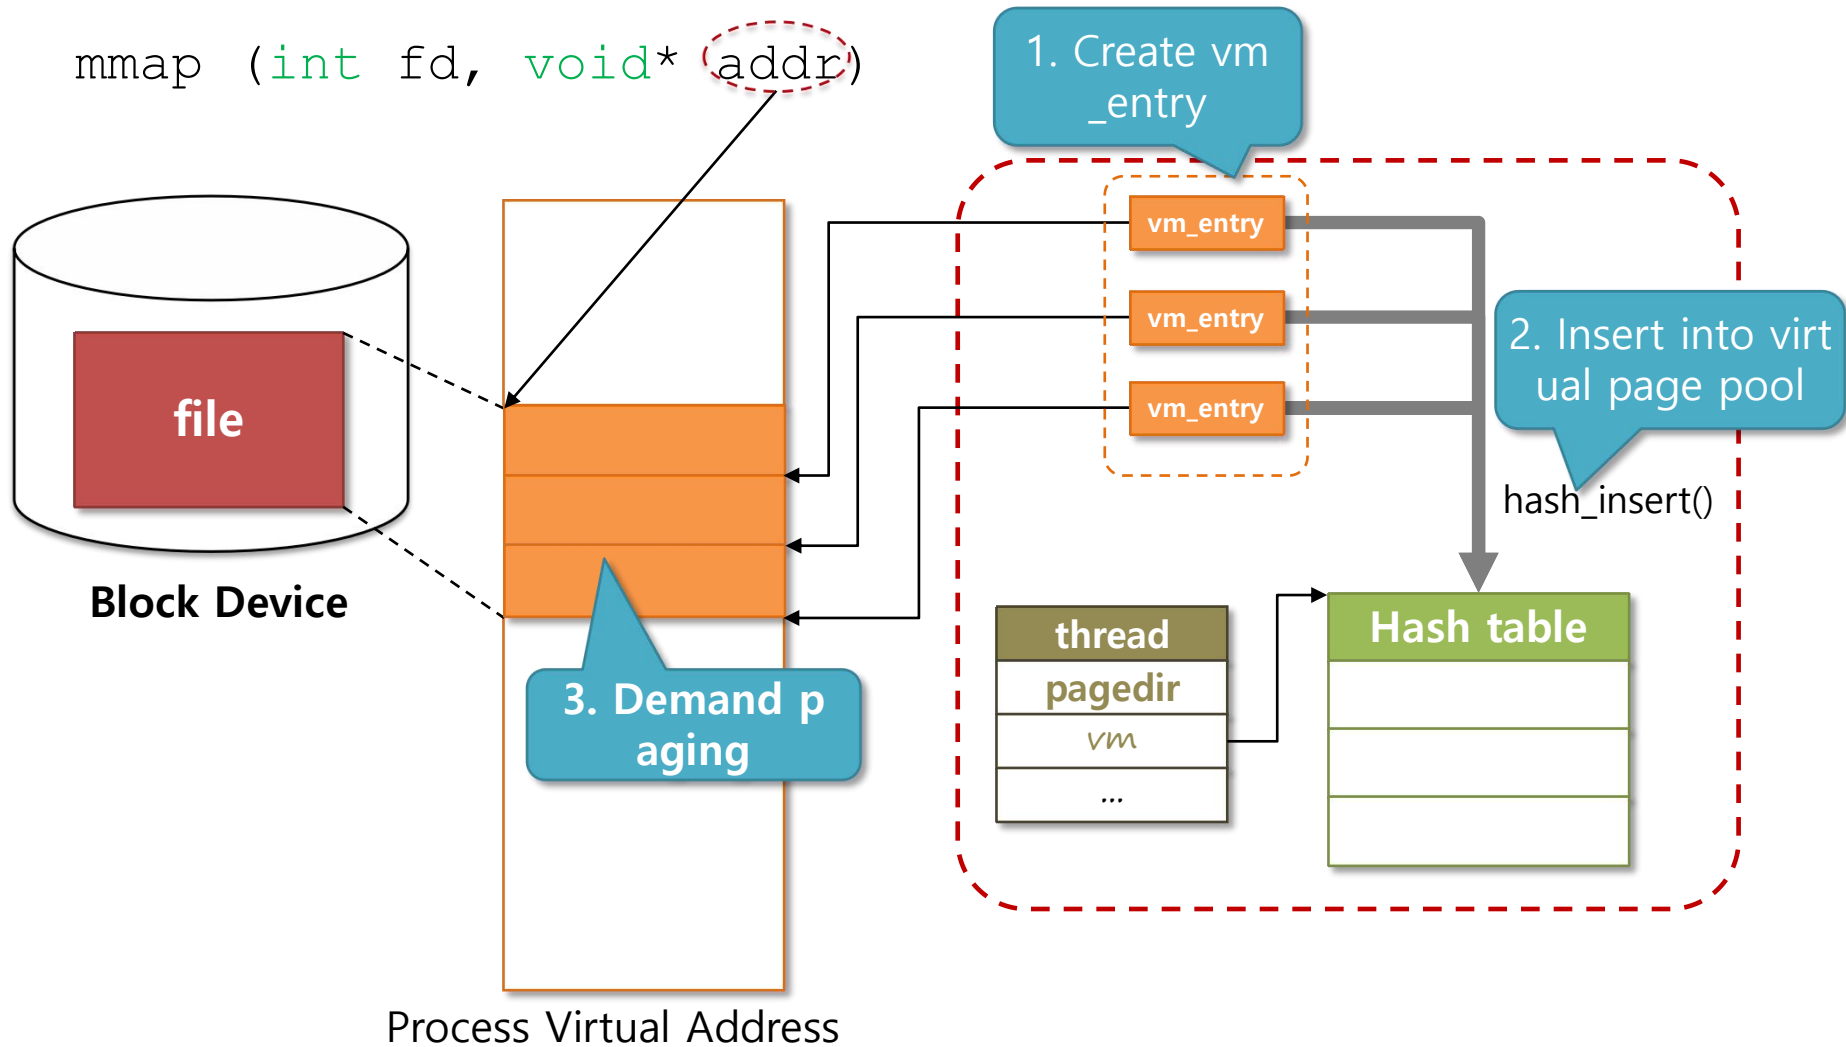
```
main (int argc, char *argv[]) {
    int i;

    for (i = 1; i < argc; i++)    {
        int fd;
        mapid_t map;
        void *data = (void *) 0x10000000;
        int size;
        fd = open (argv[i]);
        size = filesize (fd);
        map = mmap (fd, data);
        write (STDOUT_FILENO, data, size);
        munmap (map);
    }
    return EXIT_SUCCESS;
}
```





# mmap and munmap



# mmap() and munmap()

□ `int mmap(int fd, void *addr)`

- ◆ Load file data into memory by demand paging.
- ◆ `mmap()`'ed page is swapped out to its original location in the file.
- ◆ For a fragmented page, fill the unused fraction of page with zero.
- ◆ Return `mapping_id`: unique id within a process to identify the mapped file.
- ◆ Fails if
  - File size is 0.
  - Addr is not page aligned.
  - Address is already in use.
  - Addr is 0.
  - STDIN and STDOUT are not mappable..

□ `void munmap(mapping_t mapid)`

- ◆ Unmap the mappings in the `mmap_list` which has not been previously unmapped.

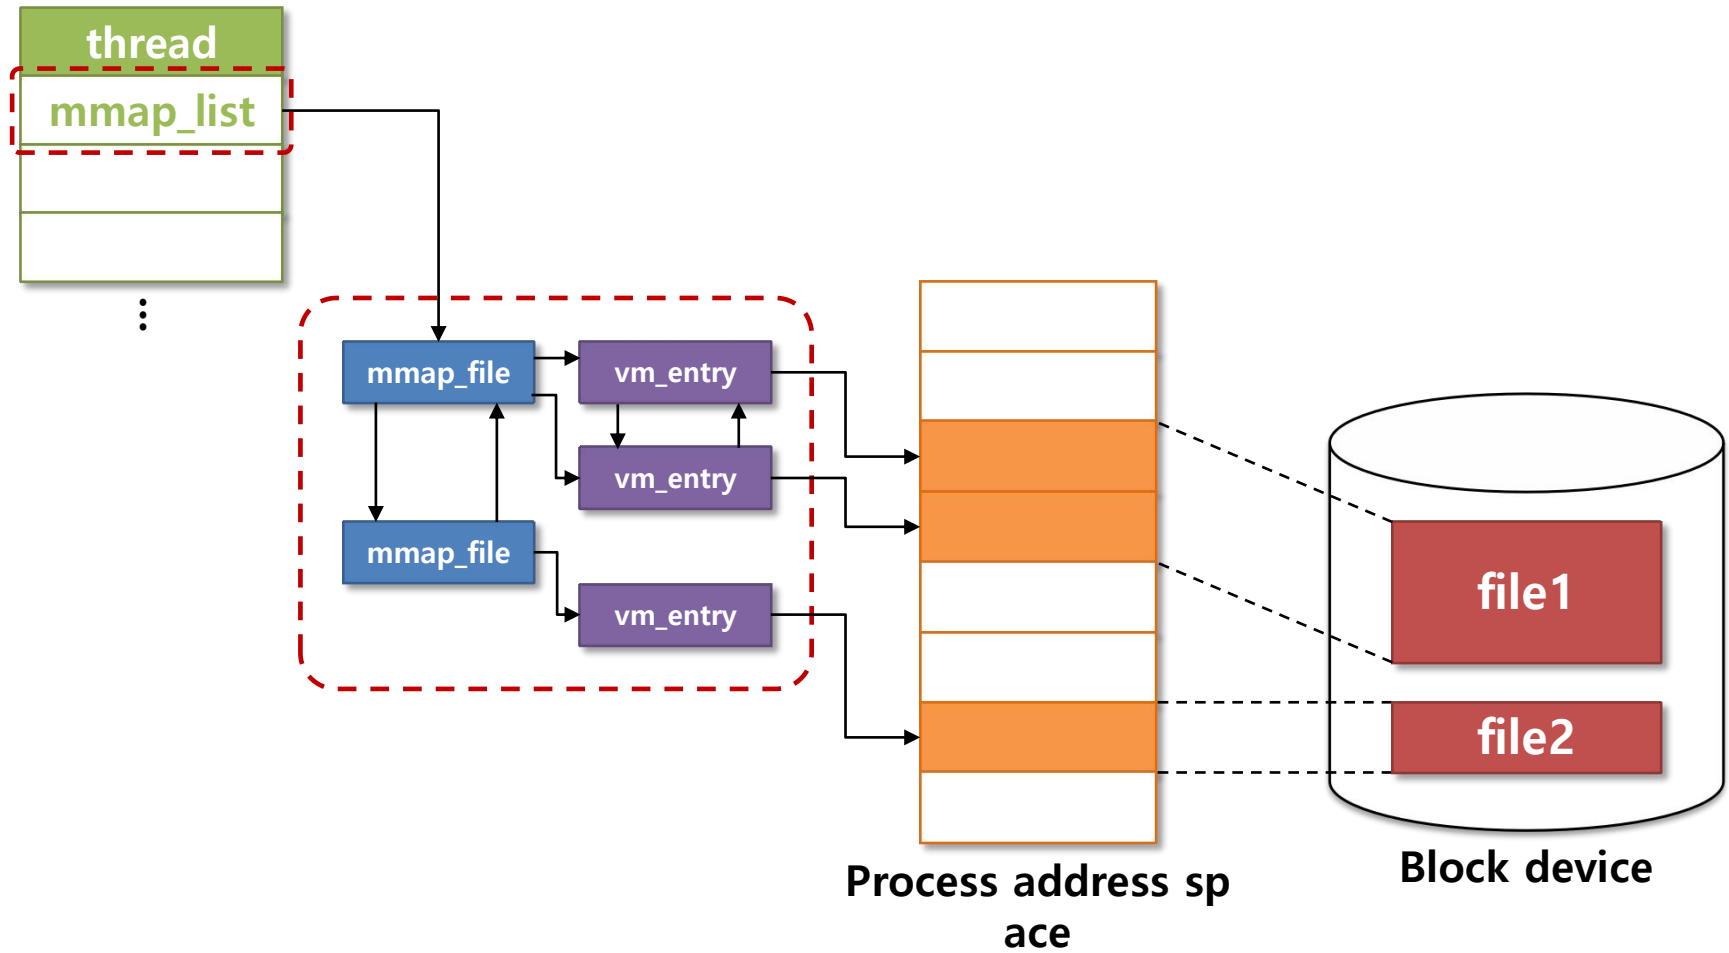
# Requirements

- ▣ All mappings of a process are implicitly unmapped when the process exits.
- ▣ When a mapping is unmapped, the pages are written back to the file.
- ▣ Upon munmap, the pages are removed from the process' virtual page list.
- ▣ Once created, mapping is valid until it is unmapped regardless of the file is closed or deleted.
- ▣ If the two or more processes map the same file, they do not have to see the consistent view.

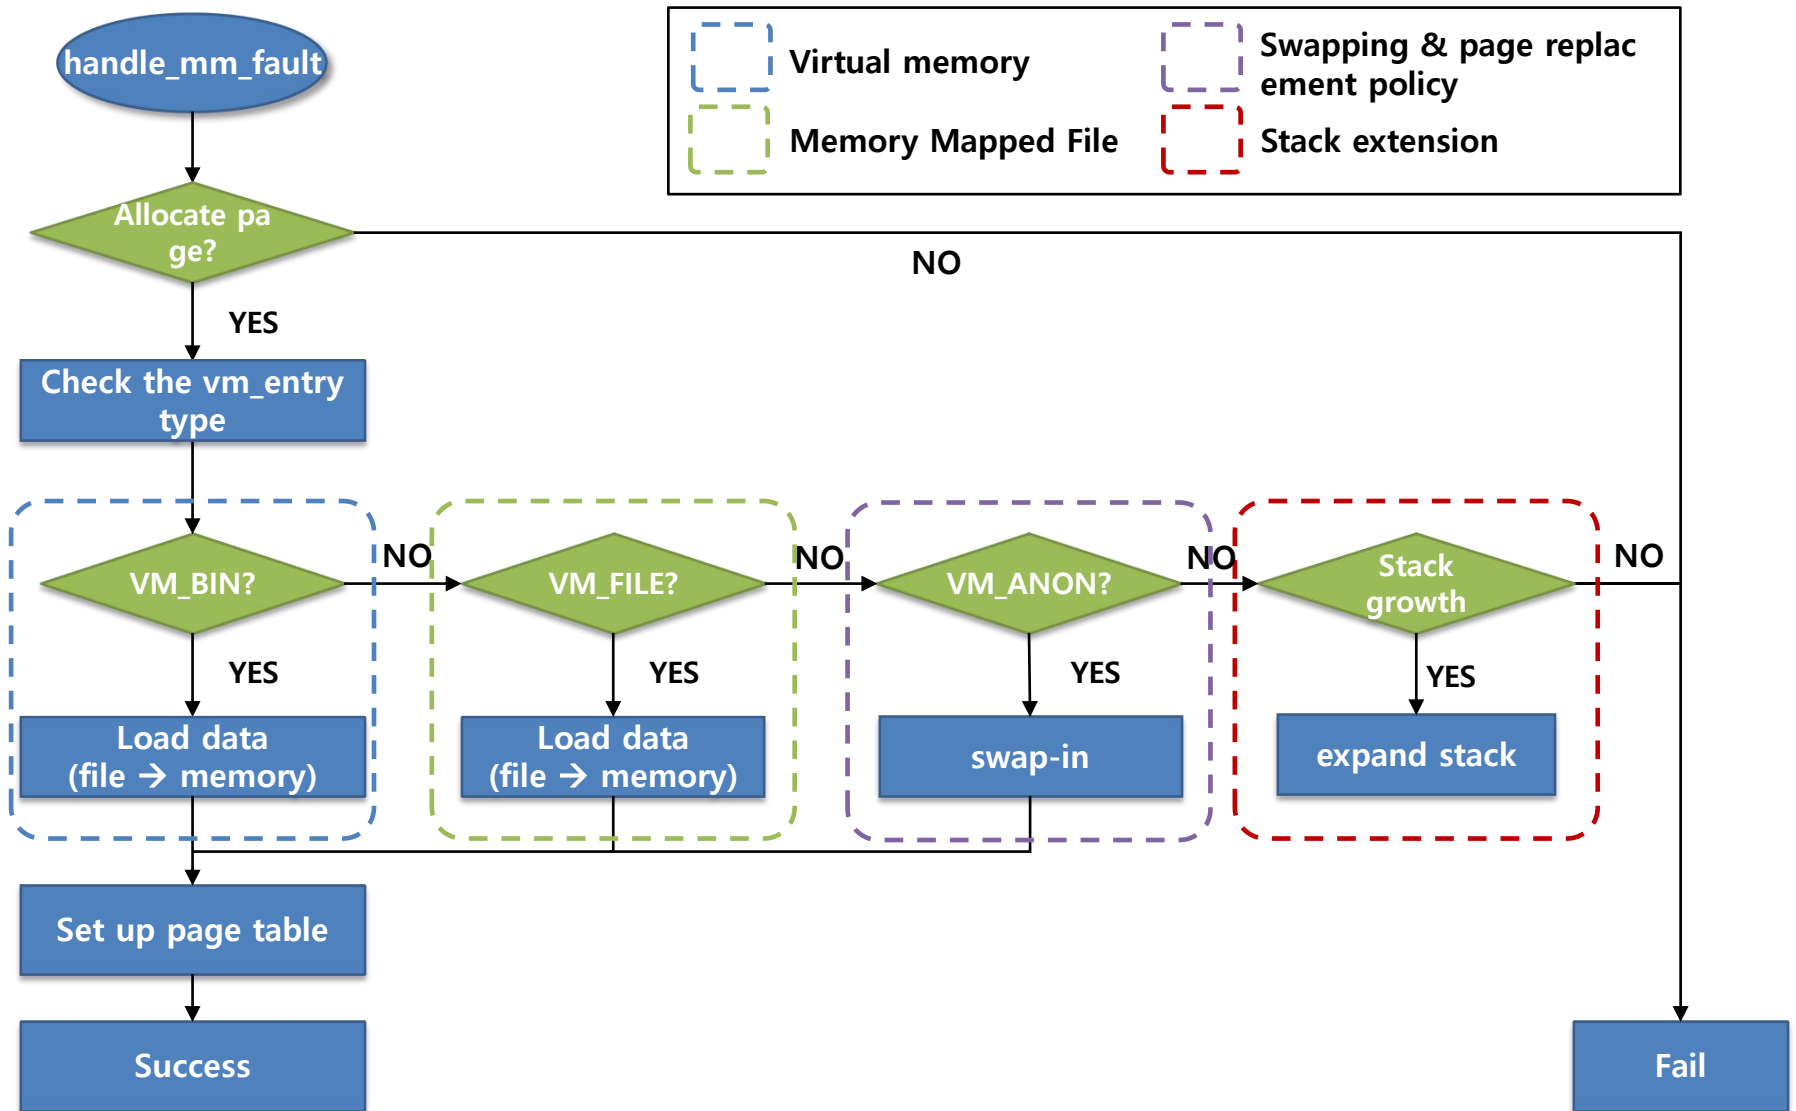
# Additional data structure and Functions to modify

- ▣ `struct mmap_file`
  - ◆ Data structure containing information from mapped files
  - ◆ mapping id
  - ◆ mapping file object
  - ◆ `mmap_file` list element
  - ◆ `vm_entry` list.
- ▣ `bool handle_mm_fault(struct vm_entry *vme)`
  - ◆ Load data if `vm_entry` type is `VM_FILE`
- ▣ `void process_exit(void)`
  - ◆ Release all `vm_entry` corresponding to `mapping_list` at the end of process.

# Managing mapped files



# Modify page fault handler for `mmap()`

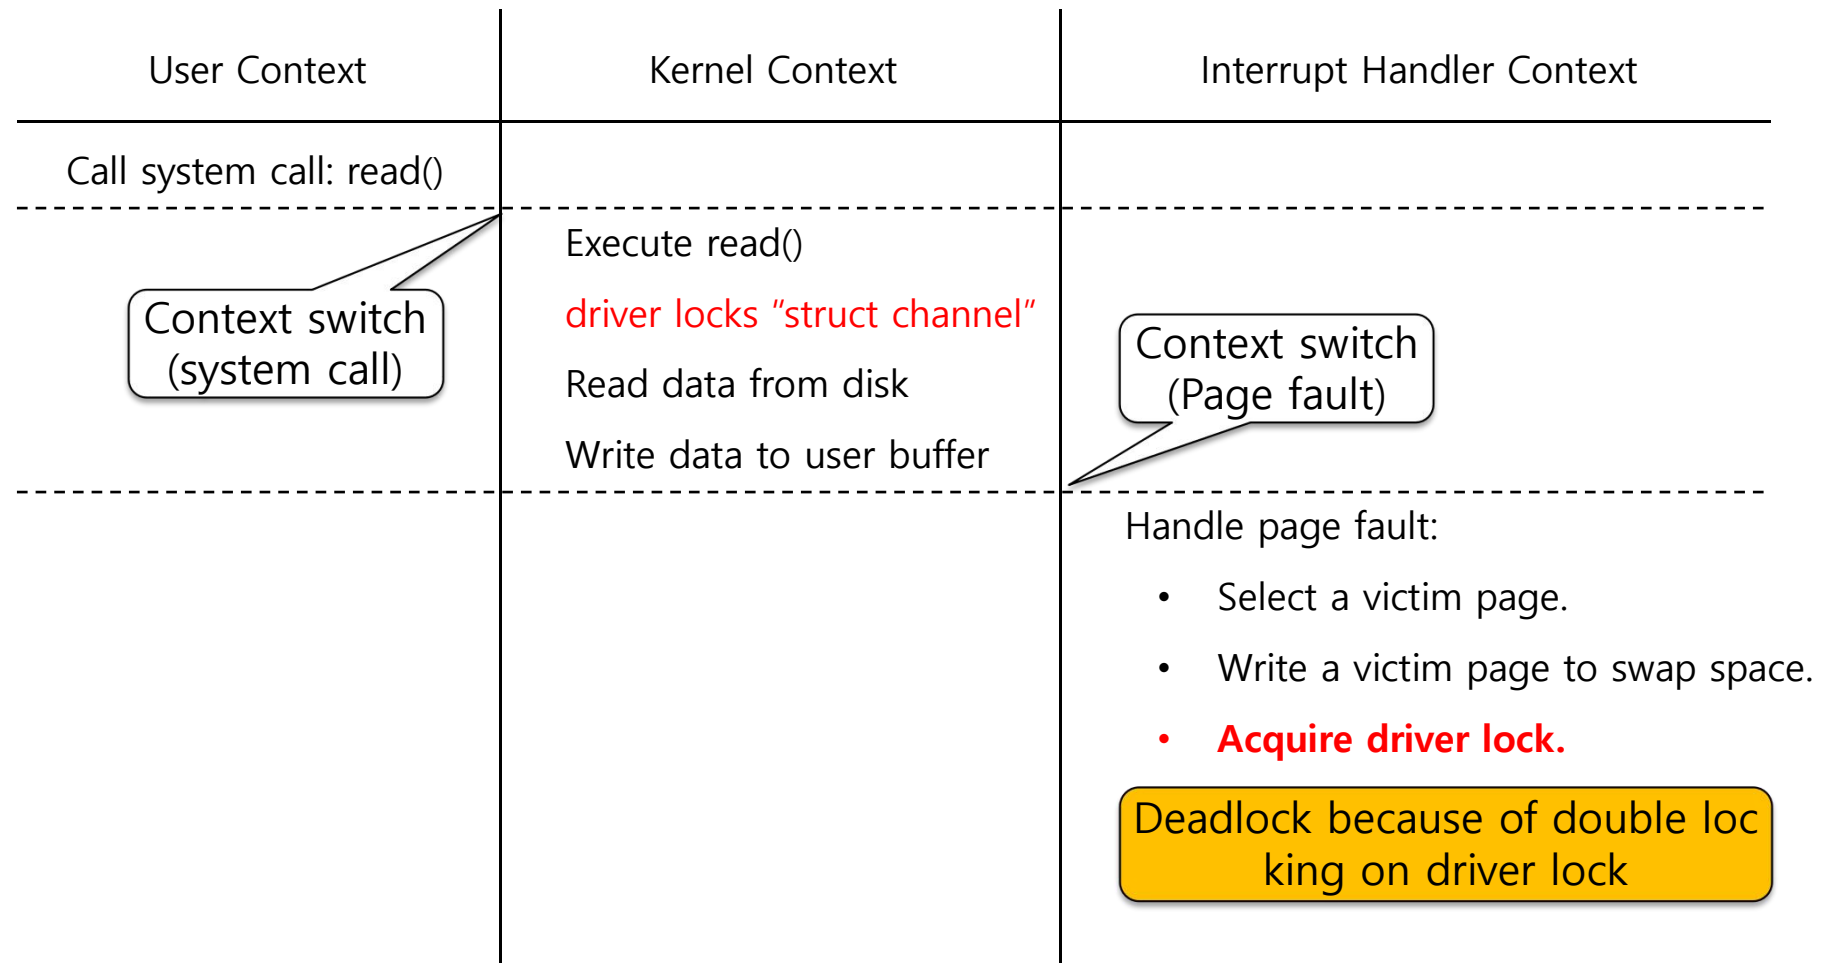


# Accessing User Address Space

---

# Why should page fault not occur in kernel code?

- ▣ A deadlock on kernel resource can occur.





# Pinning Page

- ▣ Prevent evicting the pages accessed during system call
- ▣ Define pinning flag about each physical page.
- ▣ On every system call,
  - ◆ Find the virtual page and pin the associated physical page.
  - ◆ After the system call returns and before the system call handler returns, unpin the pages
- ▣ On Swapping handler,
  - ◆ Do not select a pinned page as a victim.