

Operating System Lab

Part 4: Filesystem

KAIST EE

Youjip Won

Indexed and Extensible File

Extensible File

▣ Goal


- ◆ In original pintos, file size is fixed when it is created.
- ◆ In this project, we will modify pintos filesystem to change the file size dynamically. Maximum file size will be 8MB.

▣ Files to modify


- ◆ `pintos/src/filesys/inode.c`
- ◆ `pintos/src/filesys/inode.h`

How to allocate block when pintos create file


1. Creation of file A (Save start block address and length to inode)



2. Creation of file B (Save start block address and length to inode)



3. The size of file A can't be increased because of conflict with file B




After modification


1. Creation of file A (Save block addresses of all each blocks to inode)



2. Creation of file B (Save block addresses of all each blocks to inode)




3. Expanding the file A can be done by allocating a new block and by saving its address to inode of file A




Some details

1. The `seek()` does not change the file size nor does it allocate the blocks. It just updates `offset`.



2. When a `write()` is called and `offset` is larger than the file size, the file size is updated and blocks are allocated (Fill punch hole with zero).

After write ('A', "abc", size):



To Do's


1. Modify on-disk inode structure (`struct inode_disk`).
2. Modify code using on-disk inode.
 - ◆ Changing file offset to block address
 - `static block_sector_t byte_to_sector(const struct inode_disk *inode_disk, off_t pos)`
 - ◆ Creating new inode
 - `bool inode_create(block_sector_t sector, off_t length)`
 - ◆ Deleting an inode
 - `void inode_close(struct inode *inode)`
3. Handle extension of file.
 - ◆ `off_t inode_write_at(struct inode *inode, const void *buffer_, off_t size, off_t offset)`

On-disk inode in current pintos

- ▣ Fields for pointing blocks in current pintos
 - ◆ `start`: Start block address
 - ◆ `length`: Size of file (byte)
 - Fixed when the file is created.
 - All blocks should be continuous.

pintos/src/filesys/inode.c

```
struct inode_disk
{
    block_sector_t start; /* First data sector */
    off_t length; /* File size in bytes */
    unsigned magic; /* Magic number */
    uint32_t unused[125]; /* Not Used */
}
```





To Do 1 - Modify on-disk inode structure

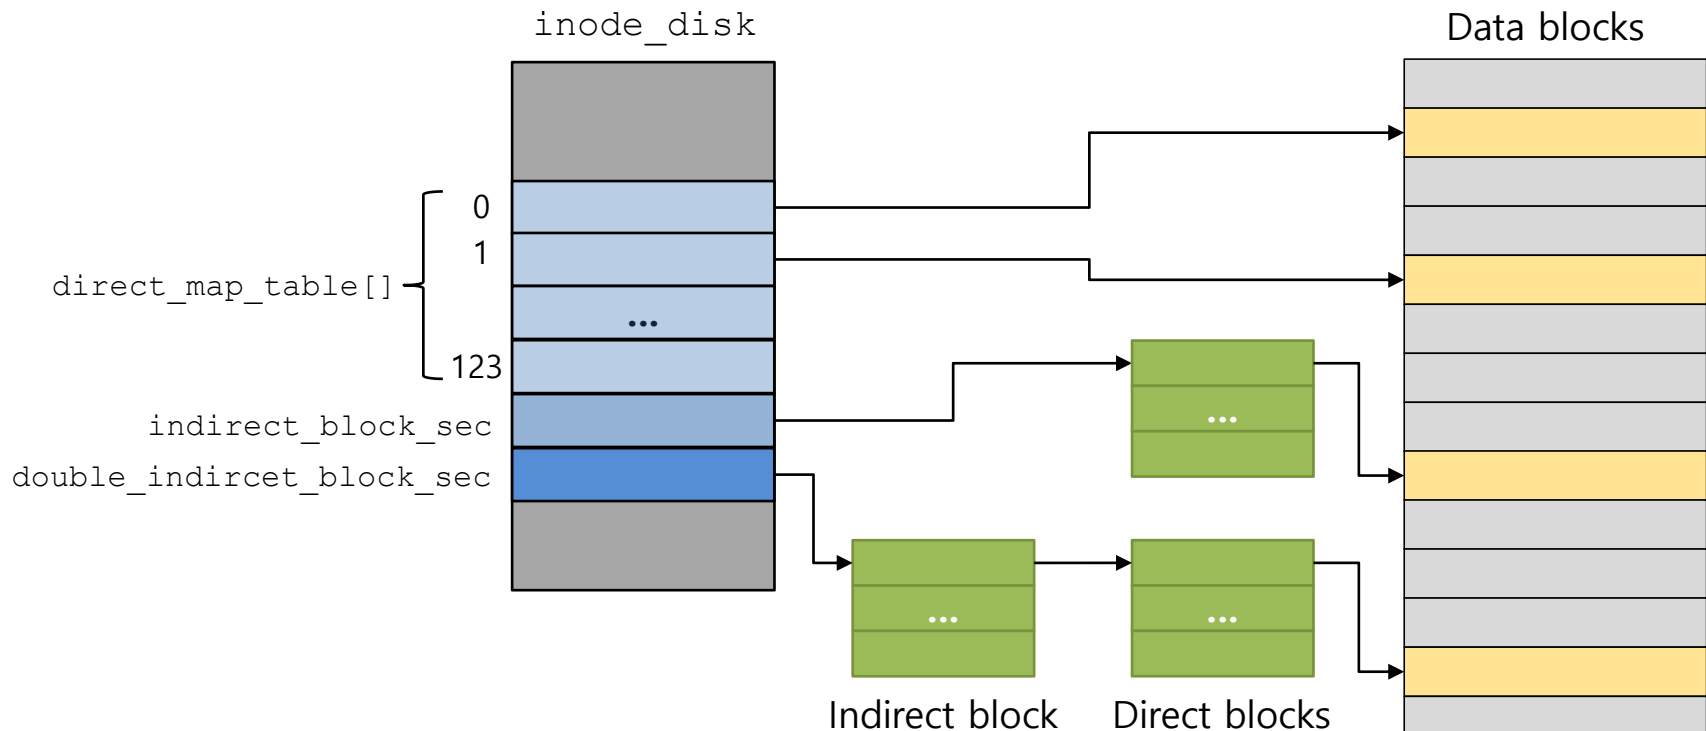
- ▣ Modify `struct inode_disk`.
 - ◆ Add block pointers of direct, indirect, double indirect.
 - ◆ We must maintain the size of "`struct inode_disk`" by adjusting the number of direct block pointers.

pintos/src/filesys/inode.c

```
struct inode_disk{  
    off_t length;           /* File size in bytes */  
    unsigned magic;         /* Magic number */  
    block_sector_t direct_map_table[DIRECT_BLOCK_ENTRIES];  
    block_sector_t indirect_block_sec;  
    block_sector_t double_indirect_block_sec;  
}
```



On-disk inode with block indexing



pintos/src/filesys/inode.c

```
struct inode_disk{
    off_t length;           /* File size in bytes */
    unsigned magic;         /* Magic number */
    block_sector_t direct_map_table[DIRECT_BLOCK_ENTRIES];
    block_sector_t indirect_block_sec;
    block_sector_t double_indirect_block_sec;
}
```

To Do 2 – compute the sector number from the file offset

- ▣ `static block_sector_t byte_to_sector(const struct inode_disk *inode_disk, off_t pos)`
 - ◆ Return block address associated an offset of inode.
- ▣ In original: just return sum of start and pos.
- ▣ Modify code in red box to use block indexing.

pintos/src/filesys/inode.c

```
static block_sector_t byte_to_sector (const struct inode *inode, off_t pos){
    ASSERT (inode != NULL);
    if (pos < inode->data.length)
        return inode->data.start + pos / BLOCK_SECTOR_SIZE;
    else
        return -1;
}
```

To Do 2 – Creating an inode

- `bool inode_create(block_sector_t sector, off_t length)`
 - ◆ Create file which have size of length.
- In original: Allocate contiguous blocks and save its start address.
- Modify code to save the block addresses of all blocks allocated.

pintos/src/filesys/inode.c

```
bool inode_create (block_sector_t sector, off_t length){
    ...
    disk_inode = calloc (1, sizeof *disk_inode);
    if (disk_inode != NULL){
        disk_inode->length = length;
        disk_inode->magic = INODE_MAGIC;
        if (free_map_allocate(sectors, &disk_inode->start)) {
            block_write(fs_device, sector, disk_inode);
            if (sectors > 0) {
                /* Fill file out by zero */
            }
            success = true;
        }
        free(disk_inode);
    }
    return success;
}
```

To Do 2 – Deleting an inode

- When it deletes an inode,
 - ◆ We have to deallocate all blocks inode have.
 - ◆ Add block deallocating code at `inode_close`.

`pintos/src/filesys/inode.c`

```
void inode_close (struct inode *inode){
    ...
    /* Release resources if this was the last opener. */
    if (--inode->open_cnt == 0){
        ...
        /* Deallocate blocks if removed. */
        if (inode->removed){
            /* Get on-disk inode structure by get_disk_inode() */
            /* Deallocate each blocks by free_inode_sectors() */
            /* Deallocate on-disk inode by free_map_release() */
        }
        ...
    }
```

Remove original code and add new code

To Do 3 – Handle extension of file

- When the file size changes,
 - ◆ Allocate a new block and update data block pointer in inode.
 - ◆ Fill the allocated blocks with zero.

pintos/src/filesys/inode.c

```
off_t inode_write_at (struct inode *inode, const void *buffer_,
                      off_t size, off_t offset) {
    ...
    /* Acquire some lock to avoid contention on inode */

    int old_length = disk_inode->length;
    int write_end = offset + size - 1;

    if (write_end > old_length - 1) {
        /* When size of file is updated, Update inode */
    }
    /* Release lock */

    while (size > 0) {
        ...
    }
}
```

Subdirectory

Overview of Subdirectories


▣ Main Goal

- ◆ Original pintos has only root directory but not the other subdirectories.
- ◆ We will implement subdirectory feature to make filesystem of pintos to have hierarchical structure.

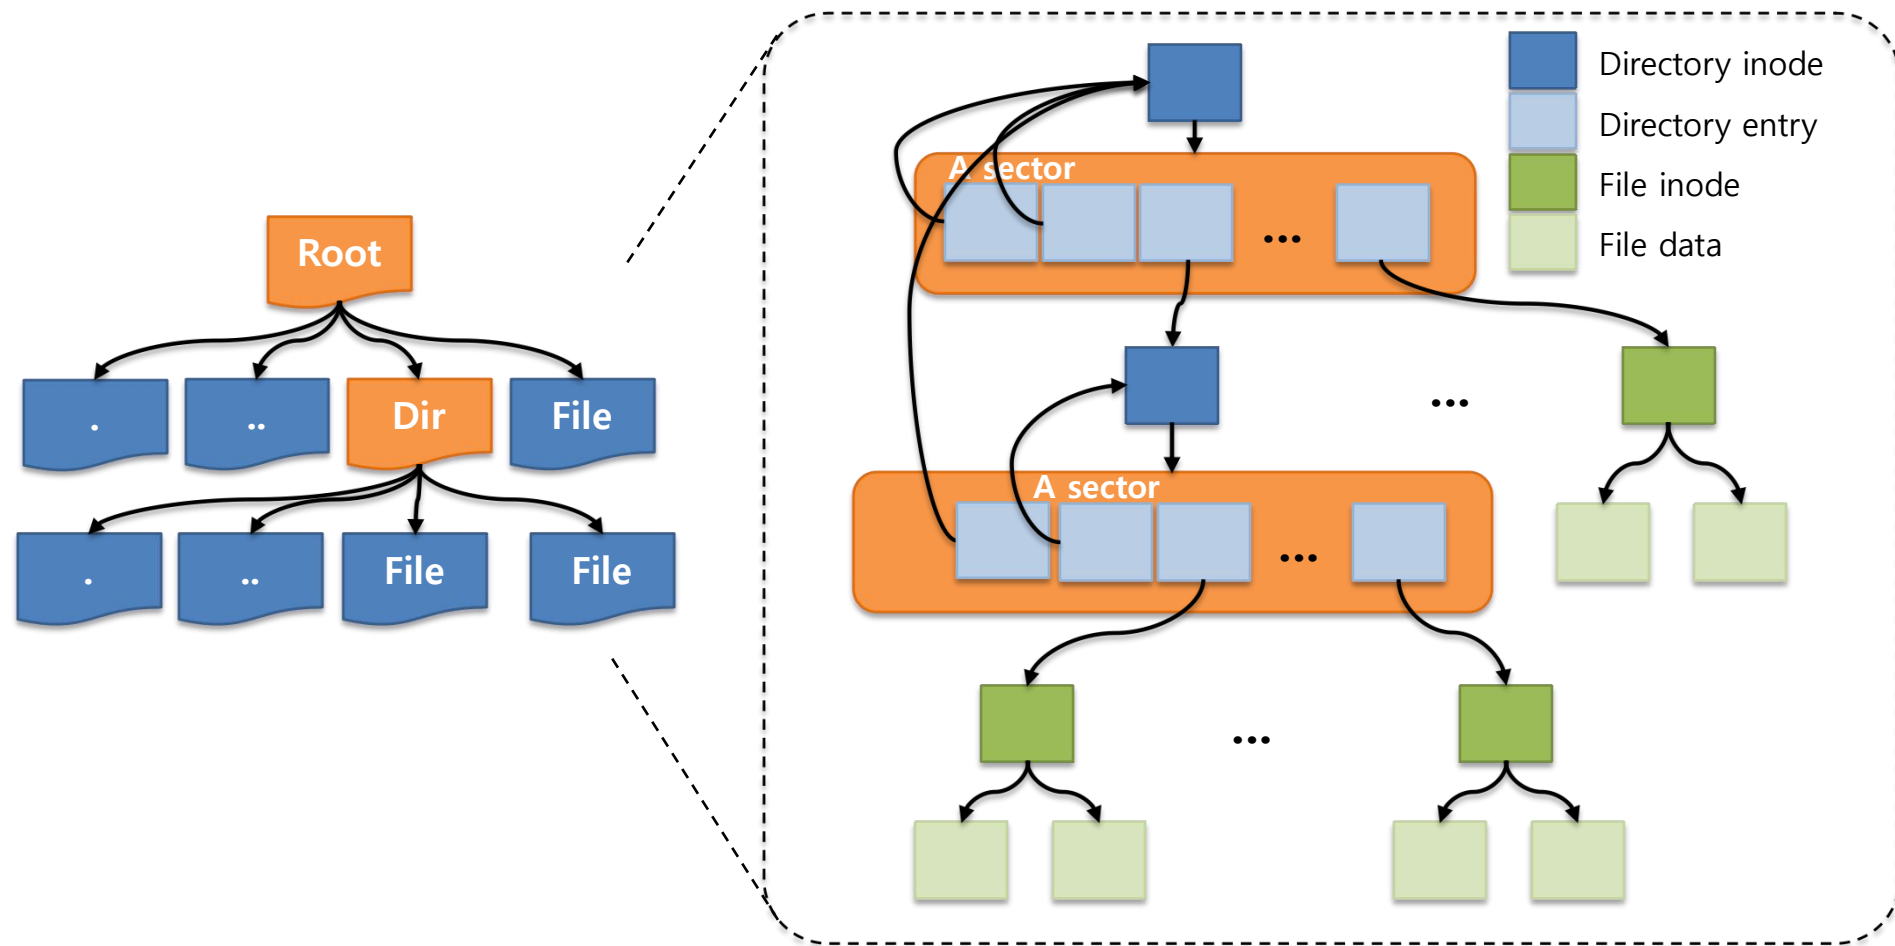
▣ Files to modify

- ◆ `pintos/src/filesys/inode.c`
- ◆ `pintos/src/filesys/filesys.*`
- ◆ `pintos/src/filesys/file.*`
- ◆ `pintos/src/filesys/directory.*`
- ◆ `pintos/src/userprog/syscall.c`

Directory structure in original pintos



Hierarchical directory structure



Requirements

- ▣ Implement Hierarchical directory structure.
 - ◆ Make directory entry can point to not only regular file but also directory.
 - ◆ Add directory entry '.' and '..'.
- ▣ Implement "current directory" for a thread.
 - ◆ Distinguish absolute path and relative path (Distinguisher is '/').
- ▣ Modify directory related functions.
 - ◆ `filesys_create()`, `filesys_open()`, `filesys_remove()`
- ▣ Add new directory related system calls.

Now, **there is a concept of path !!!**

To Do's

1. Add flag to indicate whether the inode is for regular file or for directory.
2. Define current directory pointer in `struct thread`.
3. Modify code for directory manipulation.
 1. Creating new file.
 2. Opening a file.
 3. Removing a file.
4. Add system calls for directory manipulation.
5. Add special directory entries: ".", "..".

To Do 1: Modify On-disk inode

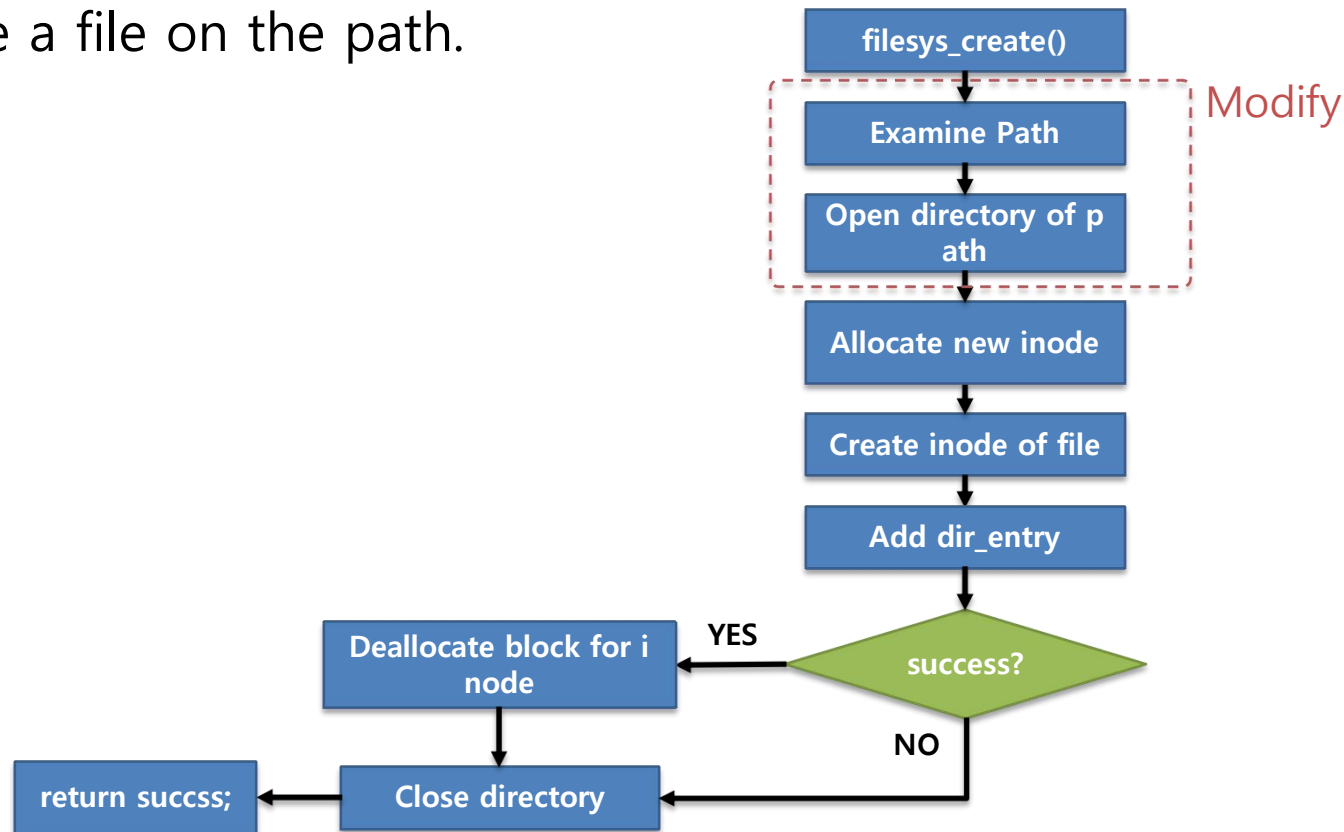
- ▣ Add a flag indicating if it is file or directory.
 - ◆ `struct inode_disk`
 - file: pintos/src/filesys/directory.c
 - Regular file(=0), directory(=1)
- ▣ When a file is created, we should set this flag properly.

To Do 2: Define current directory in `struct thread`.

- ▣ Add current directory pointer in thread.
 - ◆ `struct thread`
 - File: `pintos/src/threads/thread.h`
 - Add pointer to `struct dir` representing current directory.
- ▣ When a thread is created, the child thread inherits parent's current directory.

To Do 3: Modify algorithm of file creation

- ▣ Distinguish if it is absolute path or relative path.
- ▣ Find directory associated with path.
- ▣ Create a file on the path.



To Do 3: Modify algorithm of File creation (Cont.)

- ▣ `bool filesystem_create (const char *name, off_t initial_size)`
 - ◆ Original: Always create a file in the root directory.
 - ◆ After modification: Parse the path and create that file on that directory.
 - Distinguish absolute path and relative path and parse.
 - ◆ Add the code to set `is_dir` flag to 0 if it is regular file.
 - ◆ Add new directory entry to directory of path.

To Do 3: Opening a file

- ▣ `struct file *filesystem_open(const char *name)`
 - ◆ Original: Always find the file on the root directory.
 - ◆ After modification: Parse path, find the file on that directory, and open it.
 - Distinguish absolute path and relative path and parse.
 - When the path is absolute: Find from the root directory.
 - When the path is relative: Find from the current directory.

To Do 5: Removing a file

- ▣ `bool filesystem_remove(const char *name)`
 - ◆ Original: Always remove file from root directory.
 - ◆ After modification: Remove file from directory specified by path.
 - Distinguish absolute path and relative path and parse.
 - ◆ Check if in-memory of target file is for directory or regular file.
 - If it is directory, check it have files.
 - Remove only when directory is empty.
 - If it is file, just remove it.

To Do 6: Add system calls about directory

- ▣ `bool chdir(const char *dir)`
 - ◆ Change the current working directory of the process to `dir`.
 - ◆ Return true if successful, false on failure.
- ▣ `bool mkdir(const char *dir)`
 - ◆ Creates the directory named `dir`.
 - ◆ Returns true if successful, false on failure.
- ▣ `bool readdir(int fd, char *name)`
 - ◆ Reads a directory entry from file descriptor `fd`, which must represent a directory.
 - ◆ If successful, stores file name in `name` and return true.
 - ``.`` and `..`` should not be returned by `readdir`.
- ▣ `bool isdir(int fd)`
 - ◆ Returns true if `fd` represents a directory, false if it represents an ordinary file.
- ▣ `int inumber(int fd)`
 - ◆ Returns the inode number of the inode associated with `fd`.

To Do 7: Add special entries

- ▣ Special directory entries.
 - ◆ `'.'`: it represents itself.
 - ◆ `'..'`: it represents its parent directory.
- ▣ When a directory is created, special entries should be added.
 - ◆ even root directory created during format.
- ▣ If the user tries to remove them, system call should return fails.